

Concurrency Problem: multiple transactions working on the same shared variable, running concurrently

For correct execution, we need to impose exclusive access to the shared variable counter by Task1 and Task2.

Shared counter = 100;

Task1/Trans/Process/Thread counter = counter +10;	Task2/Trans/Process/Thread counter = counter +30;
--	--

Task1 and Task2 run concurrently. What are the possible values of counter after the end of Task1 and Task2?

Note: == means equals.

a) counter == 110 Sequence of actions T1: Reads counter == 100 T2: Reads counter == 100 T2: Writes counter == 100+30 T1: Writes counter == 100+10	b) counter == 130 Sequence of actions T1: Reads counter == 100 T2: Reads counter == 100 T1: Writes counter == 100+10 T2: Writes counter == 100+30	c) counter == 140; Sequence of actions T1: Reads counter == 100 T2: Reads counter == 100+10 T1: Writes counter == 110+30 T2: Writes counter == 110+30
--	--	--

TP monitor

The main function of a TP monitor is to **integrate other system components and manage resources**.

- 1.TP monitors manage the transfer of data between clients and servers
 - 2.Breaks down applications or code into transactions and ensures that all databases are updated properly
 - 3.It also takes appropriate actions if any error occurs
- Example: If A has just bought last product on a online shopping platform, but next second B is about to checkout the same item, such scenario needed to be managed (rolled back the B's operation) by TP monitor, to make DB's change is consistent.

TP monitor service

Heterogeneity: If the application needs access to different DB systems, local ACID properties of individual DB systems is not sufficient. Local TP monitor needs to interact with other TP monitors to ensure the **overall ACID** property. (Two-phase commit in distributed DB)

Control: If the application communicates with other remote processes, the local TP monitor should maintain the communication status among the processes to be able to recover from a crash.

Terminal management: Since many terminals run client software, the TP monitor should provide appropriate ACID property between the client and the server processes.

Presentation service: this is similar to terminal management in the sense it has to deal with different presentation (user interface) software -- e.g. X-windows

Context management: E.g. maintaining the sessions etc.

Start/Restart: There is no difference between start and restart in TP based system.

Concurrency Problem

– **Spin locks (需要 hard ware support; using atomic lock/unlock instructions)** – most commonly used

Pros: All modern processors do support some form of spin locks;

Executed using atomic machine instructions such as test and set or compare and swap; Need hardware support

Use busy waiting; Does not depend on number of processes; Very efficient for low lock contentions – all DB systems use them

Implementation of Atomic operation: **compare and swap** in spin lock for exclusive access

code may have lost values

temp = counter +1; //unsafe to increment a shared counter
counter = temp; //this assignment may suffer a lost update

Implementation of Atomic operations: test and set

```
T1
/*acquire lock*/ while (!testAndSet( &lock );
/*Xlock granted*/
/*exclusive access for T1;
counter = counter+1;
/* release lock*/
lock = 1;

T2
/*acquire lock*/ while (!testAndSet( &lock );
/*Xlock granted*/
/*exclusive access for T2;
counter = counter+1;
/* release lock*/
lock = 1;
```

Why we prefer CAS, rather than TAS? **ANS:** 1.Pro for CAS: it ask for exclusive lock only when 'write' operation is required.

So if in the scenario where multiple transactions only execute reading operation simultaneously (browsing products on online shopping platform) but not do any writing operation on it, using TAS will prevent any other transaction from reading the resources as well. 2.For reading and writing if we use TAS, which means one transaction has to start and completely finish its operations on the shared variable, then only the second transaction will be able to start (busy waiting). 3. However, if both transactions are just reading that value, reading does not conflict what happening for other transaction if both of them just reading and not making any changes to the share variable. CAS is allow multiple transaction to run concurrently for that reading part.

Concurrency Problem solution

Different ways for concurrency control

– **Dekker's algorithm (using code)** Pros:Dekker's algorithm needs almost no hardware support although it needs atomic reads and writes to main memory. That is exclusive access of one time cycle of memory access time!

Cons: The code is very complicated to implement if more than two transactions/process are involved; Harder to understand the algorithm for more than two process takes lot of storage space; Uses busy waiting: If resource has been locked by T1, then T2 cannot do anything but just keeps waiting in the loop for the resource to be released, hampers the overall performance of transaction; **Efficient if the lock contention (that is frequency of access to the locks) is low.**

Dekker's algorithm
int c1, c2, turn = 1; /* global variable */

T1 { some code T1} /* T1 wants exclusive access to the resource and we assume initially c1 = 0 */ c1 = 1; turn = 2; repeat until { c2 == 0 or turn == 1 } /* Start of exclusive access to the shared resource (successfully changed variables) */ use the resource counter = counter+1; /* release the resource */ c1 = 0; (some other code of T1)	T2 { some code T2} /* T2 wants exclusive access to the resource and we assume initially c2 = 0 */ c2 = 1; turn = 1; repeat until { c1 == 0 or turn == 2 } /* Start of exclusive access to the shared resource */ use the resource counter = counter+1; /* release the resource */ c2 = 0; (some other code of T2)
---	---

– **OS supported primitives (through interruption call)** - expensive, independent of number of processes, machine independent

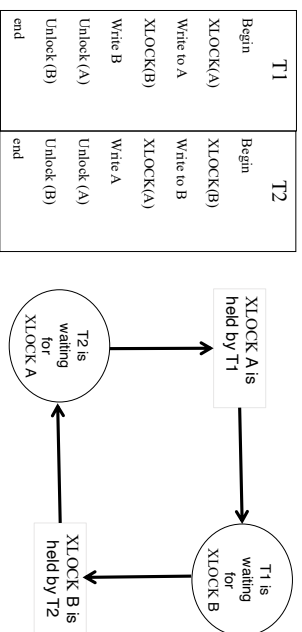
Convoy avoiding semaphore

The **Semaphores** implementation may result a long list of waiting processes (convoy), if one or couple of them are long-running processors then the one that make request later has to wait for longer time in the queue.

Solution: once one transaction is done, the lock on the shared variable gets released, every process in the waiting list is waken up, then they all re-execute the routine for acquiring semaphore (recall the resources again), which leads to the order of waiting transaction not be the same as the queue before. (i.e., the order get reshuffled), reducing the probability of long waiting queue occurs, and all current transaction can have another chance to ask disk for resources.

Deadlocks

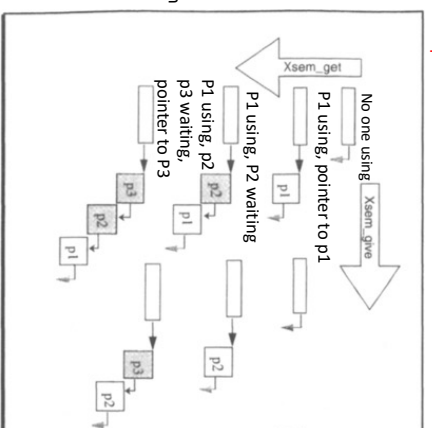
- In a deadlock, each process in the deadlock is waiting for another member to release the resources it wants.



Spin lock – Compare and Swap: improve overall # of operations that are being executed in the system=> run more operations concurrently.

```
boolean cs(int *cell, int *old, int *new)
/* the following is executed atomically*/
if (*cell == *old) { *cell = *new; return TRUE;}
else { *old = *cell; return FALSE; }
}
```

Semaphores



Deadlock - Solutions

Have enough resources so that no waiting occurs – not practical; Do not allow a process to wait, simply rollback after a certain time. This can create live locks which are worse than deadlocks; Linearly order the resources and request of resources should follow this order, i.e., a transaction after requesting i^{th} resource can request j^{th} resource if $j > i$. This type of allocation guarantees no cyclic dependencies among the transactions. (在全局保持这个顺序)

1. Pre-declare all necessary resources and allocate in a single request.
2. Periodically check the resource dependency graph for cycles. If a cycle exists – rollback (i.e., terminate) one or more transaction to eliminate cycles (deadlocks). The chosen transactions should be cheap (e.g., they have not consumed too many resources).
3. Allow waiting for a maximum time on a lock then force Rollback. Many successful systems (IBM, Tandem) have chosen this approach.

Q Many distributed database systems maintain only local dependency graphs and use time outs for global deadlocks.

Pros: 1. simple and easy to implement: we don't have to monitor state of transaction instantly for detecting deadlock cycles, which may cause high overhead. 2. Scalability: For large system with many concurrent transactions, deadlock detection algorithm is costly and complex. Timeout based strategy provides a scalable solution, allowing system to handle a large # of trans without tracking complex dependency (adapt to many system)

Cons: 1. Potential for Unnecessary Rollbacks: This approach can cause transactions to roll back even if they are not involved in a deadlock, simply because they exceeded the wait time. This can lead to wasted work and reduced efficiency, especially in cases of high contention. 2. Difficult to Set Optimal Timeout: If the timeout is set too short, transactions may frequently roll back unnecessarily, if it's too long, actual deadlocks may take too long to resolve, leading to resource contention and delayed processing. 3. Increased System Overhead Due to Rollbacks: Frequent rollbacks can increase the system's workload, as it has to repeatedly restart transactions that could have otherwise completed, impacting overall throughput and response times.

Overall, while simple and resource-efficient, the timeout-based strategy can lead to inefficiency and performance degradation in high-concurrency environments or when transactions are long-running.