



COMP90050 Advanced Database Systems

Semester 2, 2024

Lecturer: Farhana Choudhury (PhD)
Live lecture – Week 5





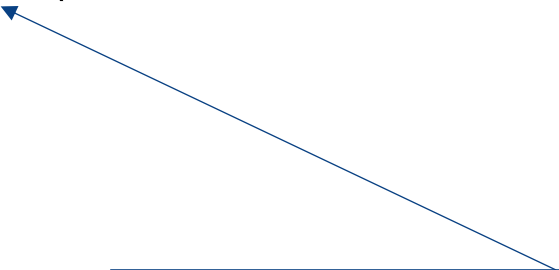
Continuation from last week's live lecture



Flat Transaction

Everything inside BEGIN WORK and COMMIT WORK is at the same level; that is, the transaction will either survive together with everything else (commit), or it will be rolled back with everything else (abort)

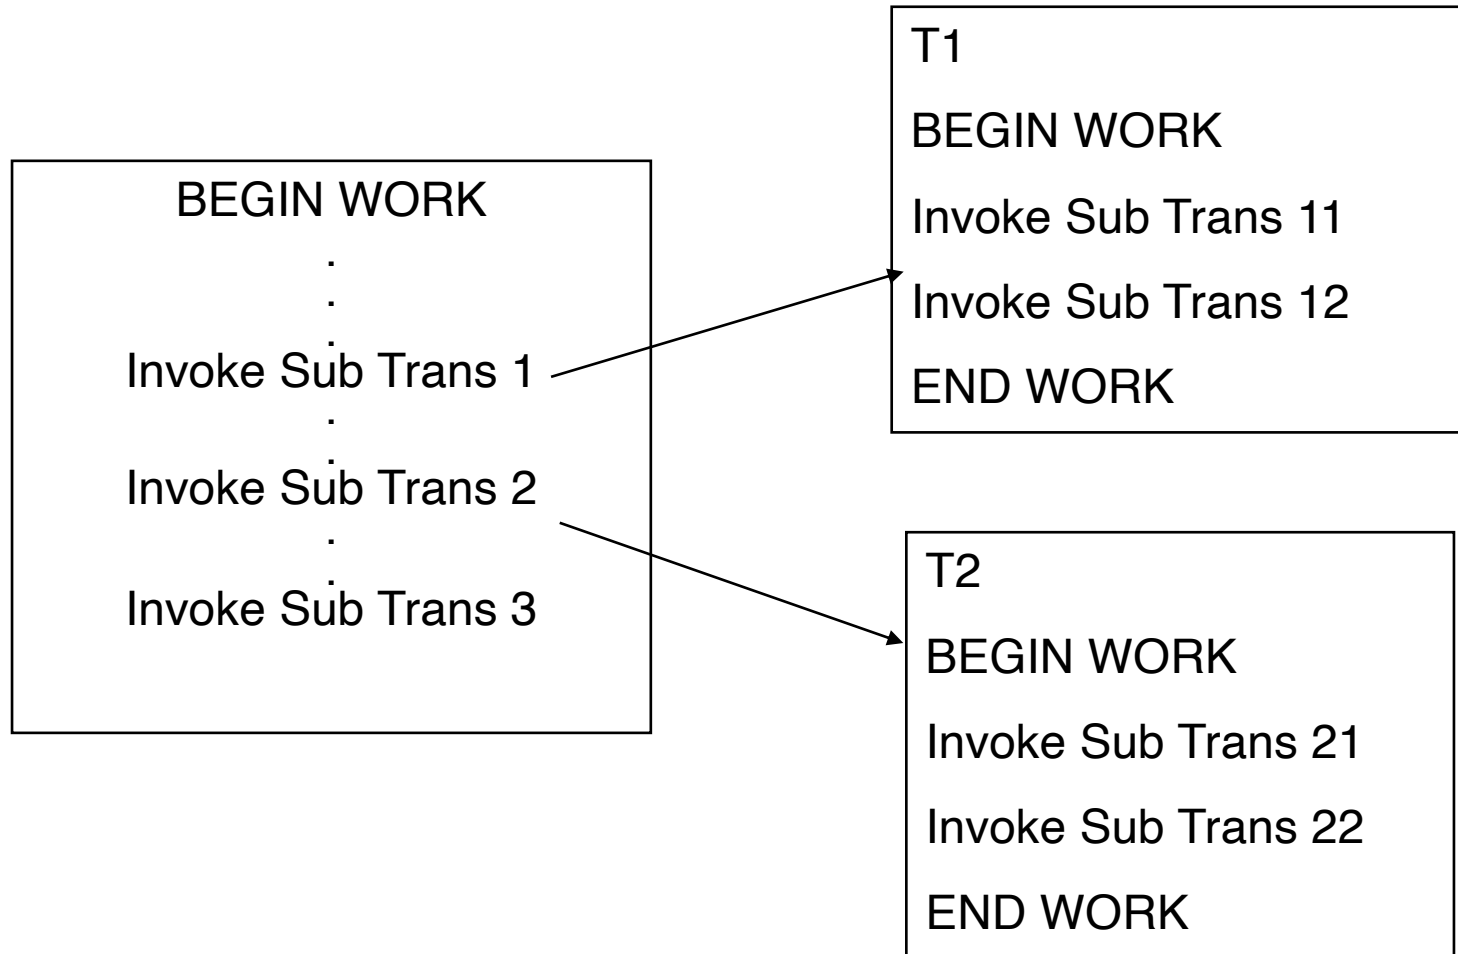
```
exec sql BEGIN WORK;  
    AccBalance = DodebitCredit(BranchId, TellerId, AcclId, delta);  
    send output msg;  
exec sql COMMIT WORK;
```



Can be a very long running transaction
with many operations

Limitations of Flat Transactions?

Nested Transactions





Nested Transactions

Commit rule

- A subtransaction can either commit or abort, however, **commit cannot take place unless the parent itself commits.**
- Subtransactions have A, C, and I properties but not D property unless all its ancestors commit.
- Commit of a sub transaction makes its results available only to its parents.

Roll back Rules

If a subtransaction rolls back, all its children are forced to roll back.

Visibility Rules

Changes made by a subtransaction are visible to the parent only when the subtransaction commits. All objects of parent are visible to its children.

Implication of this is that the **parent should not modify objects while children are accessing them.** This is not a problem as parent does not run in parallel with its children.

Advantages of nested transactions?

Time for a poll - [Pollev.com/farhanachoud585](https://pollev.com/farhanachoud585)

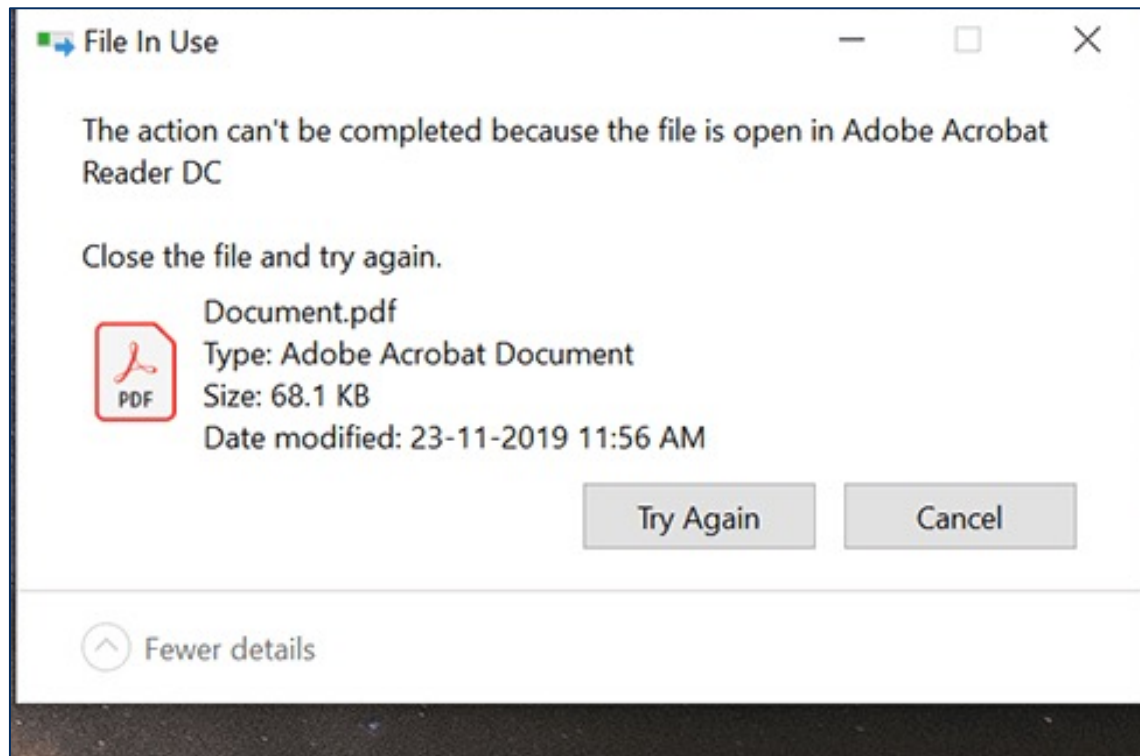


Book chapter: 4.7



This week's contents

Concurrency problem



Concurrent transactions can cause issues in Database – need concurrency control



Concurrency Problems

Transaction 1	Transaction 2	Account balance = 100;
balance = balance + 50;	balance = balance - 80;	

Both transactions are running concurrently. What are the possible outcomes if no concurrency control in place? Which ones are correct sequence of actions?

Write by a transaction gets
lost/overwritten

a) balance == 150

Sequence of actions

T1: Reads balance == 100

T2: Reads balance == 100

T2: Writes balance == 100-80

T1: Writes balance == 100+50

← not desirable
→ get extra money

b) balance == 20

Sequence of actions

T1: Reads balance == 100

T2: Reads balance == 100

T1: Writes balance == 100+50

T2: Writes balance == 100-80

← not desirable
loses money
can't last write not overwrite the first write.

c) balance == 70;

Sequence of actions

T1: Reads balance == 100

T1: Writes balance == 100+50

T2: Reads balance == 150

T2: Writes balance == 150-80

d) balance == 70;

Sequence of actions

T2: Reads balance == 100

T2: Writes balance == 100-80

T1: Reads balance == 20

T1: Writes balance == 20+50

*with a. b 不对: 我可以先 withdraw 再 deposit, 或先 deposit 再 withdraw \Rightarrow 但 These two transaction get intertwined \Rightarrow They both read \hookrightarrow at the same time, one is making change the other one is unaware of change that we made \rightarrow there is no concurrency control.

How two transaction should work on the same value, so that one does not impact the other?

one solution: no interleaving of actions \leftarrow

another solution: $\left\{ \begin{array}{l} \text{先存再取} \\ \text{先取再存} \end{array} \right.$



Concurrency Problems

Transaction 1	Transaction 2	Account balance = <u>100</u> ;
balance = balance + 50;	balance = balance - 110 ;	

Both transactions are running concurrently. What are the possible outcomes if no concurrency control in place? Which ones are correct sequence of actions?
(assume overdrawn not allowed)

c) balance == 40;

Sequence of actions

T1: Reads balance == 100

T1: Writes balance == 100+50

T2: Reads balance == 150

T2: Writes balance == 150-**110**

d) balance == 150;

Sequence of actions

T2: Reads balance == 100

T2: Writes balance == 100-110

(message on overdrawn denied)

T1: Reads balance == 100

T1: Writes balance == 100+50

Try withdrawing money
after some time

★ both of sequences are correct, although the balance values are different ★



Concurrency Control

目的 ① ②

- To resolve conflicts,
- To preserve database consistency

Different ways for concurrency control

- ...
- **Spin locks (using atomic lock/unlock instructions)** – most commonly used

General purpose locks - semaphores

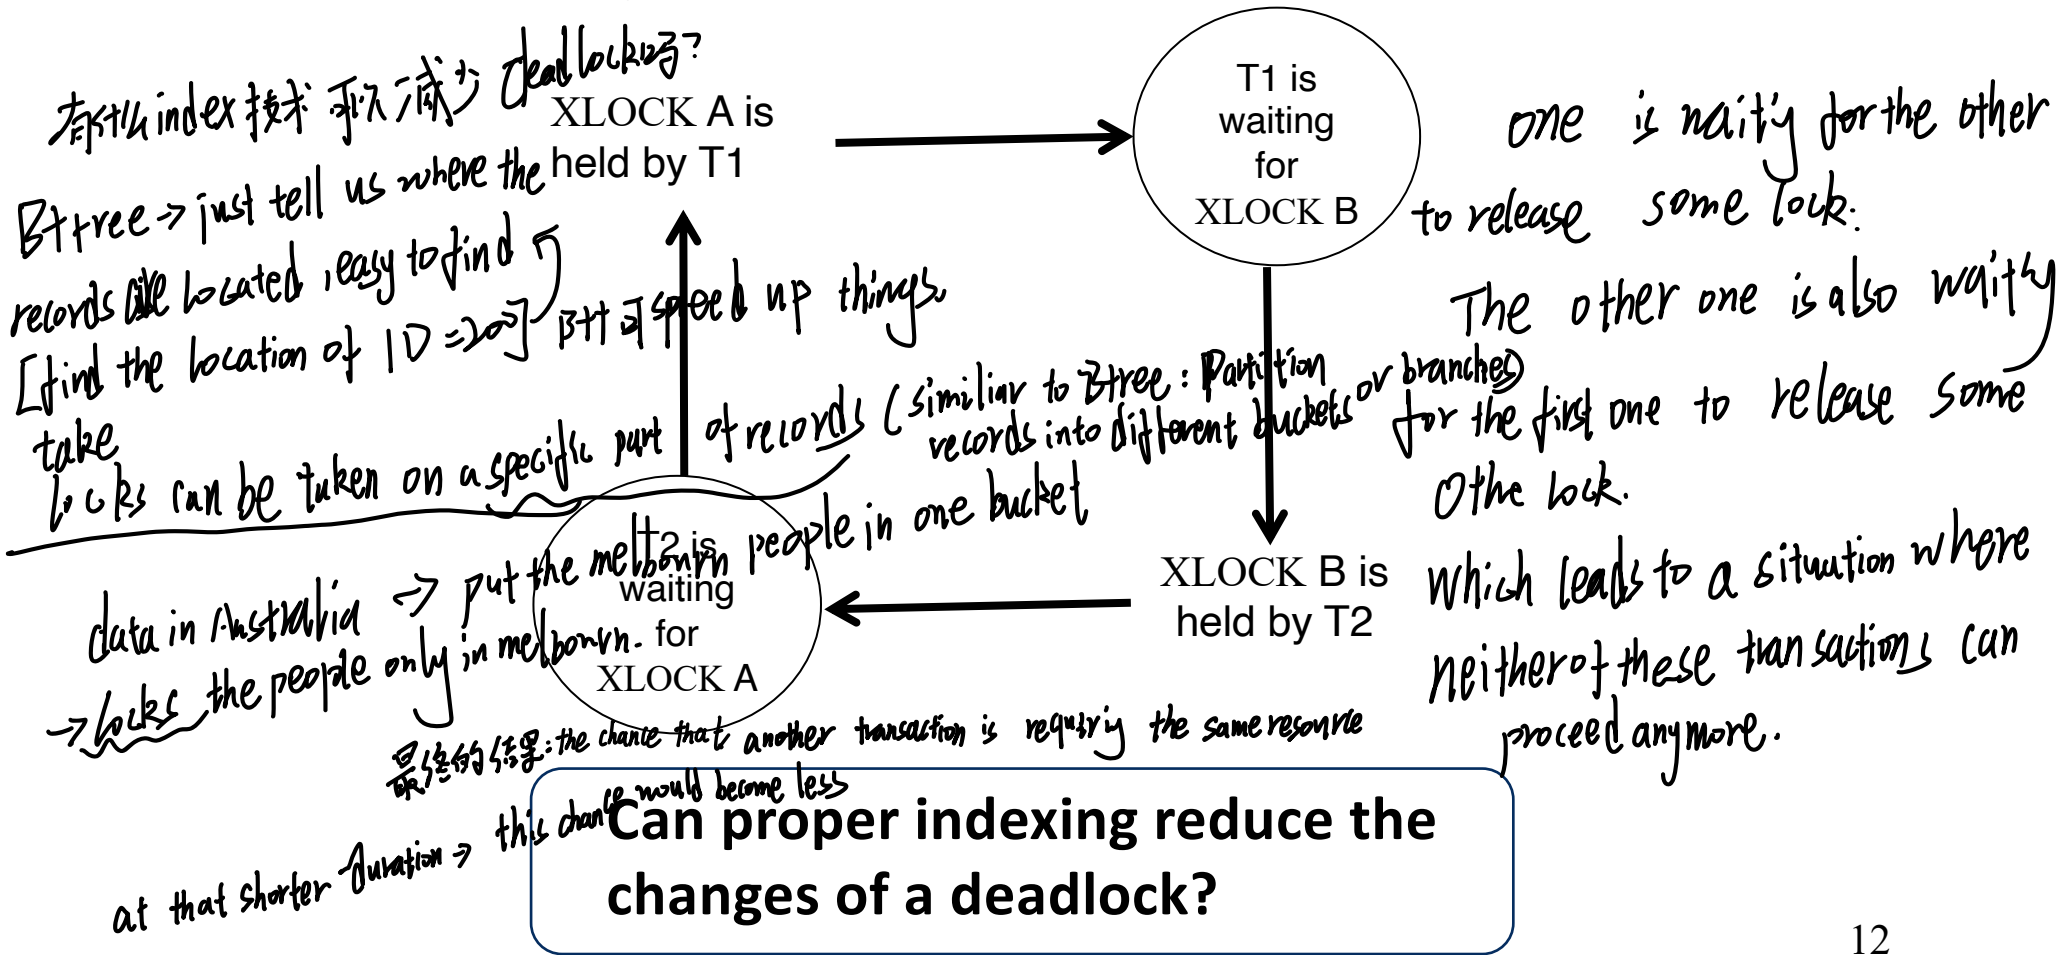


who came first who got released first.



Deadlocks ...

Deadlocks are rare, however, they do occur and the database has to deal with them when they occur



★ indexing can make query faster → the amount of time a lock will be kept that duration will become shorter ★



Deadlock avoidance/mitigation

- Pre-declare all necessary resources and allocate in a single request.
- Periodically check the resource dependency graph for cycles. If a cycle exists - rollback (i.e., terminate) one or more transaction to eliminate cycles (deadlocks). The chosen transactions should be cheap (e.g., they have not consumed too many resources).

• Allow waiting for a maximum time on a lock then force Rollback. Many successful systems (IBM, Tandem) have chosen this approach.

→ most commonly used technique → if a transaction is waiting for a long time → force one transaction to quit/rollback.

- Many distributed database systems maintain only local dependency graphs and use time outs for global deadlocks.



Concurrency problem

Multiple concurrently running transactions may cause conflicts

- **Still we try to allow concurrent runs as much as possible for a better performance, while avoiding conflicts as much as possible**

Topics for next two weeks



Core Concepts of Database management system

