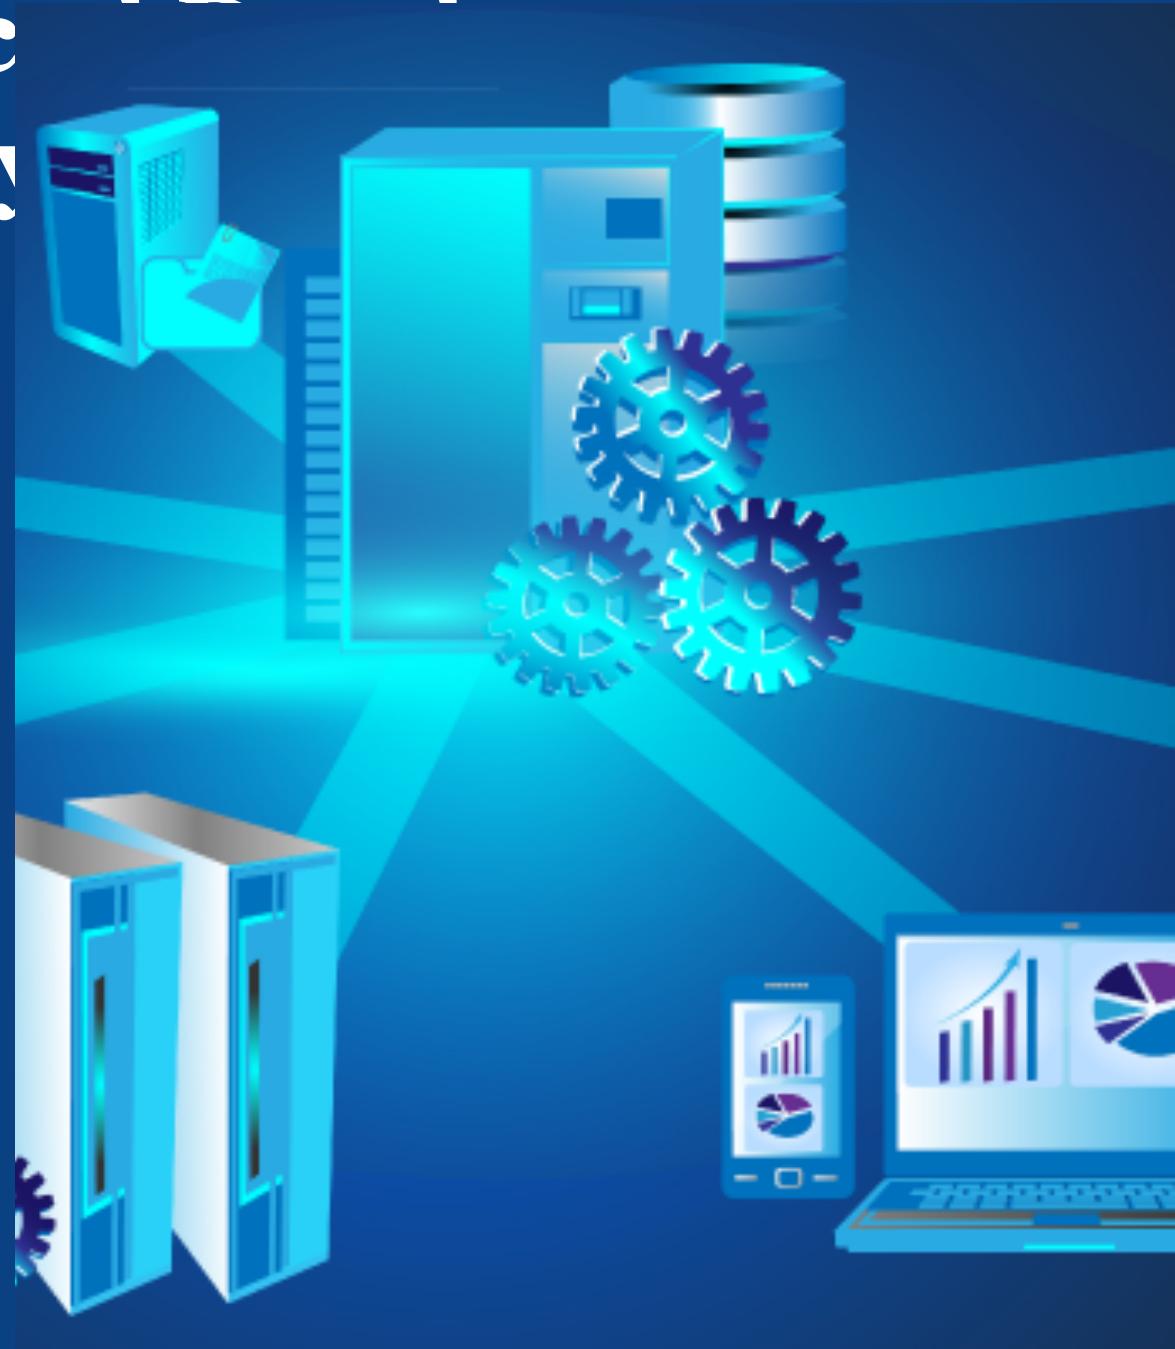




COMP90050:

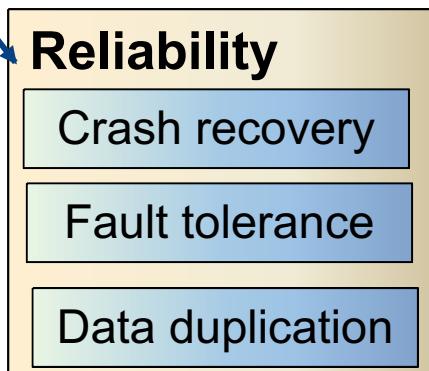
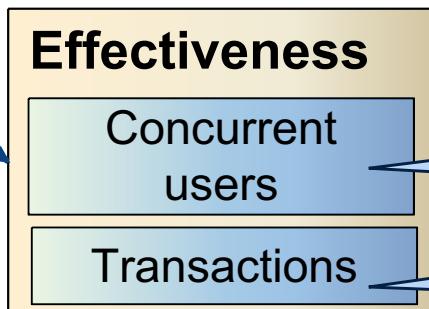
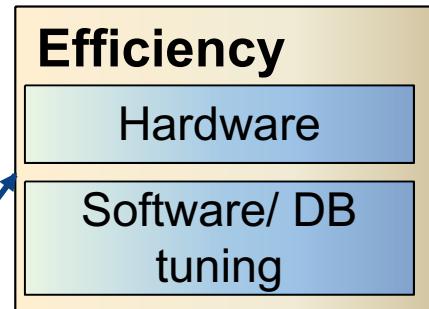
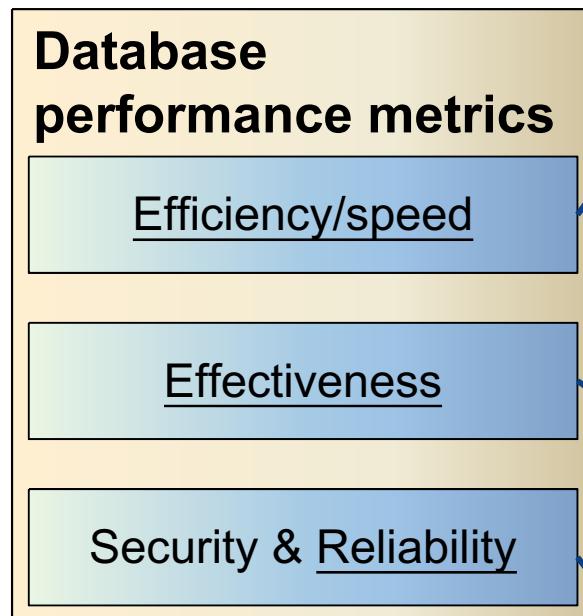
Advanced Systems

Lecturer: Farhana Choudhury (PhD)
Transaction concepts and TPM
Week 4





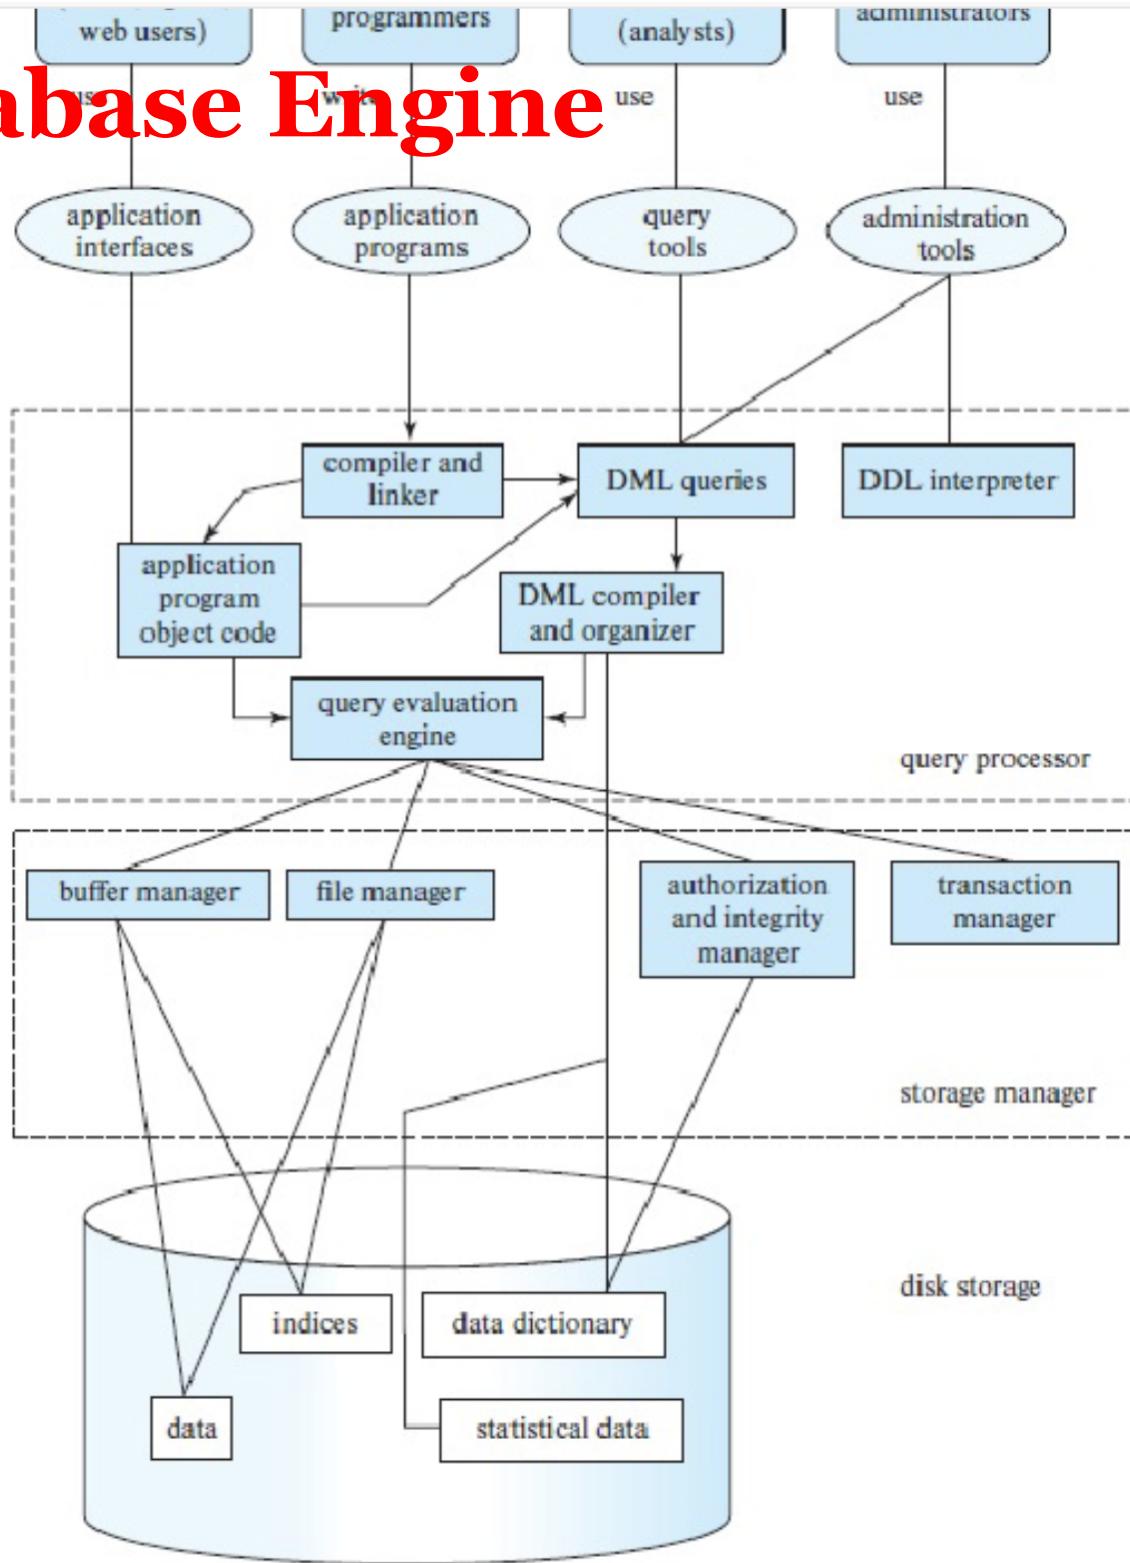
Core Concepts of Database management system



Users reading and writing over the same data

Required tasks are all done together

Database Engine





Topics

- Transaction concepts
- Different types of transactions
- Concurrent use of data and its issues
- Isolation concepts - ensuring that concurrent transactions leaves the DB in the same state as if they were executed separately
- How to achieve isolation while keeping good efficiency



Database Transactions

Transaction - A unit of work in a database

- A transaction can have any number and type of operations in it
- Either happens as a whole or not
- Transactions ideally have four properties, commonly known as ACID properties



Transaction models

ACID (Atomicity, Consistency, Isolation, Durability) properties:

Atomicity - All changes to data are performed as if they are a single operation. That is, all the changes are performed, or none of them are.

Example – A transaction that (i) subtracts \$100 if balance >100 (ii) deposits \$100 to another account
↗ (若第一件不發生，第二件發生，則結果是23)

↗ both actions will either happen together or none will happen) ↗ ↗

why bank prefer to ~~transfer~~ ACID.



Transaction models

ACID (Atomicity, Consistency, Isolation, Durability) properties:

Consistency - Data is in a ‘consistent’ state when a transaction starts and when it ends – in other words, any data written to the database must be valid according to all defined rules (e.g., no duplicate student ID, no negative fund transfer, etc.)

- What is ‘consistent’ - depends on the application and context constraints
- It is not easily computable in general
- Only restricted type of consistency can be guaranteed, e.g. serializable transactions which will be discussed later.



Transaction models ...

ACID (Atomicity, Consistency, Isolation, Durability) properties:

Isolation- transaction are executed as if it is the only one in the system.

- For example, in an application that transfers funds from one account to another, the isolation ensures that another transaction sees the transferred funds in one account or the other, but not in both, nor in neither.

Durability- the system should tolerate system failures and any **committed updates** should not be lost.



Transaction models (cont.)

Types of Actions

- **Unprotected actions** - no ACID property
- **Protected actions** - these actions are not externalised before they are completely done. These actions are controlled and can be rolled back if required. These have ACID property.
- **Real actions** - these are real physical actions once performed cannot be undone. In many situations, atomicity is not possible with real actions (e.g., firing two rockets as a single atomic action)



Let's look at a simple example that will be used to understand different types of transactions ...



Embedded SQL example in C

(Open Database Connectivity)

```
int main()
{
    exec sql INCLUDE SQLCA; /*SQL Communication Area*/
    exec sql BEGIN DECLARE SECTION;
        /* The following variables are used for communicating
           between SQL and C */

        int OrderID; /* Employee ID (from user) */
        int CustID; /* Retrieved customer ID */
        char SalesPerson[10] /* Retrieved salesperson name */
        char Status[6] /* Retrieved order status */

    exec sql END DECLARE SECTION;

    /* Set up error processing */
    exec sql WHENEVER SQLERROR GOTO query_error;
    exec sql WHENEVER NOT FOUND GOTO bad_number;
```



```
/* Prompt the user for order number */
    printf ("Enter order number: ");
    scanf_s("%d", &OrderID);
/* Execute the SQL query */
    exec sql SELECT CustID, SalesPerson, Status
                  FROM Orders
      WHERE OrderID = :OrderID // ":" indicates to refer to C variable
            INTO :CustID, :SalesPerson, :Status;
/* Display the results */
    printf ("Customer number: %d\n", CustID);
    printf ("Salesperson: %s\n", SalesPerson);
    printf ("Status: %s\n", Status);
    exit();
query_error:
    printf ("SQL error: %ld\n", sqlca->sqlcode); exit();
bad_number:
    printf ("Invalid order number.\n"); exit(); }
```



Some concepts...

Host Variables

Declared in a section enclosed by the BEGIN DECLARE SECTION and END DECLARE SECTION. While accessing these variables, they are prefixed by a colon “:”. The colon is essential to distinguish between host variables and database objects (for example tables and columns).

Data Types

The data types supported by a DBMS and a host language can be quite different. Host variables play a dual role:

- Host variables are program variables, declared and manipulated by host language statements, and
- they are used in embedded SQL to retrieve database data.

If there is no host language type corresponding to a DBMS data type, DBMS automatically converts the data. So, the host variable types must be chosen carefully.



Error Handling

The DBMS reports run-time errors to the applications program through an SQL Communications Area (SQLCA) by INCLUDE SQLCA. The WHENEVER...GOTO statement tells the pre-processor to generate error-handling code to process errors returned by the DBMS.

Singleton SELECT

The statement used to return the data is a singleton SELECT statement; that is, *it returns only a single row of data*. Therefore, the code example does not declare or use cursors.

Reference:

[http://msdn.microsoft.com/enus/library/ms714570\(VS.85\).aspx](http://msdn.microsoft.com/enus/library/ms714570(VS.85).aspx)



Flat Transaction

```
exec sql CREATE Table accounts (
```

```
    AcId           NUMERIC(9),
```

```
    BranchId      NUMERIC(9), FOREIGN KEY REFERENCES branches,
```

```
    AccBalance    NUMERIC(10),
```

```
    PRIMARY KEY(AcId));
```

```
*****
```

Everything inside BEGIN WORK and COMMIT WORK is at the same level; that is, the transaction will either survive together with everything else (commit), or it will be rolled back with everything else (abort)

everything survive together or rolled back ~~together~~ together.



Flat Transaction ...

```
exec sql BEGIN DECLARE SECTION;
    long AcclId, BranchId, TellerId, delta, AccBalance;
exec sql END DECLARATION;

/* Debit/Credit Transaction*/
DCApplication()
{read input msg;
exec sql BEGIN WORK;
AccBalance = DodebitCredit(BranchId, TellerId, AcclId, delta);
send output msg;
exec sql COMMIT WORK;
}
```



```
/* Withdraw money -- bank debits; Deposit money – bank credits */
Long DoDebitCredit(long BranchId,
                    long TellerId, long AcclId, long AccBalance, long delta){
    exec sql UPDATE accounts
        SET AccBalance =AccBalance + :delta
        WHERE AcclId = :AcclId;
    exec sql SELECT AccBalance INTO :AccBalance
        FROM accounts WHERE AcclId = :AcclId;
    exec sql UPDATE tellers
        SET TellerBalance = TellerBalance + :delta
        WHERE TellerId = :TellerId;
    exec sql UPDATE branches
        SET BranchBalance = BranchBalance + :delta
        WHERE BranchId = :BranchId;
    Exec sql INSERT INTO history(TellerId, BranchId, AcclId, delta, time)
        VALUES( :TellerId, :BranchId, :AcclId, :delta, CURRENT);
    return(AccBalance);
}
```



Let's include a check on the account balance and refusing any debit that overdraws the account.

...

```
DCAapplication(){
```

```
    read input msg;
```

```
    exec sql BEGIN WORK;
```

```
    AccBalance = DodebitCredit(BranchId, TellerId, AcclId,  
    delta);
```

```
    if (AccBalance < 0 && delta < 0){
```

账户余额不足

```
        exec sql ROLLBACK WORK; roll back.
```

```
}
```

```
else{
```

```
    send output msg;
```

账户平衡

```
    exec sql COMMIT WORK; un
```

```
}
```

提交成功

```
}
```



Limitations of Flat Transactions

Flat transactions do not model many real applications.

E.g. airline booking

BEGIN WORK

S1: book flight from Melbourne to Singapore *roll back to here*

We save point S2: book flight from Singapore to London *then start from there*

S3: book flight from London to Dublin

END WORK

destination changed \Rightarrow *roll back to the start X*

Problem: from Dublin if we cannot reach our final destination instead we wish to fly to Paris from Singapore and then reach our final destination.

If we roll back we need to re do the booking from Melbourne to Singapore *which is a waste*

we are unable to reuse previous operations if we have to roll back to the origin point. ~~waste~~ (unnecessary) completion



Limitations of Flat Transactions ...

```
IncreaseSalary()
```

```
{      real percentRaise;  
      receive(percentRaise);  
      exec SQL BEGIN WORK;  
          exec SQL UPDATE employee  
              set salary = salary*(1+ :percentRaise)  
              send(done);  
      exec sql COMMIT WORK;  
      return  
}
```

This can be a very long running transaction. Any failure of transaction requires lot of unnecessary computation.

→ To avoid redoing transactions from the beginning \Rightarrow instead I save some flags for my finish Action's \Rightarrow roll back to some previous saved actions

Flat transaction with save points

BEGIN WORK

SAVE WORK 1 give tag to

Action 1

Action 2

SAVE WORK 2

Action 3

Action 4

Action 5

SAVE WORK3

Action 6

Action 7

ROLLBACK WORK(2)

roll back will take you back to where you're declared before.

Although some actions will lost, will still save many actions that have been executed before

Action 8

Action 9

SAVE WORK4

Action 10

Action 11

SAVE WORK 5

Action 12

Action 13

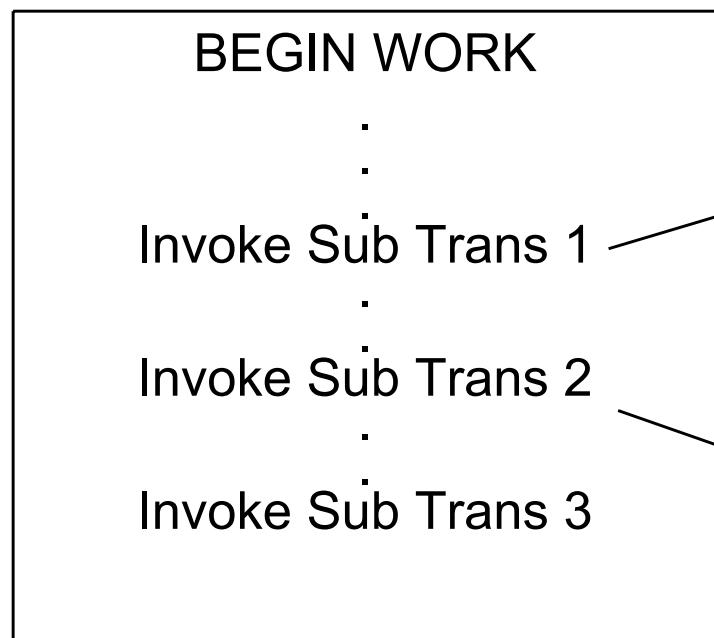
ROLLBACKWORK(5)

Want save point
要達成，得有個用意

Reading - Chained transactions



Nested Transactions



10 10
20 20
25 25
30 30



Nested Transaction Rules

Commit rule

- A subtransaction can either commit or abort, however, **(commit cannot take place unless the parent itself commits.)**
- Subtransactions have A, C, and I properties but not D property unless all its ancestors commit.
- Commit of a sub transaction makes its results available only to its parents.

Roll back Rules

If a subtransaction rolls back, all its children are forced to roll back.

Visibility Rules

Changes made by a subtransaction are visible to the parent only when the subtransaction commits. All objects of parent are visible to its children.

Implication of this is that the **parent should not modify objects while children are accessing them.** This is not a problem as parent does not run in parallel with its children.

TP monitor 的作用
~~如果~~: | 有买了最后一件商品, 又继续一个用 TP monitor 的操作 \Rightarrow 需要让那操作回滚
T₁ T₂ \Rightarrow ~~回滚~~ cart

Transaction Processing Monitor

The main function of a TP monitor is to *integrate* other system that ~~database change~~ is consistent.

- TP monitors manage the transfer of data between clients and servers
- breaks down applications or code into transactions and ensures that all databases are updated properly
- It also takes appropriate actions if any error occurs



TP monitor services

ACID Heterogeneity: If the application needs access to different DB systems, local ACID properties of individual DB systems is not sufficient. Local TP monitor needs to interact with other TP monitors to ensure the overall ACID property. A form of 2 phase commit protocol must be employed for this purpose (will be discussed later).

ACID Control communication: If the application communicates with other remote processes, the local TP monitor should maintain the communication status among the processes to be able to recover from a crash.
(Distributed + 2 phase 也有, 也有)



TP Services ...

0 **Terminal management:** Since many terminals run ACID client software, the TP monitor should provide appropriate ACID property between the client and the server processes.

Presentation service: this is similar to terminal management in the sense it has to deal with different presentation (user interface) software -- e.g. X-windows

Context management: E.g. maintaining the sessions etc.

Start/Restart: There is no difference between start and restart in TP based system.

Chain of request: make a deposit to an account → transfer some amount to another account

faster
helps data retrieval

Session ~~1.1.2014~~
e.g. In session → fix student records → if make a request for changing just
one particular student's record ⇒ so the ~~entire~~ ^{top send} or just a single row, ^{retriever} & update ~~DB~~ easily
efficient make ^{or} without the need of refreshing whole table.

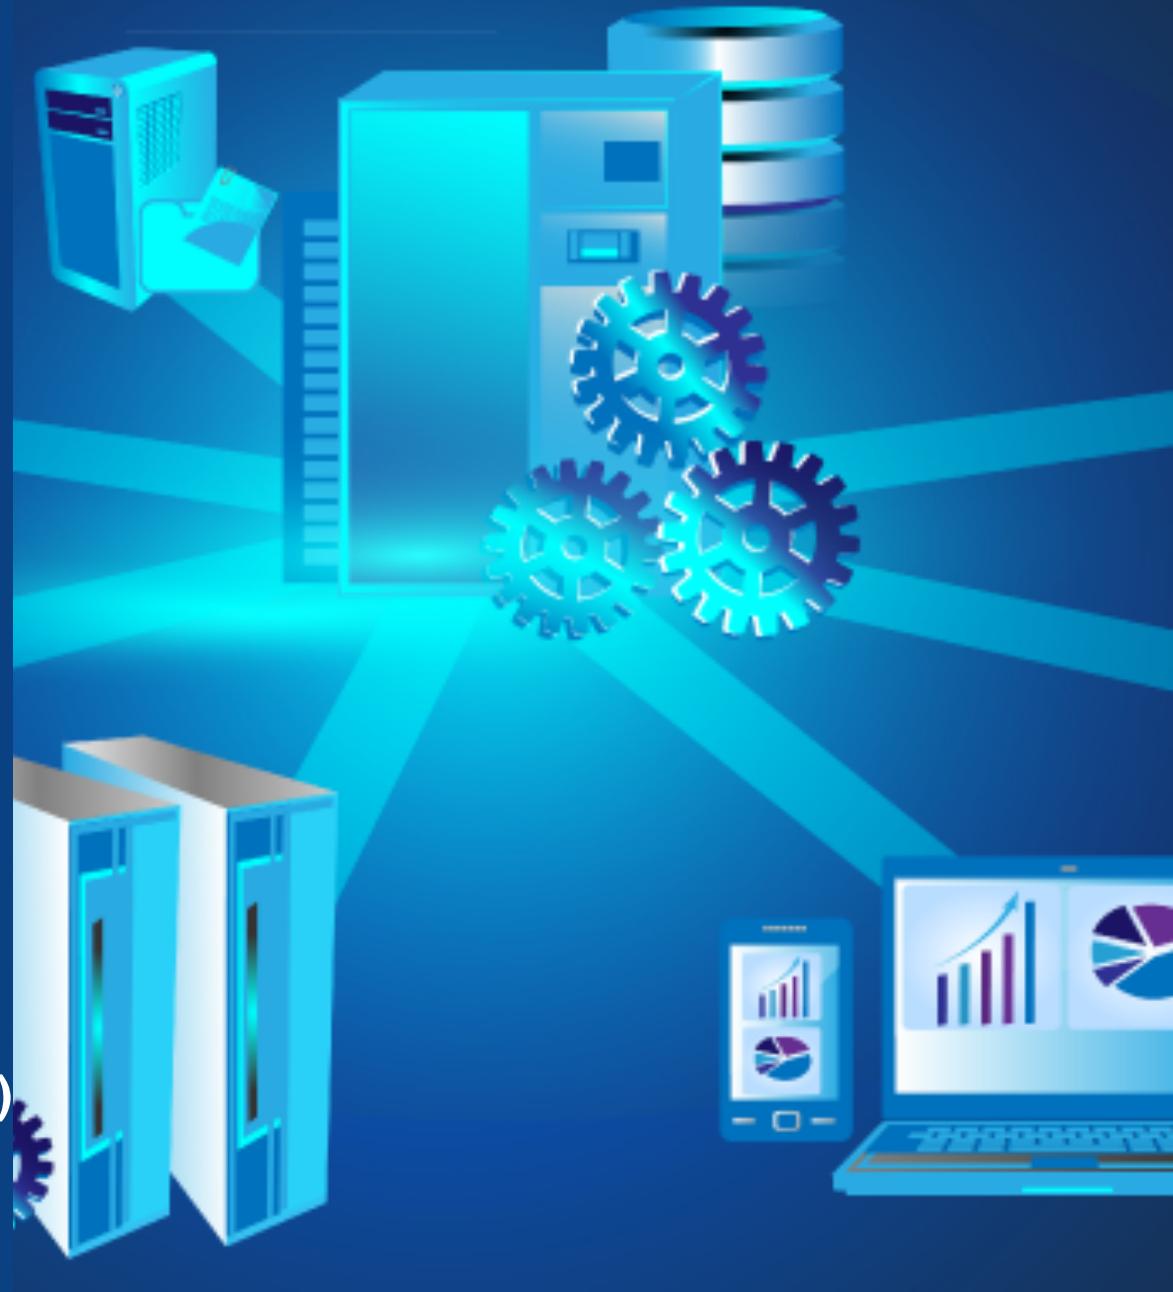


COMP90050 Advanced Database Systems

Semester 2, 2024

Lecturer: Farhana Choudhury (PhD)

Live lecture – Week 4





Continuation from last week's live lecture



Effect of indexes – do they always improve performance?

Group activity - <https://tinyurl.com/49ASF2RC>

| ID | Name | Department |
|----|------|------------|
| 1 | Jane | Comp. Sci |
| 2 | John | Biology |

| ID | Position | Salary |
|----|--------------------|--------|
| 1 | Lecturer | 75,000 |
| 1 | Research assistant | 40,000 |
| 2 | Senior lecturer | 82,000 |

There are indexes constructed on:

- B+tree on IDs of both tables
- Hash index on Department
- Bitmap index on position

What will be the performance of data insertion, deletion, and updates, compared to no indexing (or just on IDs)?

① B+ tree is preferred choice when performing range queries, (e.g., finding a list of users' names within a specific range.)

② Hash indexes are highly effective for equality-based queries, where you are looking for an exact match of a key.

① B+tree on IDs of both tables ② Hash index on department ③ bitmap index on position

Q query 1: Inserting records for a new lecturer (不用 V.S 不用)
Inserting records means that we need to update those indexes for ID: 用 ~~B+~~ tree 的信息

容易 straightforward (empty space for insert), 但有时需要 restructure/reorganize the index.
BFR.

不止① (只用 index), index 也一样, 也会被更新
index → updated as well

∴ all four indexes need to be updated ⇒ because those corresponding info needs to be inserted ⇒ it is overhead extra.

But without using index ⇒ we just insert directly into the table now no extra index overhead.
(No index faster)

Q query 2: updating the salary of instructor with ID 2 用 V.S 不用:
⇒ we just need to go to that 'id' position in the table ⇒ we don't have to go to other table ⇒ just going to the ~~table~~ for a particular ID ⇒ simply use B+tree on the ID much much faster than non-index. if we don't have B+tree we have to go through those all records even if the IDs were sorted ⇒ we still have to do binary search (more time-consuming)

Q query 3: Find the total number of research assistants 用 V.S 不用
用 C there is just 1 and 0 to match with corresponding column's value ⇒ finding the RA is just equivalent to count the number of ones on a particular position on bitmap ⇒ very fast..

Q query 4: Finding the name of all instructors who are in 'Biology' department. 用 V.S 不用
Hash index will just go to bucket ⇒ Then we can quickly find the instructor in Biology department (in the same bucket) we don't have to ~~find~~ other instructors who work in Computer Science or ECE. (faster to focus on exactly those records that meet this criteria. (quickly filter based on the department name))

~~Query 5: Find person for which the salary is greater than 30000. Explain.~~

~~using any join at least one right table , not ID.~~

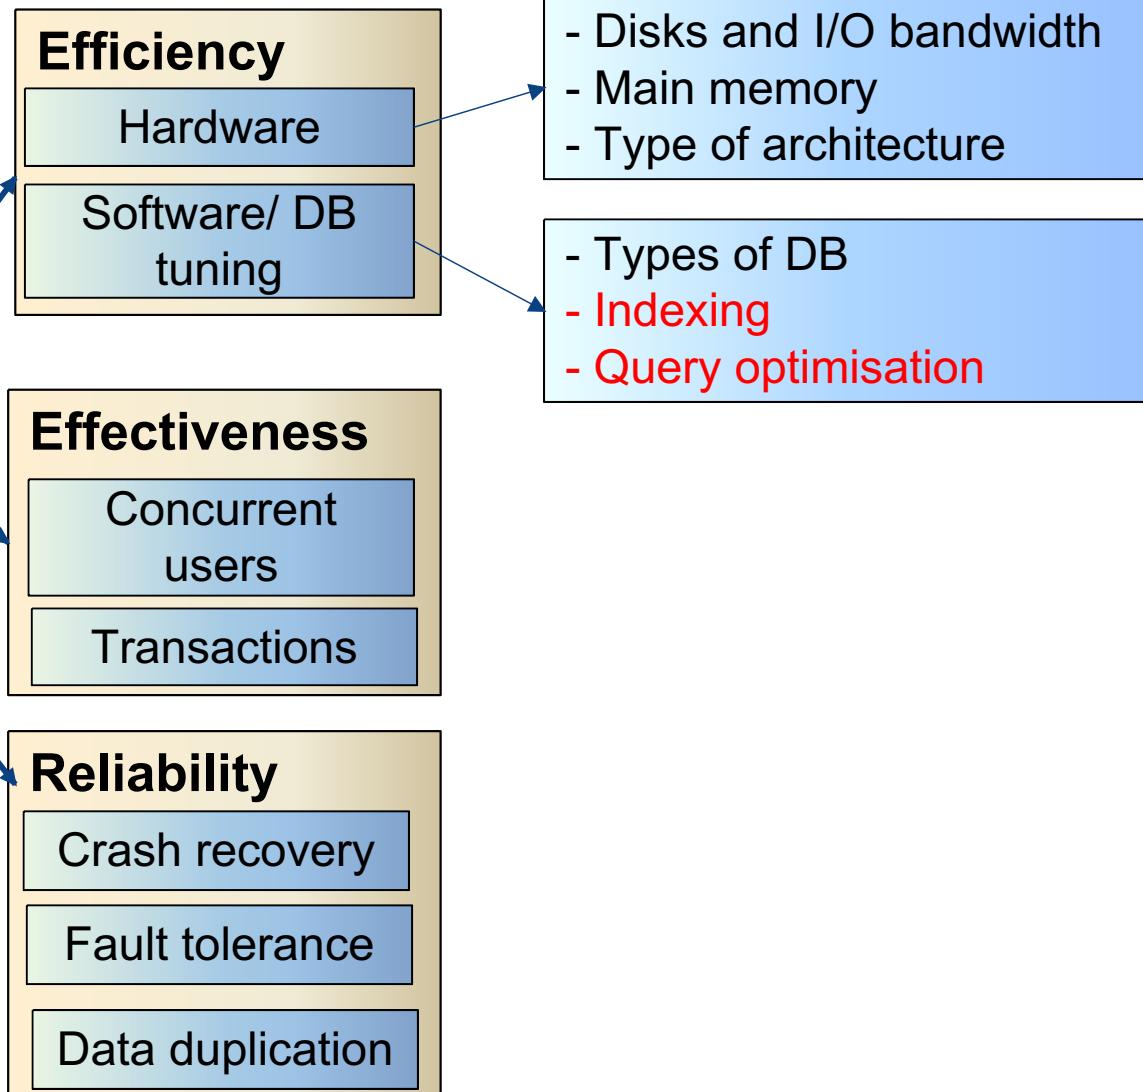
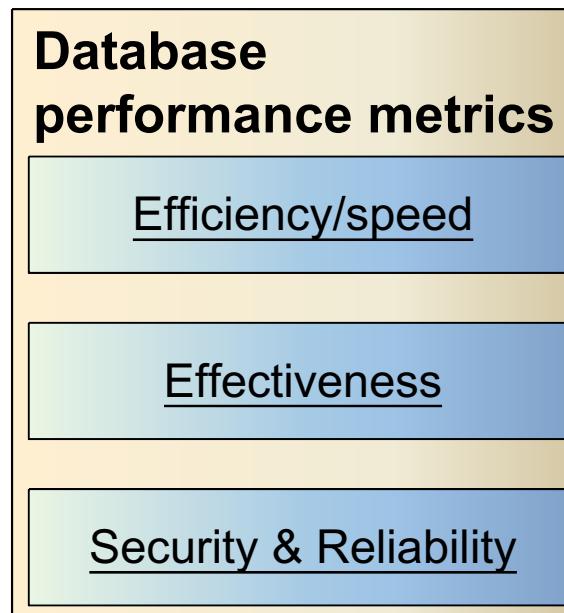
Main Cost is search through all the records in this table to make comparison one by one \Rightarrow hardly all those four index ~~is~~ not helpful for comparison.
such huge volume of Put on cheat sheet



What you need to do now?

- n Given a data set, when uploading to the DBMS
 - | Find the potential query types
 - | Research what indices that particular DBMS would have for that data type
 - | Research for what queries you would better do on what index
 - | Create index if you have large data
 - | Monitor performance
 - | Tune or create other indices

Summary of last week





Discussion on this week's topics

- Some terminologies on transactions
- Types of transactions
 - Flat transactions
 - Flat transactions with save points
 - Nested transactions
- Transaction processing monitor

Embedded SQL example in C

(Open Database Connectivity)

```
int main()
{
    exec sql INCLUDE SQLCA; /*SQL Communication Area*/
    exec sql BEGIN DECLARE SECTION;
        /* The following variables are used for communicating
           between SQL and C */

        int OrderID; /* Employee ID (from user) */
        int CustID; /* Retrieved customer ID */
        char SalesPerson[10] /* Retrieved salesperson name */
        char Status[6] /* Retrieved order status */

    exec sql END DECLARE SECTION;

/* Set up error processing */
    exec sql WHENEVER SQLERROR GOTO query_error;
    exec sql WHENEVER NOT FOUND GOTO bad_number;
```

Proper error handling is important!



What can go wrong in this example?

(Open Database Connectivity)

```
int main()
{
    ....//code to get data from database
    Printf("%d", (get_salary(EMPID)/800000);

/* Set up error processing */
    ↓
    exec sql WHENEVER DIVIDE_BY_ZERO GOTO inf_error;
    exec sql WHENEVER NOT_PERMITTED GOTO permission_err;
    exec sql WHENEVER NOT_FOUND GOTO bad_number;
    按序check error (order 很重要)
```

Time for a poll - Pollev.com/farhanachoud585



Proper error handling is important!

$a = \text{get_salary}(\text{EMPID}) / 8000$

if $a=1$ then do $a/0$ # $a=1$ 时 只有当 salary 是 80000 时

这个永远报错 (divide-by-zero) .

我们可以用这样的 run 所有数都除以 80000

Transaction Processing Monitor

Used in large scale business

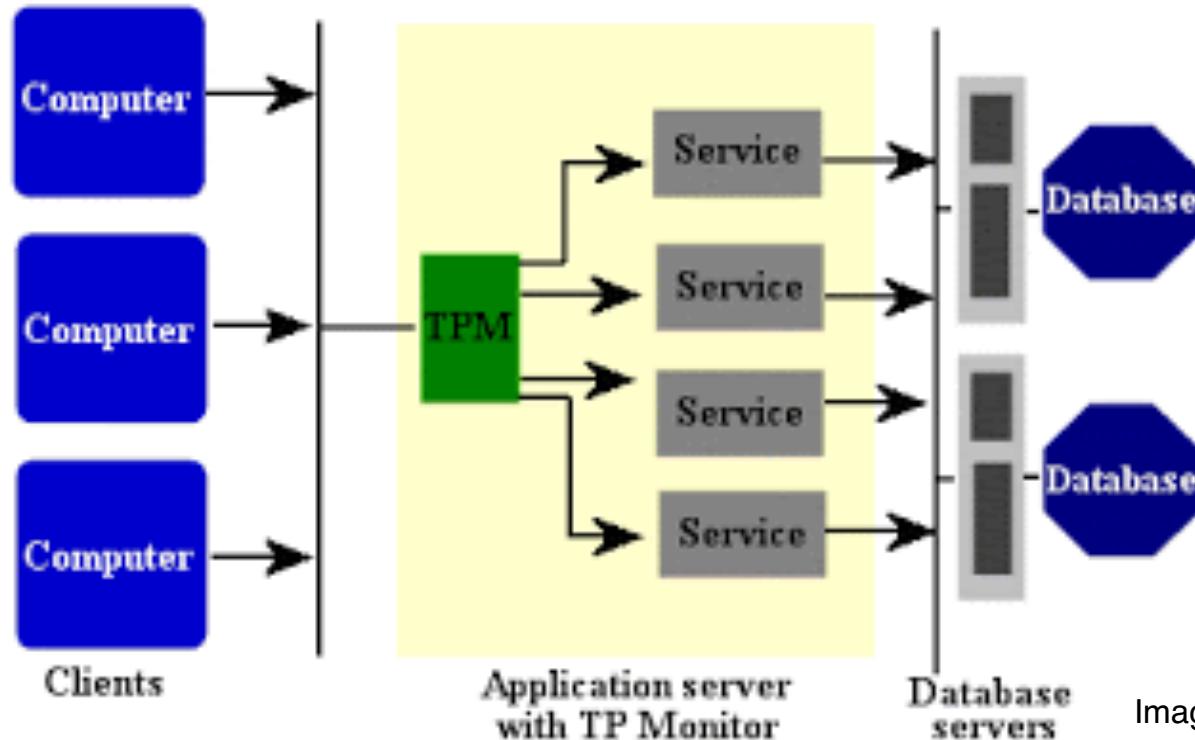


Image source: <http://3.bp.blogspot.com>

Integrates other system components and manage resources.

- Manages the transfer of data between clients and servers
- breaks down applications or code into transactions and ensures that all databases are updated properly
- It also takes appropriate actions if any error occurs



Flat Transaction

Everything inside BEGIN WORK and COMMIT WORK is at the same level; that is, the transaction will either survive together with everything else (commit), or it will be rolled back with everything else (abort)

```
exec sql BEGIN WORK;  
    AccBalance = DodebitCredit(BranchId, TellerId, AcctId, delta);  
    send output msg;  
exec sql COMMIT WORK;
```

Can be a very long running transaction
with many operations

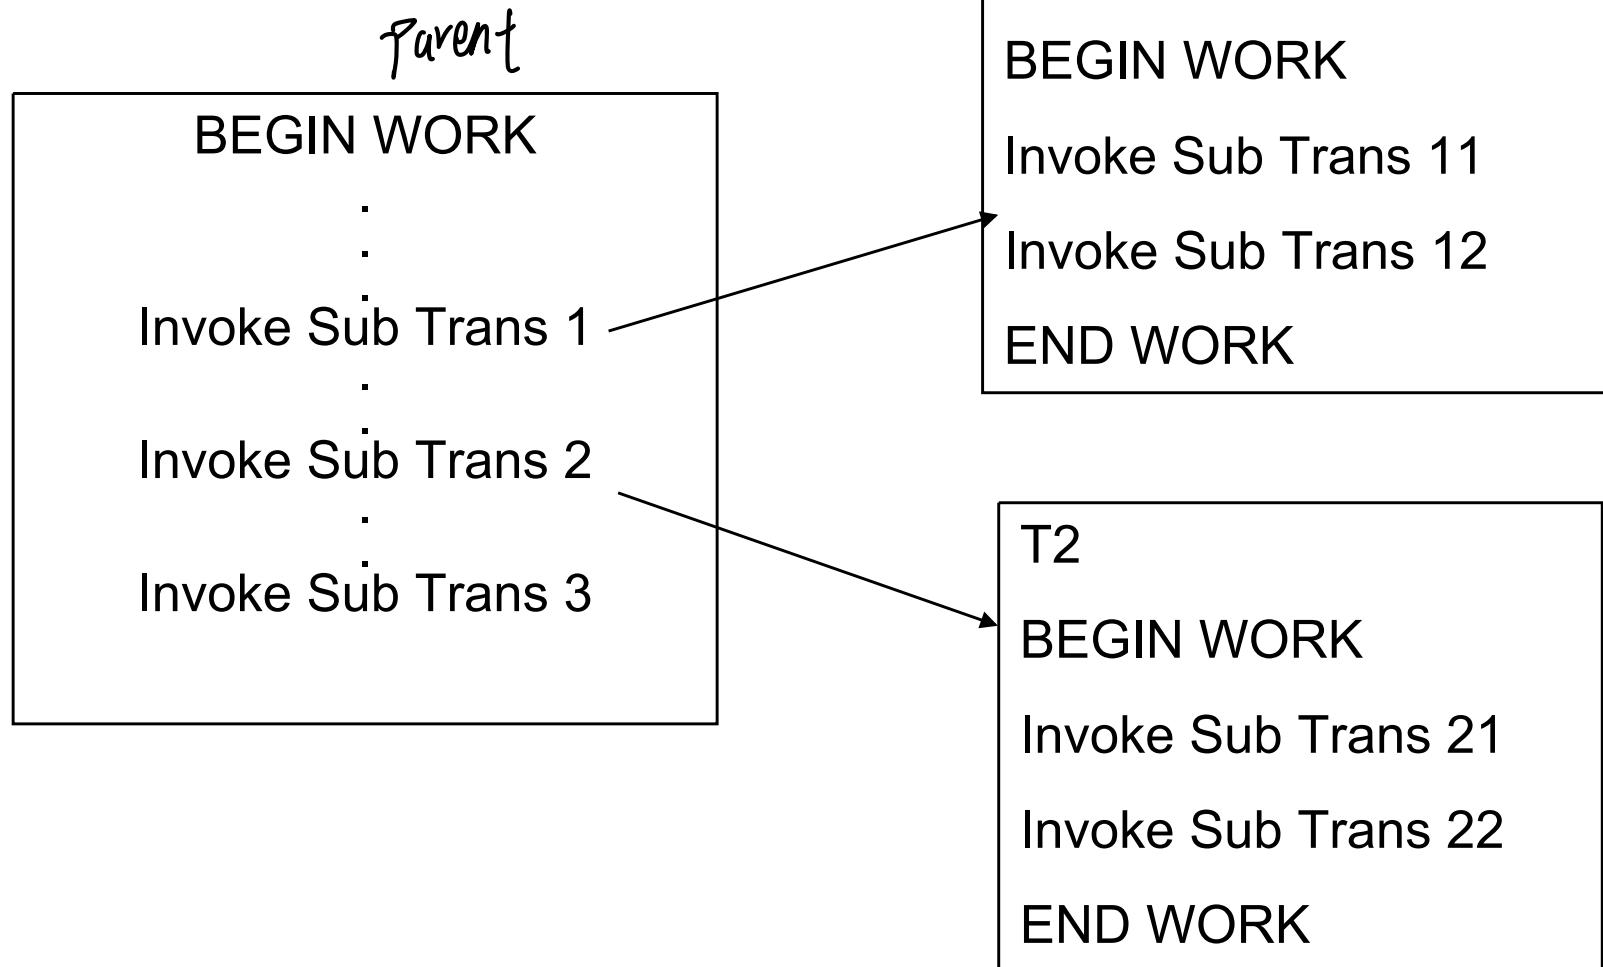
Limitations of Flat Transactions?

Time for a poll - Pollev.com/farhanachoud585

If something happen in the middle of a long funny transaction
and it rolls back , we have to start again, all the things were done will be wasted (waste
time . computation power, etc.)



Nested Transactions



Nested Transactions

Commit rule

~~(AC)~~ but ~~no durability~~ ~~for sub trans~~

- A subtransaction can either commit or abort, however, **commit cannot take place unless the parent itself commits.** 除非父节点 commit 了子节点才能 commit.
- Subtransactions have A, C, and I properties but not D property unless all its ancestors commit.
- Commit of a sub transaction makes its results available only to its parents.

Roll back Rules

If a subtransaction rolls back, all its children are forced to roll back.

Visibility Rules

Changes made by a subtransaction are visible to the parent only when the subtransaction commits. All objects of parent are visible to its children.

Implication of this is that the **parent should not modify objects while children are accessing them.** This is not a problem as parent does not run in parallel with its children.

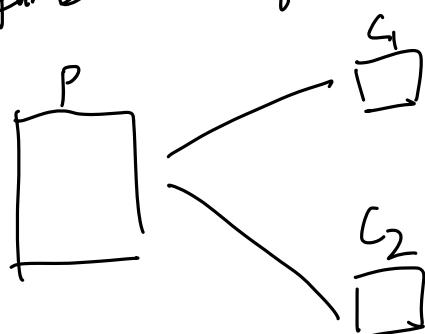
What are the advantages?

~~pros: Unlike flat transaction . every thing gets rolled back.~~

~~child~~ ~~parent~~

- ① Subtransaction can run also in parallel if there are no conflicts between them, making the whole transaction much faster (for example divide data into several task chunks)
- ② If some transaction rolled back, it won't impact others. Enhance modularity.
- ③ Sometimes we want to roll back purposefully \Rightarrow sometimes we want users to input something \Rightarrow account balance, amount of cash the person wants to withdraw from their account \Rightarrow which might make account balance negative \Rightarrow we just withdraw that part of the transaction \Rightarrow (maybe before that the user made some other changes which were perfectly fine, and won't affect those changes.)

So for some specific input or under some specific condition \Rightarrow we just want to roll back part of transaction without impacting others \Rightarrow nested transaction allows you ~~to~~ to design such functionality.



if C_1 has satisfied, we don't need to execute this part

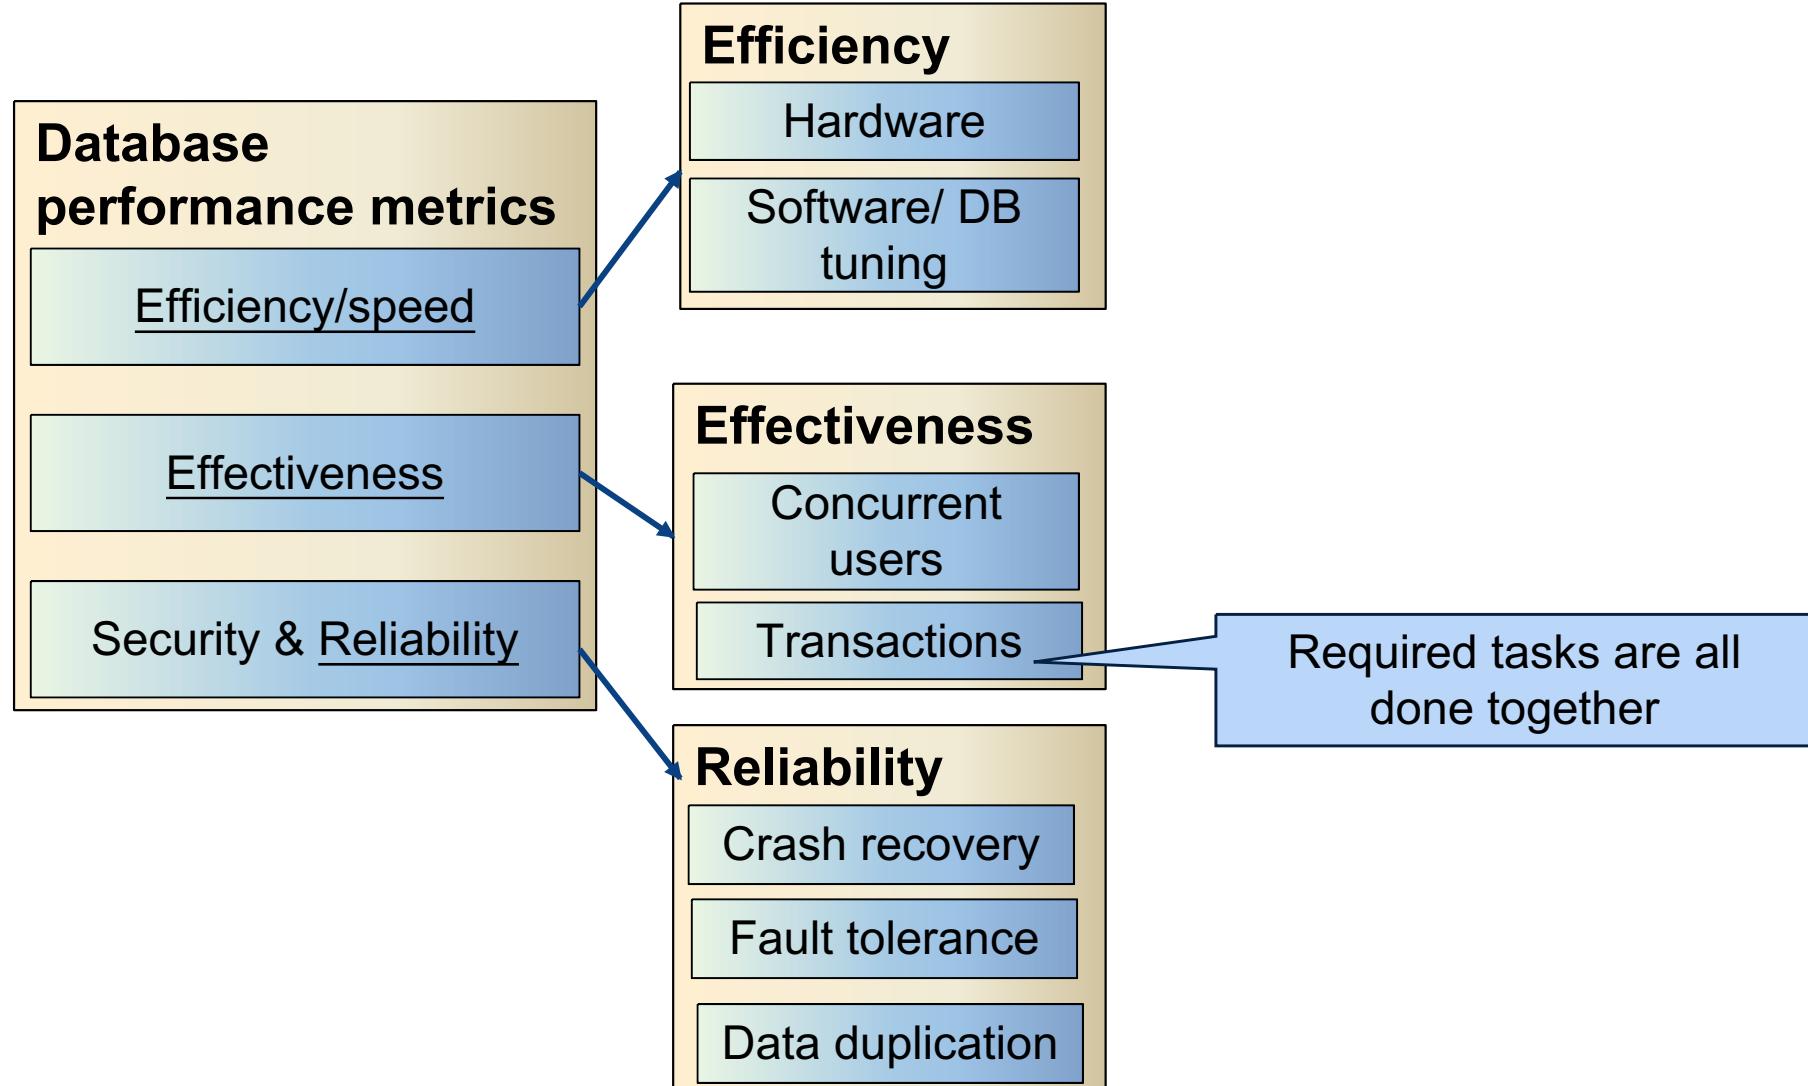
~~if~~ C_2 is wrong, we can just let C_2 roll back

such sort of functionality is possible for nested transactions

~~fine tuning~~ fine tuning roll back is possible.

roll back a specific part of transaction while keeping others unimpacted.

Core Concepts of Database management system

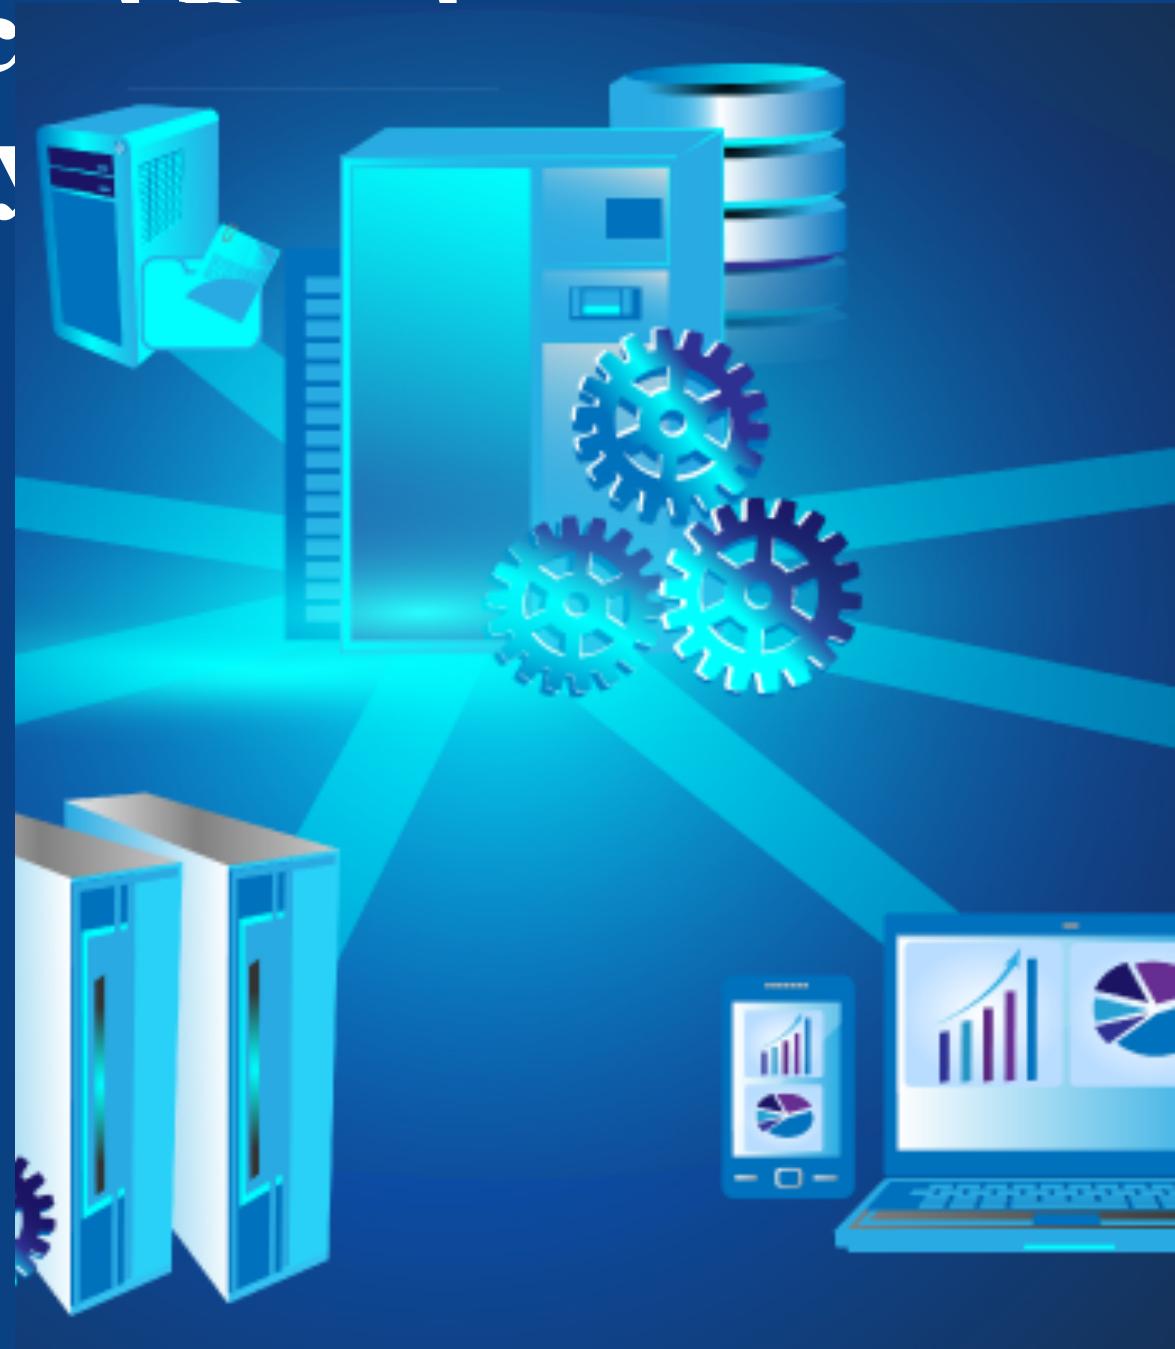




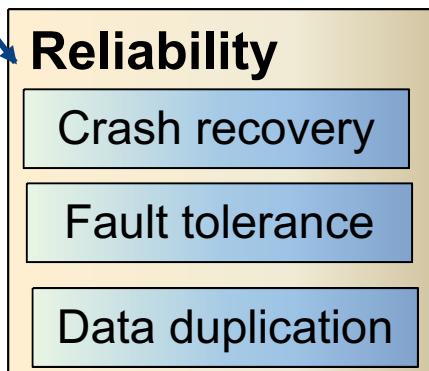
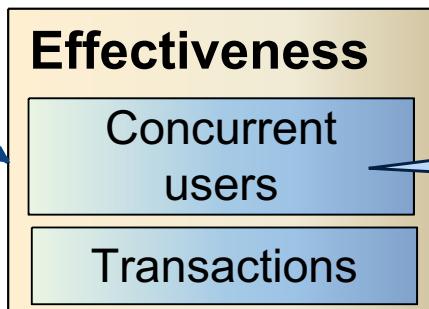
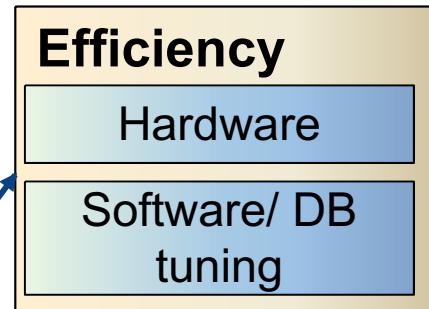
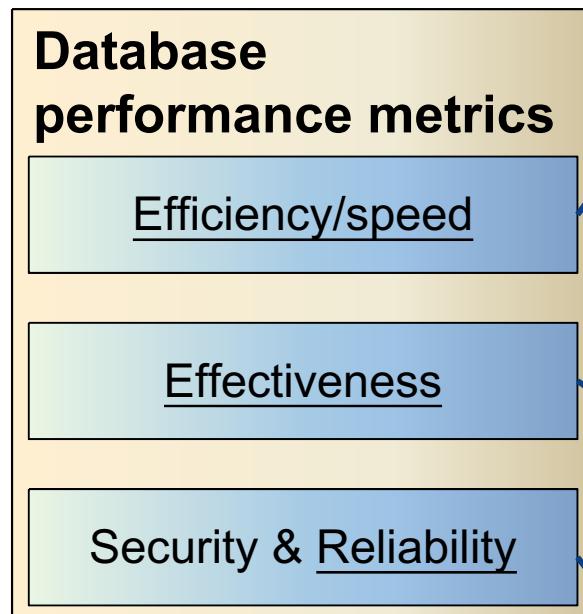
COMP90050:

Advanced Systems

Lecturer: Farhana Choudhury (PhD)
Concurrency control
Week 5



Core Concepts of Database management system



Users reading and writing over the same data

Concept of locks
Implementation of exclusive access
Semaphore
Deadlocks



Concurrency Problems

Shared counter = 100;

Task1/Trans/Process/Thread
counter = counter +10;

Task2/Trans/Process/Thread
counter = counter +30;

Task1 and Task2 are running concurrently. What are the possible values of counter after the end of Task1 and Task2?

- a) counter == 110;
- b) counter == 130;
- c) counter == 140;

For correct execution we need to impose exclusive access to the shared variable counter by Task1 and Task2.



Concurrency Problems

Shared counter = 100;

Task1/Trans/Process/Thread
counter = counter +10;

Task2/Trans/Process/Thread
counter = counter +30;

Task1 and Task2 run concurrently. What are the possible values of counter after the end of Task1 and Task2?

Note: == means equals.

a) counter == 110

Sequence of actions

T1: Reads counter == 100

T2: Reads counter == 100

T2: Writes counter == 100+30

T1: Writes counter == 100+10

b) counter == 130

Sequence of actions

T1: Reads counter == 100

T2: Reads counter == 100

T1: Writes counter == 100+10

T2: Writes counter == 100+30

c) counter == 140;

Sequence of actions

T1: Reads counter == 100

T1: Writes counter == 100+10

T2: Reads counter == 110

T2: Writes counter == 110+30

different versions

Time



Concurrency Control

- To resolve conflicts
- To preserve database consistency

Different ways for concurrency control

- **Dekker's algorithm (using code)** - needs almost no hardware support, but the code is very complicated to implement for more than two transactions/processes
- **OS supported primitives (through interruption call)** - expensive, independent of number of processes, machine independent
- **Spin locks (using atomic lock/unlock instructions)** – most commonly used



Concurrency control: Implementation of exclusive access

Dekker's algorithm

```
int c1, c2, turn = 1; /* global variable*/
```

```
T1
{ some code T1}

/* T1 wants exclusive access to the resource
   and we assume initially c1 == 0 */

c1 = 1; turn = 2;
repeat until { c2 == 0 or turn == 1}
/* Start of exclusive access to the
   shared resource (successfully
   changed variables) */
use the resource
counter = counter+1;
/* release the resource */
c1 = 0;
{some other code of T1}
```

```
T2
{ some code T2}

/* T2 wants exclusive access to the resource
   and we assume initially c2 == 0 */

c2 = 1; turn = 1;
repeat until { c1 == 0 or turn == 2}
/* Start of exclusive access to the
   shared resource */
use the resource
counter = counter+1;
/* release the resource */
c2 = 0 ;
{some other code of T2}
```



Implementation of exclusive access

- Dekker's algorithm
 - needs almost no hardware support although it needs atomic reads and writes to main memory, That is exclusive access of one time cycle of memory access time!
 - **the code is very complicated to implement if more than two transactions/process are involved**
 - harder to understand the algorithm for more than two process
 - takes lot of storage space
 - uses busy waiting waiting for the resources to be released by the other transaction
 - efficient if the lock contention (that is frequency of access to the locks) is low

↳ low
frequency of access to the locks is

*if the resource has been by other transactor \rightarrow T₁ cannot do anything but just to wait \rightarrow keep repeating this loop and just wait
 \rightarrow busy waiting \rightarrow one transaction is just waiting in a loop for the resource to be released \Rightarrow it cannot do anything at that time hampers the overall performance of the transaction.*



Implementation of exclusive access

- OS supported primitives such as lock and unlock
 - through an interrupt call, the lock request is passed to the OS
 - need no special hardware
 - are very expensive (several hundreds to thousands of instructions need to be executed to save context of the requesting process.)
→ register status (CPU)
 - do not use busy waiting and therefore more effective
- All modern processors do support some form of spin locks.



lock: a request to access to have exclusive access to a shared resource.
after using that resource the process or transaction is released
⇒ e.g. shared database table or rows

Implementation of exclusive access

Spin Locks

Executed using atomic machine instructions such as test and set or swap

- need hardware support – should be able to lock bus (communication channel between CPU and memory + any other devices) for two memory cycles (one for reading and one for writing). During this time no other devices' access is allowed to this memory location.
- use busy waiting (~~if shared resources is currently being used by another transaction, it has to wait in the loop and wait until that resource is released by the other process~~)
- algorithm does not depend on number of processes
- are very efficient for low lock contentions – all DB systems use them

Spin locks: Once get exclusive lock on the resource, only that process of the transaction will be able to read and write on that resource ⇒ no other device or transaction will be able to access that particular memory location that is being currently read or written.



Implementation of Atomic operations: test and set

testAndSet(int *lock)

```
{ /* the following is executed atomically, memory bus can be locked for up  
    to two cycles (one for read and for writing*)
```

```
if (*lock == 1){ * lock = 0; return (true)}  
else return (false);
```

if not, acquire the resource for other processes

exclusive access
can not access

→ check current status; [是] ⇒ 状态 ⇒ 读

Using test and set in spin lock for exclusive access

```
int lock = 1; % initial value
```

T1

/*acquire lock*/

```
while (!testAndSet( &lock )
```

/*Xlock granted*/

//exclusive access for T1;

counter = counter+1; ↗

/* release lock */

lock = 1;

Why this is important? 有 process (標準) → 好了 (no
需要 hardware support) T2

T2

/*acquire lock*/

while (!testAndSet(&lock)) { if it's not

/*Xlock granted*/

//exclusive access for T2;

counter = counter+1;

/ release lock */*

lock = 1;



Implementation of Atomic operation: compare and swap

The following code may have lost values

```
temp = counter + 1; //unsafe to increment a shared counter
```

```
counter = temp; //this assignment may suffer a lost update
```

第一行执行同时有其他线程修改了 counter

的值

即 Counter 值无效了 (Not valid).

CPU 第一：先读到的是什么

为了防止这种现象保证 Atomicity.

⇒ compare and swap 用上。

(也是用 exclusive access)

a) counter == 110

Sequence of actions

T1: Reads counter == 100

T2: Reads counter == 100

T2: Writes counter == 100+30

T1: Writes counter == 100+10

Using compare and swap in
spin lock for exclusive access



Implementation of Atomic operation: compare and swap

The following code may have lost values

```
temp = counter + 1; //unsafe to increment a shared counter
```

```
counter = temp; //this assignment may suffer a lost update
```

Instead, we can use the ~~atomic~~ operation of compare and swap instruction

```
boolean cs(counterint *cell, int *old, int *new)  
/* the following is executed atomically */  
if (*cell == *old) { *cell = *new; return TRUE; }  
else { /* old = *cell; return FALSE; */ }
```

~~如果修改了共享变量 old 为 new [0] 改成 [0]~~
~~→ 检查是否被其他线程修改了 old [0] → 取新的值~~
→ So that it can be used later to make further change.

~~old 为 120 → 说明值已被另外修改。~~

temp = counter;

do

new = temp+1;

~~若返回 false 则进入忙等待~~
~~返回 true 则一直忙等待~~

Using compare and swap in

while(!cs(&counter,&temp,&new)); spin lock for exclusive access

inside loop: keep checking if this value is safe ⇒ means it has not been changed by others



Why we prefer to use compare and swap, rather than test and set?
pros for compare and swap: ask for exclusive lock only when is required.

Semaphores

Semaphores derive from the corresponding mechanism used for trains: a train may proceed through a section of track only if the semaphore is clear. Once the train passes, the semaphore is set until the train exits that section of track.

② For ready and write if use test and set \Rightarrow which means one transaction has to start and completely finish its work on that shared variable \Rightarrow then only the second

Try to Get(track), wait if track not clear

transaction will be able to start.

③ However, if both transactions are just ready that value-ready does not conflict what happening for other transaction if both of them just ready and not making any change to the shared variable, ~~no~~ compare and

swap will allow multiple transactions to run ~~concurrently~~ for that ready part.



If Get(track) successful, use it (no other train will be able to use it now)

So it will improve the performance ~~as~~ \Rightarrow if multiple transactions are able to run concurrently. \Rightarrow the overall number of operations that are being executed in the system $\xleftarrow{\text{improve}}$ run more operations together

Compare and Swap

for write part

just check if it is changed by someone \Rightarrow whether it's safe to change now, just do this check and not prevent other transaction from reading

\Rightarrow then this compare and swap is ~~more~~ more suitable option.

↳ this requires transaction to fetch \Rightarrow ~~lock~~ still needed?

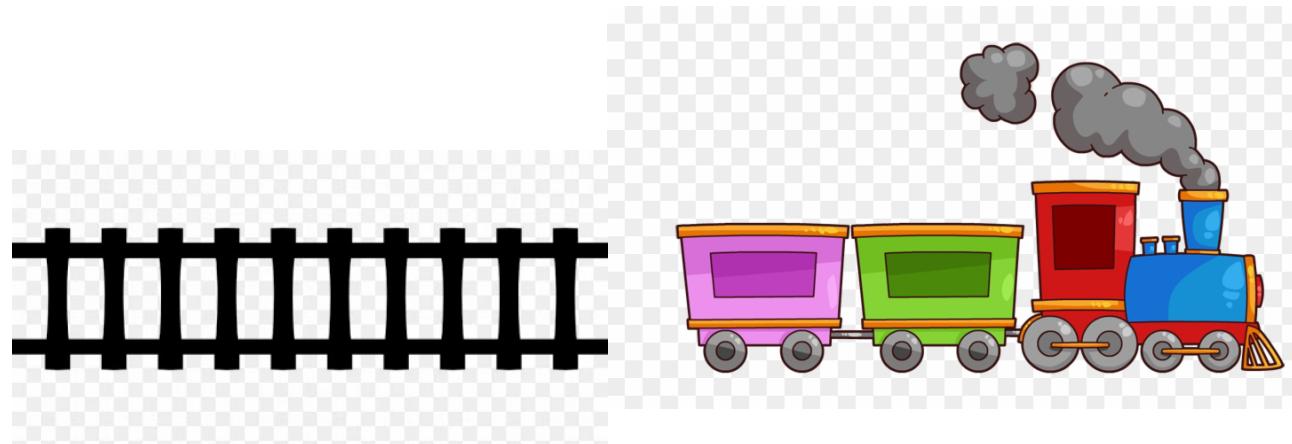


Semaphores

Semaphores derive from the corresponding mechanism used for trains: a train may proceed through a section of track only if the semaphore is clear. Once the train passes, the semaphore is set until the train exits that section of track.

release after using

Release(track) after using
(so that others can use it)





Semaphores

Computer semaphores have a `get()` routine that acquires the semaphore (perhaps waiting until it is free) and a dual `give()` routine that returns the semaphore to the free state, perhaps signalling (waking up) a waiting process.

Semaphores are very simple locks; indeed, they are used to implement general-purpose locks.

若资源未用
需要等待

当操作正在执行的 transaction, track 可用,



Implementation of Exclusive mode Semaphore

Pointer to a queue of processes

If the semaphore is busy but there are no waiters, the pointer is the address of the process that owns the semaphore.

If some processes are waiting, the semaphore points to a linked list of waiting processes. The process owning the semaphore is at the end of this list.

After usage, the owner process wakes up the oldest process in the queue (first in, first out scheduler)

- the process owns the shared resource
pointer is a mechanism maintaining the queue of processes in the waiting.

这个队列在尾部



Implementation of Exclusive mode Semaphore

type long PID

type struct Process{

 PID pid; /*process ID*/

 PCB * sem_wait; /* waiting process are put in the queue */

} PCB;

PID MyPID (void); /* returns the caller's process ID */

PCB * MyPCB(void)

/* returns pointer to caller's process descriptor */

void wait (void); /* suspends calling process */

void wakeup(PCB * him) wakes him.



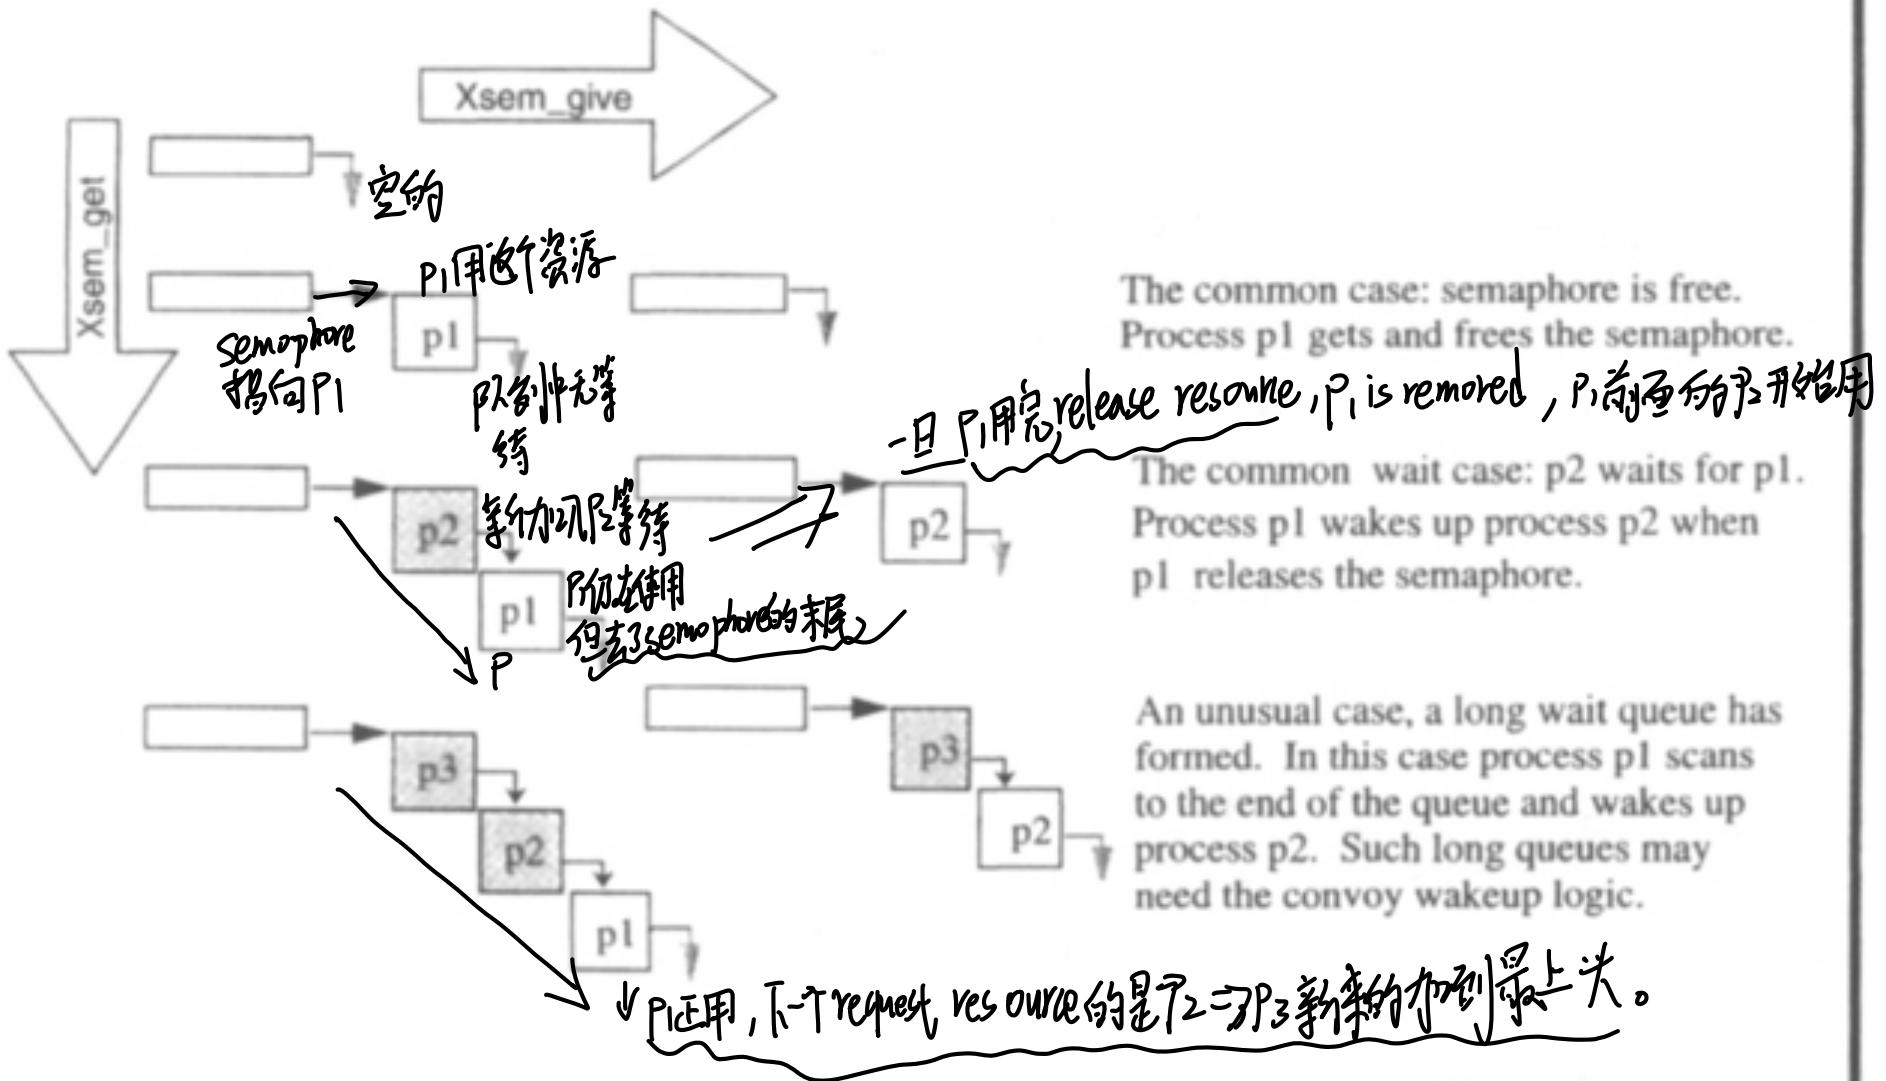
Implementation of exclusive lock semaphore operations

```
void lock(Xsemaphore *sem){  
    PCB * new = myPCB();  
    PCB * old  = NULL;  
    /*put the process in the semWait queue*/  
    do{  
        new ->semWait = old;  
    } while(! cs(sem, &old, &new));  
    /*if queue not null then must wait for a  
     * wakeup from the semaphore owner*/  
    if (old != NULL) wait() % OS call  
    return;  
}
```

```
type struct Process{  
    PID pid;  
    PCB * sem_wait;  
}PCB;  
typedef PCB *Xsemaphore  
Void initialise( Xsemaphore *sem)  
{*sem = NULL; return;}  
  
void unlock(Xsemaphore *sem){  
    PCB * new = NULL;  
    PCB * old  = myPCB();  
    if(cs(sem, &old, &new)) return;  
        while(old->semWait  
        !=myPCB()){  
            old = old->semWait  
        }  
        old->semWait = NULL;  
        /*wake up the oldest process (first in,  
         * first out scheduler)*/  
        wakeup(old); % OS Call  
}
```



每次只能有一个访问，其他的放进队列等待，一个用完，下一个用。



Convoy avoiding semaphore

The previous implementation may result a long list of waiting processes – called convoy

To avoid convoys, a process may simply free the semaphore (set the queue to null) and then wake up every process in the list after usage.

In that case, each of those processes will have to re-execute the routine for acquiring semaphore.

problems: long list of waiting processes \Rightarrow convoy

if one of them or couple of them are long-running processes,
then the one that make request later has to wait for longer time
in the queue.



Solution: once one transaction is done \Rightarrow lock released \Rightarrow every process in the waiting list is
woken up \Rightarrow recall the resources again \Rightarrow order of the waiting in next queue will not



they all

Convoy avoiding semaphore

```
void lock(Xsemaphore *sem){  
    PCB * new = myPCB();  
    PCB * old = NULL;  
    do{  
        new ->semWait = old;  
    }while( ! cs(sem, &old, &new));  
    if(old != NULL) {  
        wait(); lock(sem);  
    }  
    return;  
}
```

void unlock(Xsemaphore *sem){

PCB * new = NULL;
 PCB * old = myPCB();
 while(!cs(sem, &old, &new));

**while(old->semWait !=
myPCB()){**

new = old->semWait;

**old->sem_wait =
NULL;**

wakeup all the
process in queue
wakeup(old);

Once it is weaken up \Rightarrow it can again ask for this process
which will go back from the beginning of this

return;
lock function,

be the same
as the queue before.
(reshuffle \Rightarrow if they all wake up and
all require the resource again)
long waiting queue does not occur
(ensure that we have another chance to ask
for resources)



ensure long waiting queue do not occur and
this LDDP is responsible for that.

A quick summary

Concurrency problems – why we need concurrency control

Implementation of exclusive access –

- Dekker's algo
- OS primitives
- **Spin locks** - Atomic operations to get and release locks

{ Different ways of how
the exclusive access
can be granted.

→ if other processors that are also requesting for that lock the semaphore

Semaphores - get lock, release lock, maintain queue of processes will maintain those

Avoiding long queues in semaphores

↳ hijacking → reshuffling processors.

~~Concurrency problems: When multiple processors or transactions are running concurrently~~

~~and they are working on the same shared resources or the same shared data.~~

問題在那: The order of their execution can effect the final output of that shared resource
and it will have consistency on the value of those shared resources.

which is not desirable - to have consistency in the form of those shared resources
the processors or transactions need some exclusive access on those resources

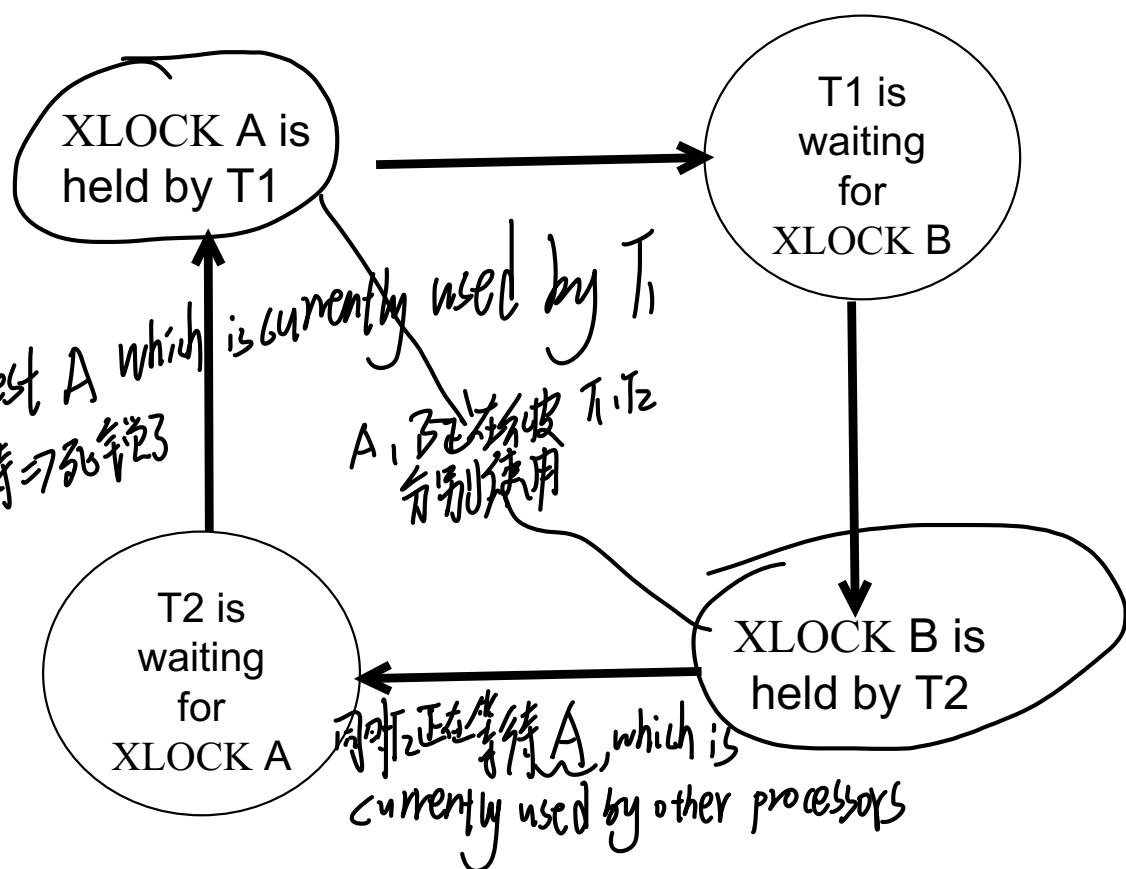
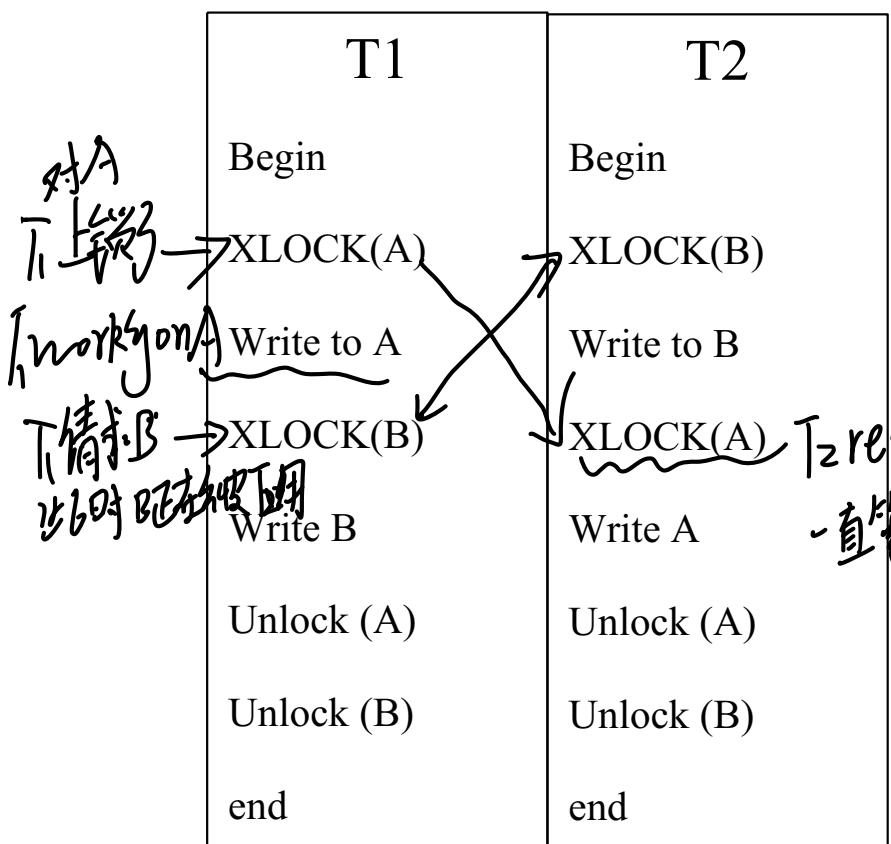


Deadlocks

In a deadlock, each process in the deadlock is waiting for another member to release the resources it wants.

A very simple Deadlock

Resource Dependency Graph





这个loop: 每个transaction 都在等待其他transaction释放资源，同时自己holding the lock of the resources which ~~is~~ are currently requested by other transactions

Deadlocks

So that no process can proceed to the next set of operations \Rightarrow the operation

Solutions: will not be executed \Rightarrow because ~~all~~ each of processes are just waiting

- Have enough resources so that no waiting occurs – not practical
- Do not allow a process to wait, simply rollback after a certain time. This can create live locks which are worse than deadlocks.
- Linearly order the resources and request of resources should follow this order, i.e., a transaction after requesting i^{th} resource can request j^{th} resource if $j > i$. This type of allocation guarantees no cyclic dependencies among the transactions.

- ① \rightarrow 7. 不知道哪两个 transaction 同时请求了同一个资源，造成死锁。
- ② $i_1 \rightarrow i_2 \rightarrow \dots \rightarrow i_n$: 一个线性顺序。如果一个process \Rightarrow 到达deadlock \Rightarrow rollback
- ③ N 个资源 \Rightarrow put linear order on these resources \Rightarrow if a transaction requests for a particular resource i \Rightarrow then it can only request another resource that is in a later position in that linear order resource $[j > i]$



graph-structure

2 3
5 4

Deadlocks

graph
节点
节点
节点
节点

Pa: Holds resources at level i and request resource at level j which are held by Pb. $j > i$

P_a hold i request j

edges 表示要请求的资源.

Pq: Holds resources at level g and request resource at level l which are held by Pd. $l > g$

P_q hold g also request l

Pb: Holds resources at level j and request P_b hold j resource at level k which are held by Pc. $k > j$

request k

Pc: Holds resources at level k and request resource at level l which are held by Pd $l > k$

P_c hold k request l

Pd: Holds resources at level l and is currently running.
P_d hold l

$l > k > j > i \text{ and } l > g$.

We cannot have loops. The dependency graph can be a tree or a linear chain and hence cannot have cycles.

X 表示请求
P_d 持有 g
因为 $l > g$, 优先级高



这样可避免 deadlock, 但 expensive \Rightarrow 若一个资源被释放, 需要被 maintained in that dependency relationship.

Deadlock avoidance/mitigation

或者 (进阶) 通过 cycle \Rightarrow 有死锁的 transaction
或 rollback -by-

\Rightarrow 要选择好要 terminate 的 transaction
 \rightarrow 选取代价小的 (or the transaction just start)

- Pre-declare all necessary resources and allocate in a single request.
不知道它们之间哪些是先决条件 (we cannot know beforehand that)
- Periodically check the resource dependency graph for cycles. If a cycle exists - rollback (i.e., terminate) one or more transaction to eliminate cycles (deadlocks). The chosen transactions should be cheap (e.g., operations. they have not consumed too many resources).
what resources need when they require executed to many otherwise, it has not
- Allow waiting for a maximum time on a lock then force Rollback. Many successful systems (IBM, Tandem) have chosen this approach. cheap & easy
- Many distributed database systems maintain only local dependency graphs and use time outs for global deadlocks.

Main full dependency graph for distributed database is very complex \Rightarrow distributed database

only maintain local dependency graph for the protocol and resources in that particular node of that distributed system. 若是 multiple nodes \Rightarrow timeout approach.



Deadlocks ...

Deadlocks are rare, however, they do occur and the database has to deal with them when they occur

What is the probability of a deadlock occurrence? - tutorial



Summary

Concurrency problems – why we need concurrency control

Implementation of exclusive access – Dekker's algo, OS primitives, spin locks

Semaphores - get lock, release lock, maintain queue of processes

Atomic operations to get and release locks

Semaphore queue, avoiding long queues

Deadlocks



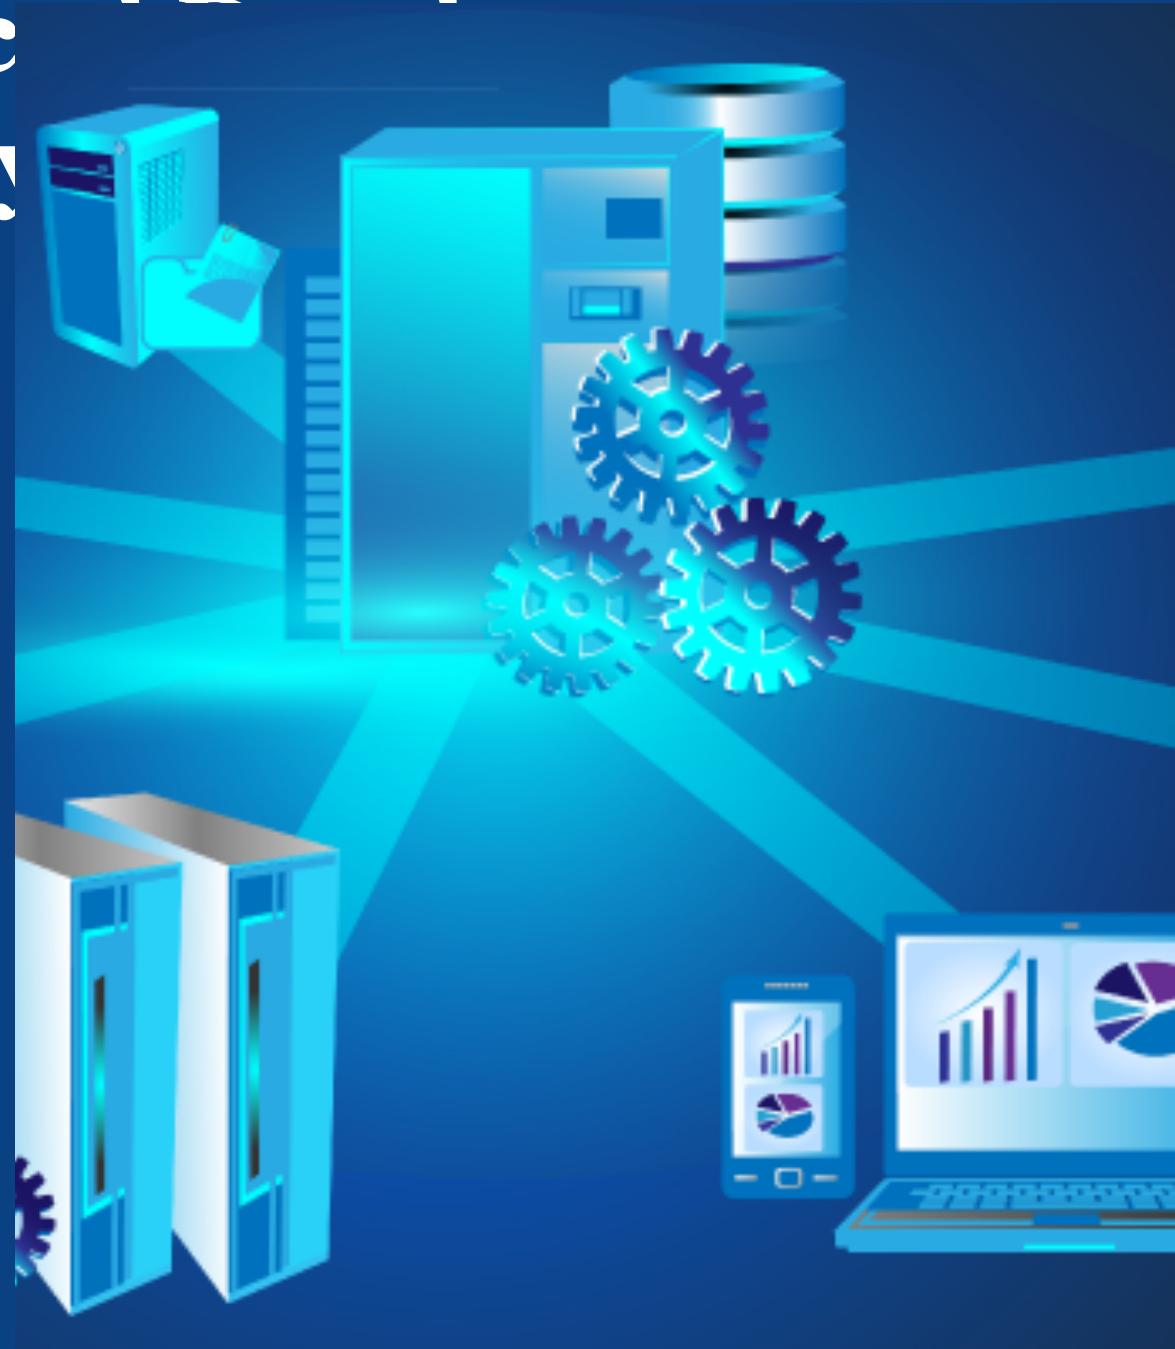
COMP90050:

Advanced Systems

Lecturer: Farhana Choudhury (PhD)

Isolation concepts

Week 6





What we have seen in previous lecture

- Problems may arise from multiple concurrent transactions
- Concurrency control is needed
 - Take exclusive access to shared resources to handle concurrency problems

In this lecture we will see the concurrency control more formally and in more details



Isolation Concepts – I in ACID

Isolation ensures that concurrent transactions leaves the database in the same state as if the transactions were executed separately.

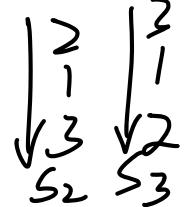
Isolation guarantees consistency, provided each transaction itself is consistent.

isolation property: if there are concurrent transactions \Rightarrow those concurrent transactions should leave the data base in the same state as if these concurrent transaction were executed separately.

T_1 —
 T_2 —

run them one by one \rightarrow separate execution order

T_3 —
 S_1



3个不同的状态应该和 S_1, S_2, S_3 的不匹配

The current state should not have different state than a



Isolation – expected output example

transaction
State run separately. *(guarantee the consistency of data base)*

Shared counter = 100;

Task1/Trans/Process/Thread
counter = counter +10;

Task2/Trans/Process/Thread
counter = counter +30;

Task1 and Task2 run concurrently.

a) counter == 110

Sequence of actions

T1: Reads counter == 100

T2: Reads counter == 100

T2: Writes counter == 100+30

T1: Writes counter == 100+10

b) counter == 130

Sequence of actions

T1: Reads counter == 100

T2: Reads counter == 100

T1: Writes counter == 100+10

T2: Writes counter == 100+30

c) counter == 140;

Sequence of actions

T1: Reads counter == 100

T1: Writes counter == 100+10

T2: Reads counter == 110

T2: Writes counter == 110+30

Alternatively, T2 executing before T1 will also have the same state



Isolation Concepts ...

We can achieve isolation by sequentially processing each transaction - generally not efficient and provides poor response times.

We need to run transactions concurrently with the following goals:

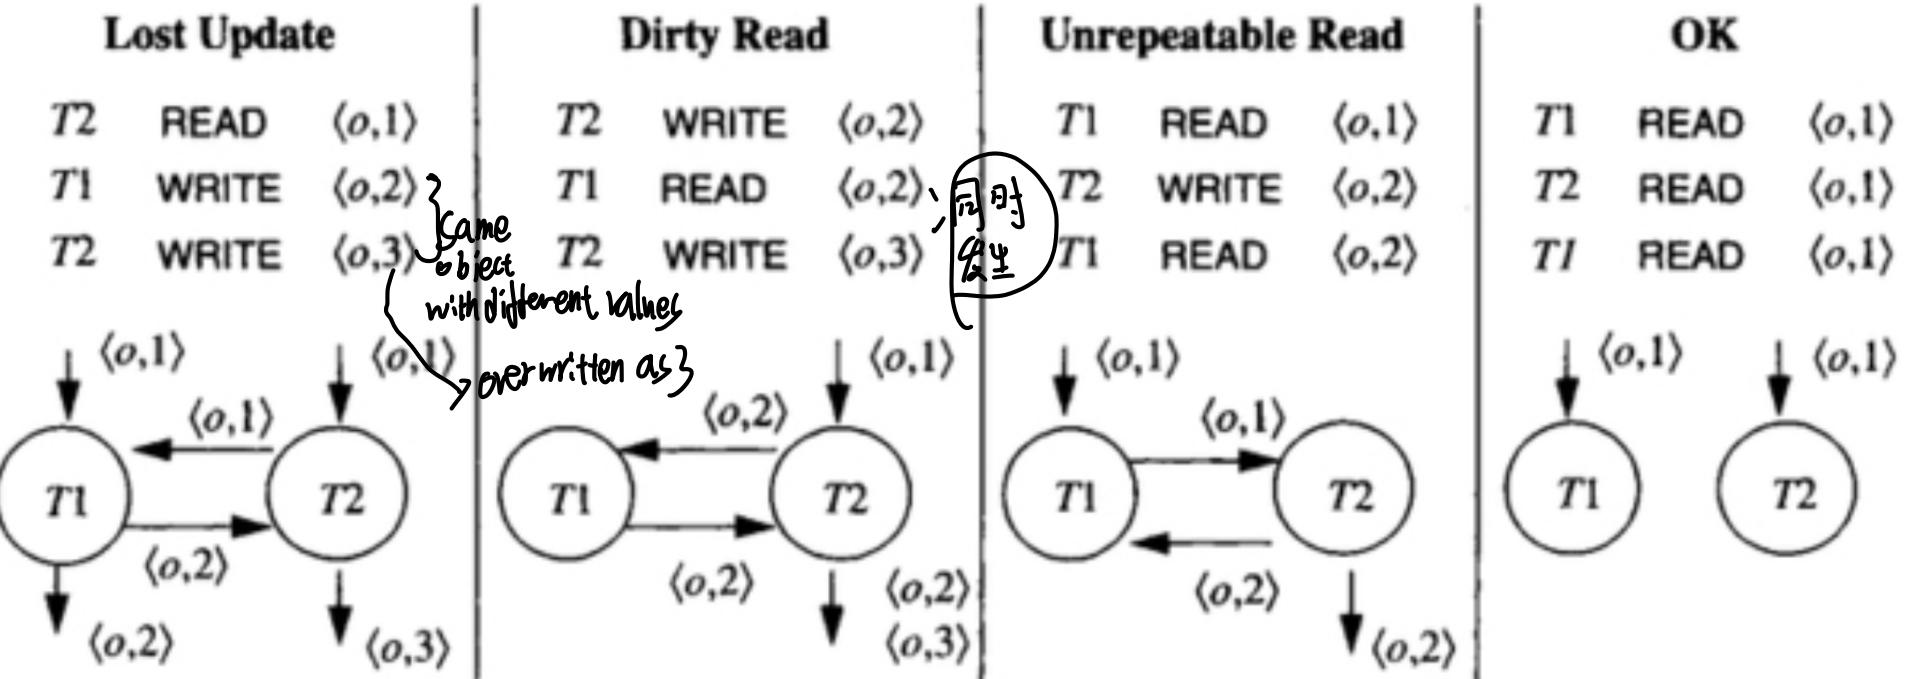
- concurrent execution should not cause application programs (transactions) to malfunction.
- Concurrent execution should not have lower throughput or bad response times than serial execution.

execute sequentially \Rightarrow no concurrency but low efficiency \Rightarrow ~~慢~~ ^{耗時} poor response time

To achieve isolation we need to understand dependency of operations

State of running transactions concurrently = state of running transactions separately,

Possible dependencies





How can we find the dependencies?

Given a set of transactions, how can we determine which transaction depends on which other transaction?

set of records

Dependency model

- I_j : set of inputs (objects that are read) of a transaction T_j ^{read by}
 - O_i : set of outputs (objects that are modified) of a transaction T_i
- Note O_j and I_j are not necessarily disjoint that is $(O_j \cap I_j \neq \emptyset)$

Given a set of transactions , Transaction T_j has no dependency on any particular transaction T_i if -

$$O_i \cap (I_j \cup O_j) = \text{empty for all } i \neq j$$

none of this objects are \rightarrow all the objects that are either read or written by T_j

This approach cannot be planned ahead as in many situations inputs and outputs may be state dependant/not known in prior.

if that's true T_j is transaction has no dependency on any of this other transaction T_i

the set of objects
 whether it's ready or writey \Rightarrow it will never be modified by any other transaction
 rather T_j is operating
 $O_i \cap (I_j \cup O_j) = \text{empty for all } i \neq j$
 output: modified
 \nearrow Input if they are not disjoint $\Rightarrow I_j \cap O_i$ not hold

| | ready T ₁ | T ₂ | T ₃ | T ₄ |
|--|---|---|---|---|
| Trans In \cup Out (I _i \cup O _i) | O _i \cap (I ₁ \cup O ₁) | O _i \cap (I ₂ \cup O ₂) | O _i \cap (I ₃ \cup O ₃) | O _i \cap (I ₄ \cup O ₄) |
| T ₁ | O ₁ | \emptyset | \emptyset | \emptyset |
| T ₂ | \emptyset | O ₂ | \emptyset | \emptyset |
| T ₃ | \emptyset | \emptyset | O ₃ | \emptyset |
| T ₄ | \emptyset | \emptyset | \emptyset | O ₄ |

4 trans
 run
 concurrently

T₁, ..., T₄ 互斥且 those not have any relationship with
 any other transactions.

∂ : output modification
 I : input ready

If the inputs and outputs of the concurrent transactions are not disjoint, the following dependencies can occur –

The transaction execution sequences

T_1 READ $\langle o,1 \rangle$

T_2 WRITE $\langle o,2 \rangle$

T_1 WRITE $\langle o,2 \rangle$

T_2 READ $\langle o,2 \rangle$

T_1 WRITE $\langle o,2 \rangle$

T_2 WRITE $\langle o,3 \rangle$

The dependency graphs

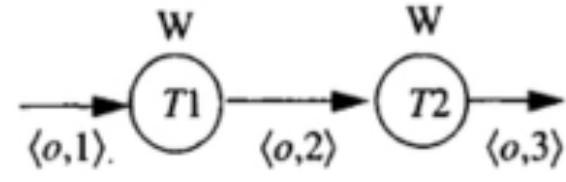
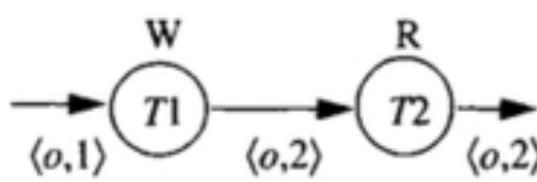
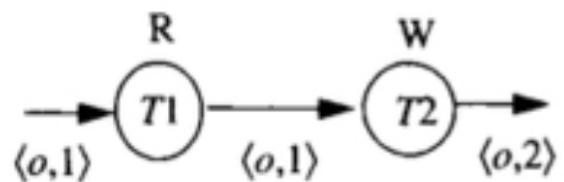
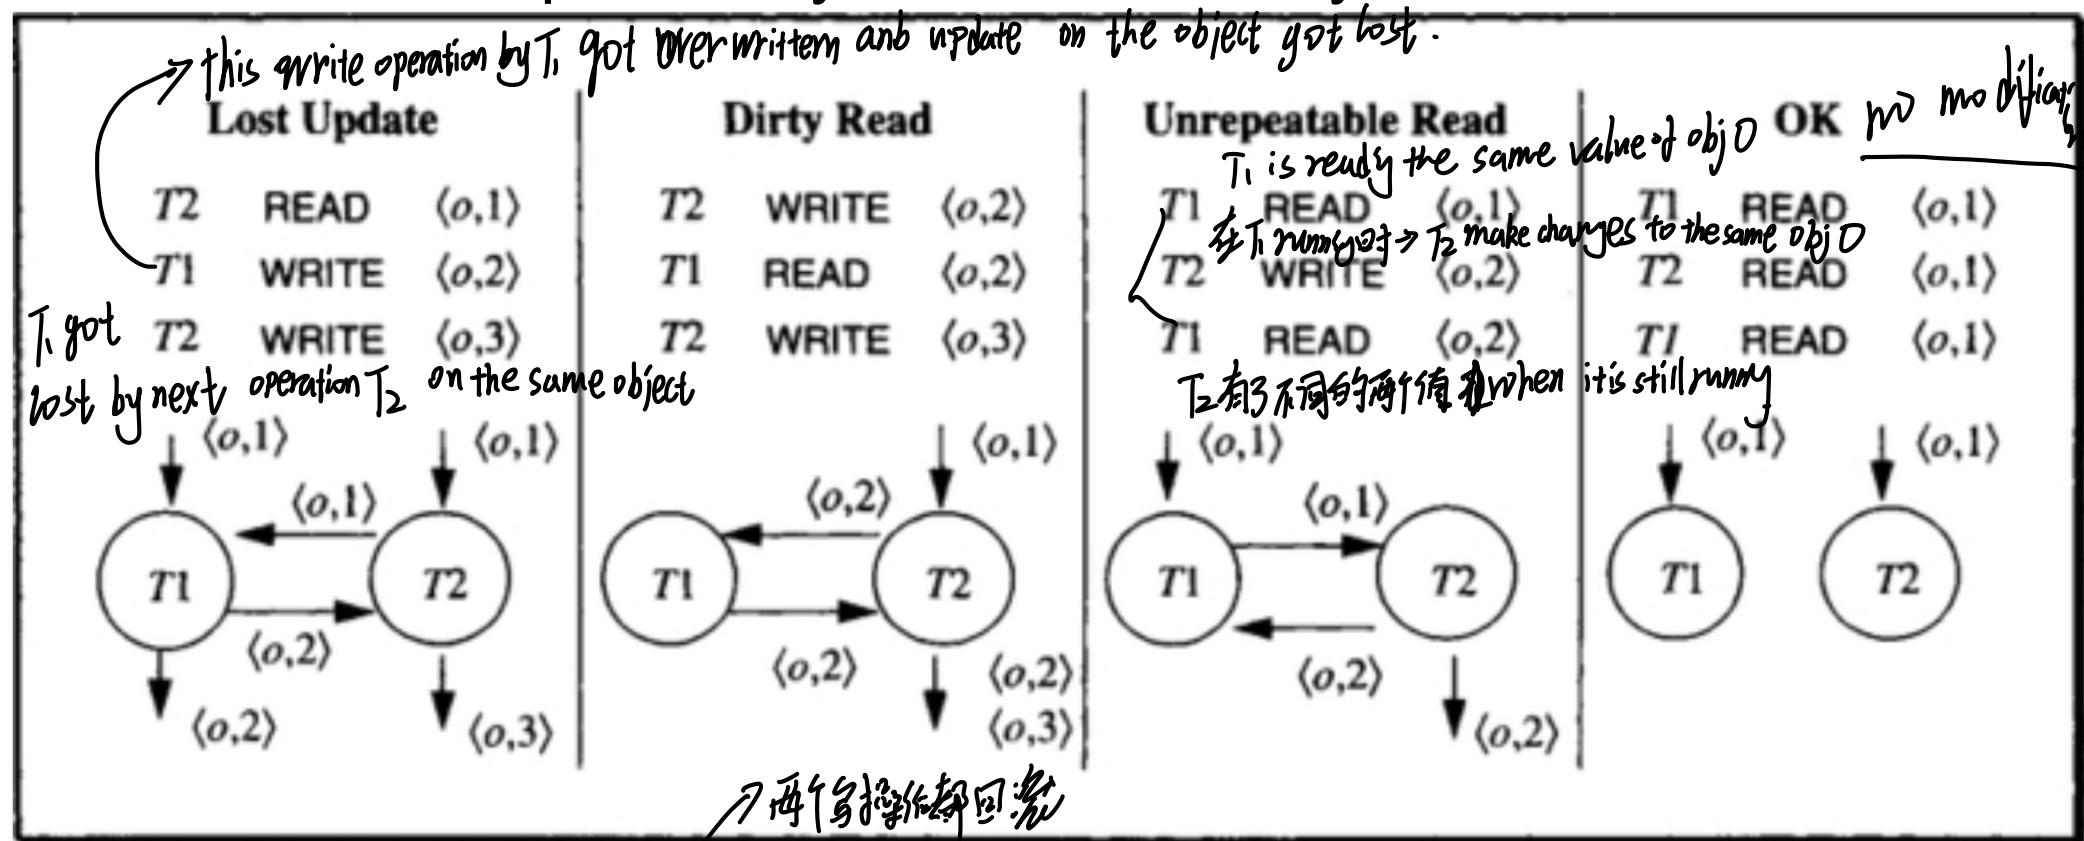


Fig 7.1 in the main reference book

Read-Read dependency do not affect isolation

Dependencies

When dependency graph has cycles then there is a violation of isolation and a possibility of inconsistency.



% T2 did not see the update of T1

What if T2 aborts?
 T1's read will be invalid (inconsistency)

The value of o changes by another transaction T2 while T1 is still running (inconsistency) if T2 aborts ⇒ last read was lost.



Formal definition of dependency

What are history of transaction is sequence of actions performed by the Transaction

Let H is a history sequence of tuples of the form (T, action, object).

Let T1 and T2 are transactions in H. If T1 performs an action on an object O, then T2 performs an action on the same O, and there is no write action in between by another transaction on O – then T2 depends on T1.

T₂ has dependency on T₁ based on object O

Formally, the dependency of T2 on T1 (T_1, O, T_2) exists in history H if there are indexes i and j such that $i < j$, $H[i]$ involves action a_1 on O by T_1 , (i.e., $H[i] = (T_1, a_1, O)$) and $H[j]$ involves action a_2 on O by T_2 (i.e., $H[j] = (T_2, a_2, O)$) and there are no other $H[k] = (T', \text{WRITE}, O)$ for $i < k < j$

Dependency graph : Transactions are nodes, and object labels the edges from the node T_i to T_j if (T_i, O, T_j) is in $\text{DEP}(H)$.

T₁ make decisions based on its ready
any inconsistency $\frac{T_1}{T_2}$ $\frac{T_2}{T_1}$
 $\frac{T_1}{T_2}$ $\frac{T_2}{T_1}$ $\frac{T_1}{T_3}$ $\frac{T_3}{T_1}$ $\frac{T_1}{T_4}$ $\frac{T_4}{T_1}$



Dependency relations

We focus on the dependency in three scenarios

- $a_1 = \text{WRITE}$ & $a_2 = \text{WRITE}$;
- $a_1 = \text{WRITE}$ & $a_2 = \text{READ}$;
- $a_1 = \text{READ}$ & $a_2 = \text{WRITE}$ (dependency as T1 may read again after a2).



Dependency relations - equivalence

$\text{DEP}(H) = \{(T_i, O, T_j) \mid T_j \text{ depends on } T_i\}$.

Given two histories H_1 and H_2 contain the same tuples, H_1 and H_2 are equivalent if $\text{DEP}(H_1) = \text{DEP}(H_2)$

(tuple-元组, 操作不同 但 $\text{DEP}(H_1) = \text{DEP}(H_2) \Leftrightarrow$ equivalent)

This implies that a given database will end up in exactly the same final state by executing either of the sequence of operations in H_1 or H_2

E.g., sequentially look into this action:
first object O_1 another transaction performing action on the O_1

$H_1 = \langle (T_1, R, O_1), (T_2, W, O_5), (T_1, W, O_3), (T_3, W, O_1), (T_5, R, O_3), (T_3, W, O_2), (T_5, R, O_4), (T_4, R, O_2), (T_6, W, O_4) \rangle$
 T_1 , read on O_1 , T_3 write on O_1 , T_5 read on O_3 , transaction T_3 on O_1 has dependency $\Rightarrow T_3$ has dependency on T_1 , based on the object O_1

$\text{DEP}(H_1) = \{\langle T_1, O_1, T_3 \rangle, \langle T_1, O_3, T_5 \rangle, \langle T_3, O_2, T_4 \rangle, \langle T_5, O_4, T_6 \rangle\}$
dependency relationship of this history
relationship on T_1 based on O_1

$H_2 = \langle (T_1, R, O_1), (T_3, W, O_1), (T_3, W, O_2), (T_4, R, O_2), (T_1, W, O_3), (T_2, W, O_5), (T_5, R, O_3), (T_5, R, O_4), (T_6, W, O_4) \rangle$

$\text{DEP}(H_2) = \{\langle T_1, O_1, T_3 \rangle, \langle T_1, O_3, T_5 \rangle, \langle T_3, O_2, T_4 \rangle, \langle T_5, O_4, T_6 \rangle\}$

$\boxed{\text{DEP}(H_1) = \text{DEP}(H_2)}$

$$H = \{ (T1, R, O1), (T3, W, O3), (T3, W, O1), (T2, R, O3), (T1, R, O1), (T2, W, O2), (T1, R, O2), (T1, R, O3), (T2, R, O2) \}$$

$$DEP(H) = \{ \langle T_1, O_1, T_3 \rangle, \langle T_3, O_3, T_2 \rangle, \langle T_2, O_2, T_1 \rangle \}.$$

E.g.,

$$H1 = \{ (T1, R, O1), (T2, W, O5), (T1, W, O3), (T3, W, O1), (T5, R, O3), (T3, W, O2), (T5, R, O4), (T4, R, O2), (T6, W, O4) \}$$

$$DEP(H1) = \{ \langle T_1, O_1, T_3 \rangle, \langle T_1, O_3, T_5 \rangle, \langle T_3, O_2, T_4 \rangle, \langle T_5, O_4, T_6 \rangle \}$$

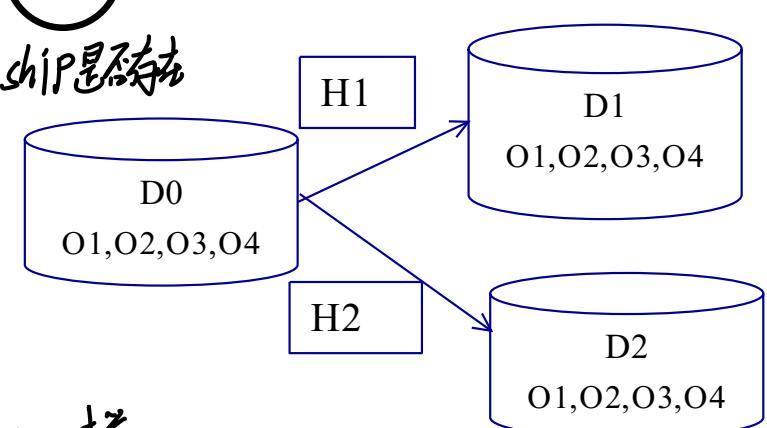
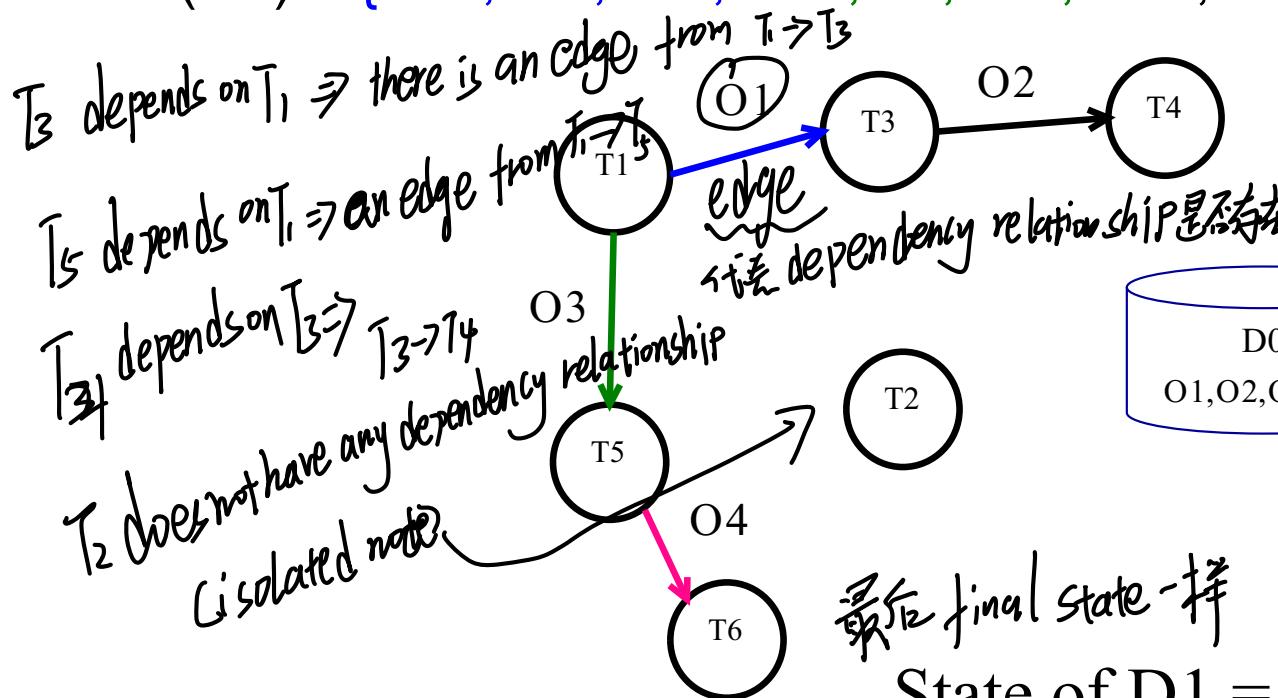
Dependency relations - equivalence

$H1 = \langle (T1, R, O1), (T2, W, O5), (T1, W, O3), (T3, W, O1), (T5, R, O3), (T3, W, O2), (T5, R, O4), (T4, R, O2), (T6, W, O4) \rangle$

$H2 = \langle (T1, R, O1), (T3, W, O1), (T3, W, O2), (T4, R, O2), (T1, W, O3), (T2, W, O5), (T5, R, O3), (T5, R, O4), (T6, W, O4) \rangle$

$\text{DEP}(H1) = \{ \langle T1, O1, T3 \rangle, \langle T1, O3, T5 \rangle, \langle T3, O2, T4 \rangle, \langle T5, O4, T6 \rangle \}$

$\text{DEP}(H2) = \{ \langle T1, O1, T3 \rangle, \langle T1, O3, T5 \rangle, \langle T3, O2, T4 \rangle, \langle T5, O4, T6 \rangle \}$



State of D1 = State of D2

Isolated history

A history is said to be isolated if it is equivalent to a serial history (as if all transactions are executed serially/sequentially)

A serial history is history that is resulted as a consequence of running transactions sequentially one by one. N transactions can result in a maximum of $N!$ serial histories.

If T_1 precedes T_2 ,
 it is written as $T_1 \ll T_2$. T_1 先发生

$$\text{Before}(T) = \{T' \mid T' \ll T\}$$

$$\text{After}(T) = \{T' \mid T \ll T'\}$$

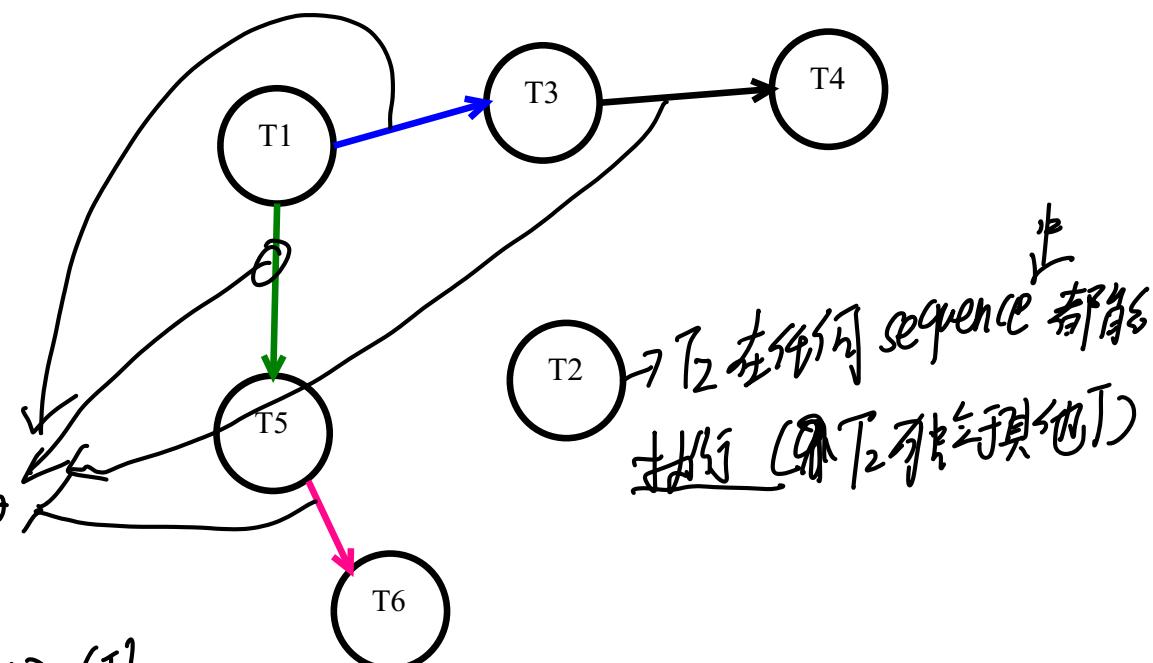
before(T_1) = { T_2, T_3, T_4, T_5 }

$$\text{E.g. } \text{After}(T_1) = \{T_5, T_6, T_3, T_4\}$$

$$\text{After}(T_3) = \{T_4\} \quad \text{before}(T_3) = \{T_1\}$$

$$\text{After}(T_5) = \{T_6\}$$

$$\text{After}(T_3) = \{T_4\}$$





Isolation Concepts ...

A transaction T' is called a wormhole transaction if

$$T' \in \text{Before}(T) \cap \text{After}(T)$$

That is $T << T' << T$. This implies there is a cycle in the dependency graph of the history. Presence of a wormhole transaction implies it is not isolated (\Rightarrow not a serial schedule).

\because history cannot be isolated \Rightarrow \exists equivalent serial history.

A history is serial if it runs one transaction at a time sequentially, or equivalent to a serial history. Serial history \Rightarrow run one by one sequentially

A serial history is an **isolated** history.

Wormhole theorem: A history is isolated if and only if it has no

wormholes. Wormhole: \exists cycle in dependency graph?

isolated history: \nexists wormhole, \exists wormhole \Rightarrow not isolated history



serial history isolated history \Leftrightarrow no wormhole.

We will now introduce a new type of lock -

SLOCK (shared lock) that allows other transactions to ~~read~~, but
~~not write/modify~~ the shared resource

The wormhole transaction concept will be useful in a later topic!



To grant lock or not to...

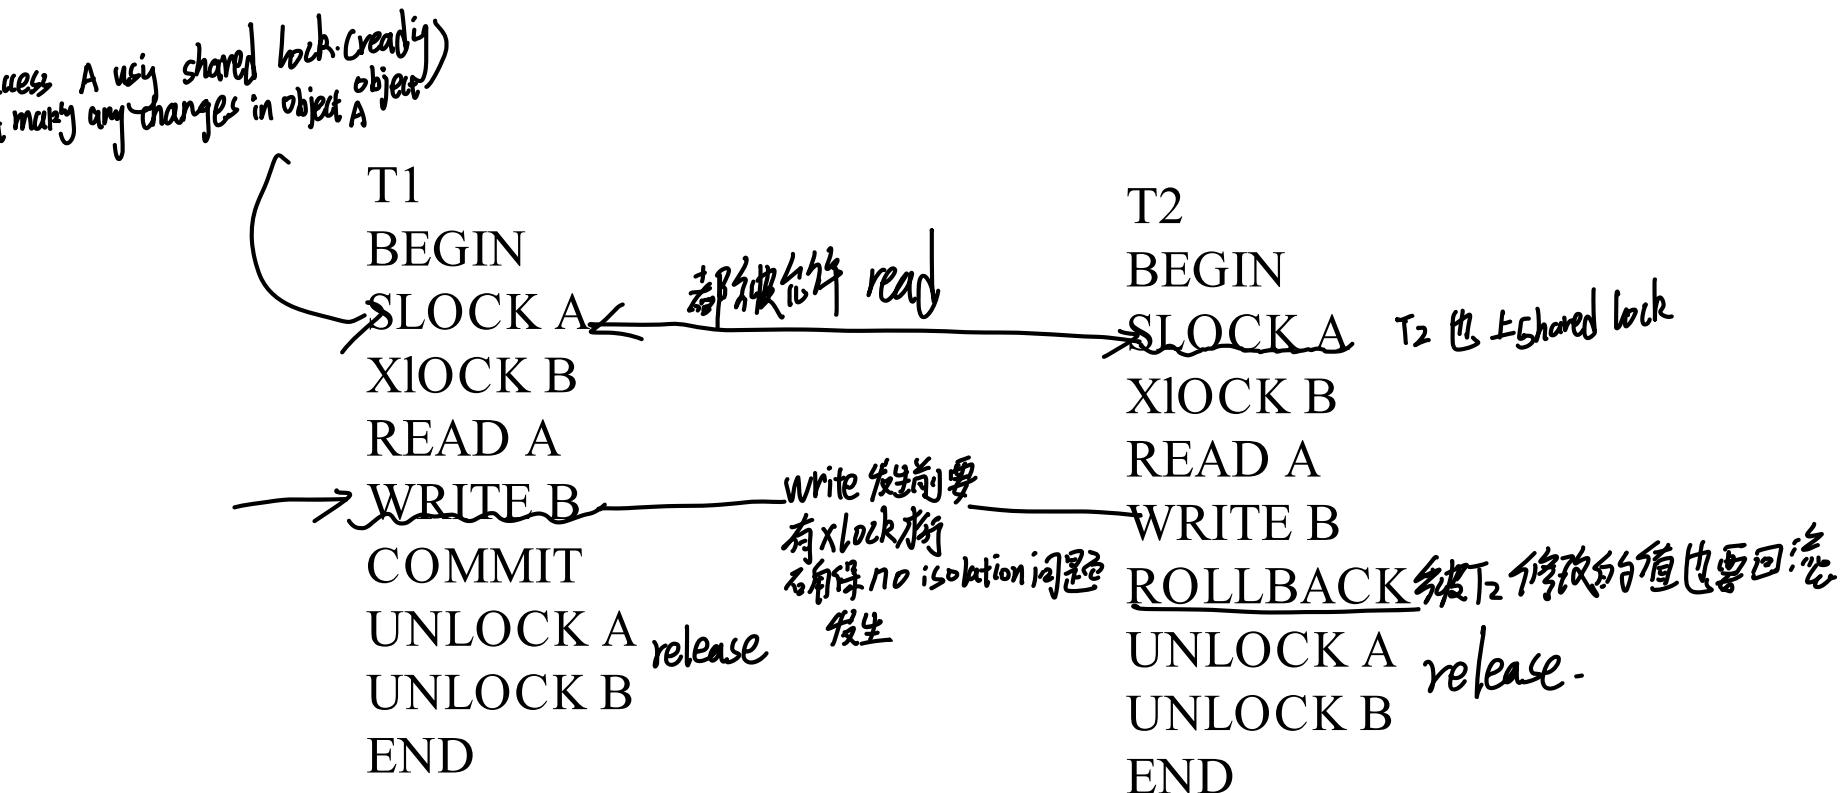
A lock on an object should not be granted to a transaction while that object is locked by another transaction in an **incompatible mode**.

Lock Compatibility Matrix

| Current Mode | Mode of Lock | | |
|---|--|---|--|
| | Free | Shared | Exclusive |
| Shared request (SLOCK) Used to block others writing/modifying | Compatible Request granted immediately Changes Mode from Free to Shared | Compatible Request granted immediately Mode stays Shared | Conflict Request delayed until the state becomes compatible Mode stays Exclusive |
| Exclusive request (XLOCK) Used to block others reading or writing/modifying | Compatible Request granted immediately Changes Mode from Free to Exclusive | Conflict Request delayed until the state becomes compatible Mode stays Shared | Conflict Request delayed until the state becomes compatible Mode stays Exclusive |

When to use what type of lock

Actions in Transactions are: READ, WRITE, XLOCK, SLOCK,
UNLOCK, BEGIN, COMMIT, ROLLBACK





Isolation Concepts ...

BEGIN, END , SLOCK, XLOCK can be ignored as they can be automatically inserted in terms of the corresponding operations

E.g. if a transaction ends with a COMMIT, it is replaced with:

{UNLOCK A if SLOCK A or XLOCK A appears in T for any object A}.

(That is to simply release all locks)

Commit = UNlock A

Similarly ROLLBACK can be replaced by

Write(undo) A = rollback

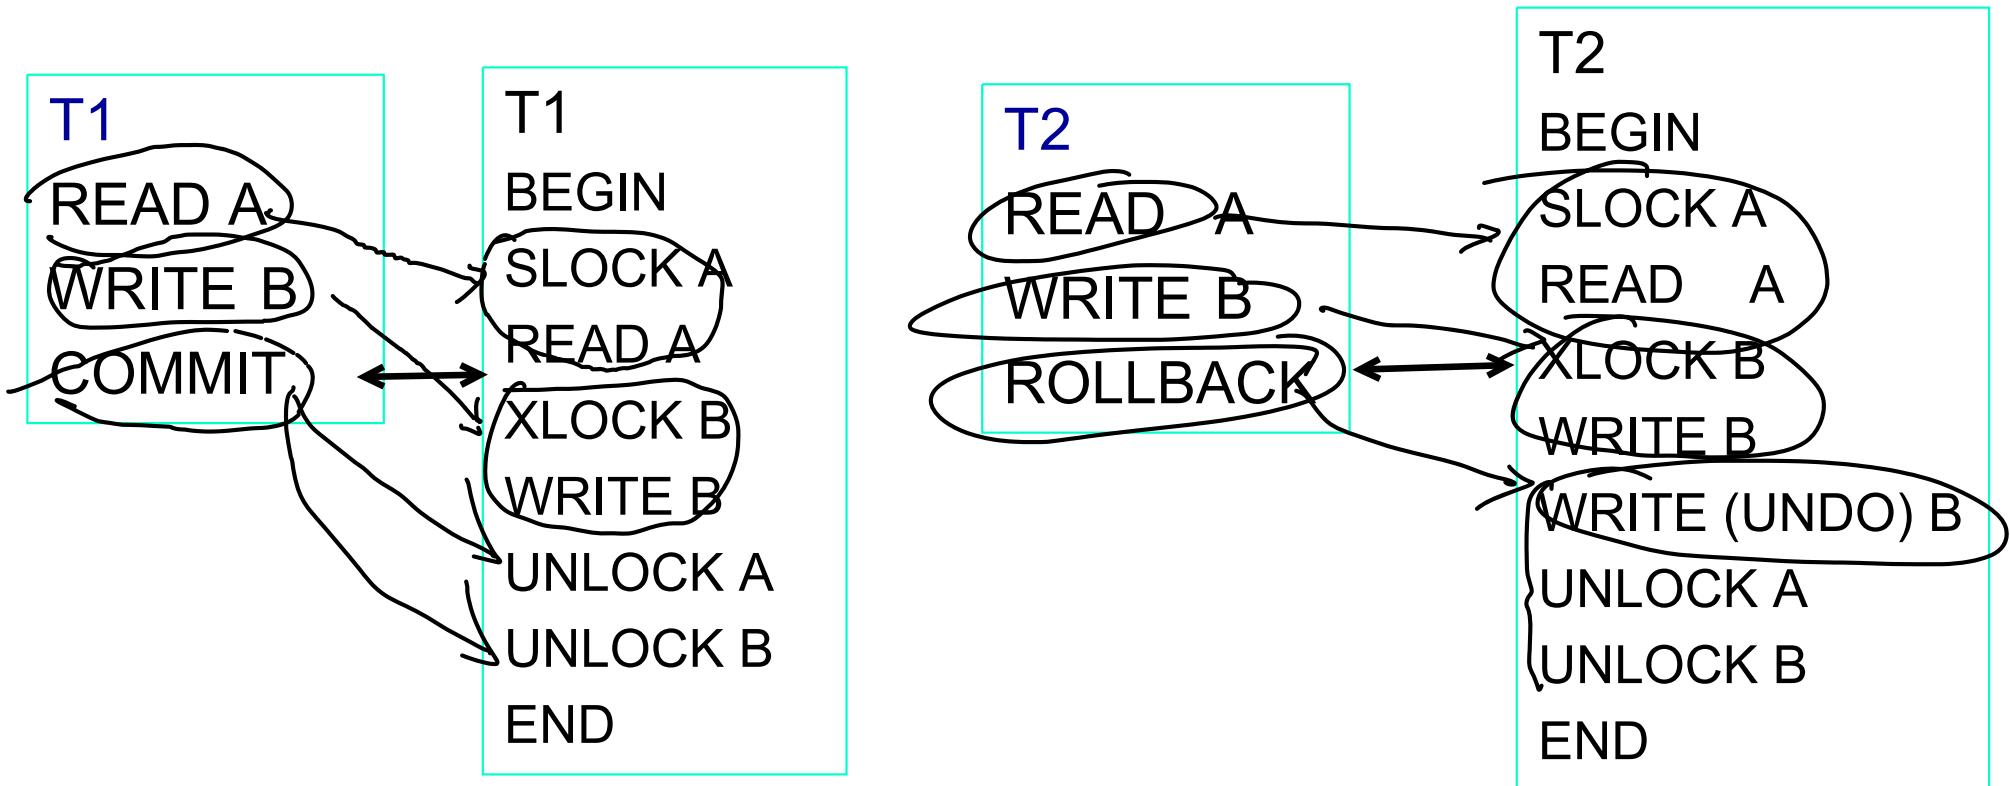
{WRITE(UNDO) A if WRITE A appears in T for any object A}

{ UNLOCK A if SLOCK A or XLOCK A appears in T for any object A}.



Isolation Concepts ...

We can replace previous transaction sequences by:





Isolation Concepts ...

well-formed transaction: $\text{① Read} \rightarrow \text{lock}$ or $\text{② Write} \rightarrow \text{lock}$. ③ Unlock ~~在 lock~~

Well-formed transactions: A transaction is well formed if all READ, WRITE and UNLOCK operations are covered by appropriate LOCK operations

Two phase transactions: A transaction is two phased if all LOCK operations precede all its UNLOCK operations. $\xrightarrow{\text{all lock 在 all unlock}} \quad \xrightarrow{\text{所有 lock 在所有 unlock 之前}}$

ObjA - lock
ObjB - lock

✓ \Rightarrow

objA - lock
objA - unlock

unlock before lock (right)

X

ObjA - unlock
ObjB - unlock

objB - lock
objB - unlock

not a two phase transaction



Isolation Theorems

Summary:

A transaction is a sequence of READ, WRITE, SLOCK, XLOCK actions on objects ending with COMMIT or ROLLBACK.

A transaction is **well formed** if each READ, WRITE and UNLOCK operation is covered earlier by a corresponding lock operation.

A history is **legal** if it does not grant conflicting grants. *locks*

A transaction is **two phase** if its all lock operations precede its unlock operations.



Isolation Theorems

Locking theorem: If all transactions are **well formed** (READ, WRITE and UNLOCK operation is covered earlier by a corresponding lock operation) and **two-phased** (locks are released only at the end), then any **legal** (does not grant conflicting grants) history will be isolated.

Locking theorem (Converse): If a transaction is **not well formed** or is **not two-phase**, then it is possible to write another transaction such that it is a wormhole.
any not well formed or not two-phase transaction is a wormhole.

Rollback theorem: An update transaction that **does an UNLOCK** and **then does a ROLLBACK** is **not two phase**.

When a transaction unlocks (releases) a resource before making a final decision (commit or roll back), it leaves the possibility open for other transactions to access and modify that object. This disrupts the isolation property because other transactions could see intermediate potentially inconsistent result of data which goes against the ACID properties.

Degrees of Isolation

highest level of isolation

Degree 3: A Three degree isolated Transaction has no lost updates, and has repeatable reads.

This is “true” isolation. all of operations have their preceding appropriate lock instructions

Lock protocol is two phase and well formed.

It is sensitive to the following conflicts:

write->write; write ->read; read->write

if it's just ready but not writing → shared locks is good; if modification is done → exclusive lock comes after lock operation. needs to be taken.

Degree 2: A Two degree isolated transaction has no lost updates and no dirty reads.

Lock protocol is two phase with respect to exclusive locks and well formed with respect to Reads and writes. (May have Non repeatable reads.)

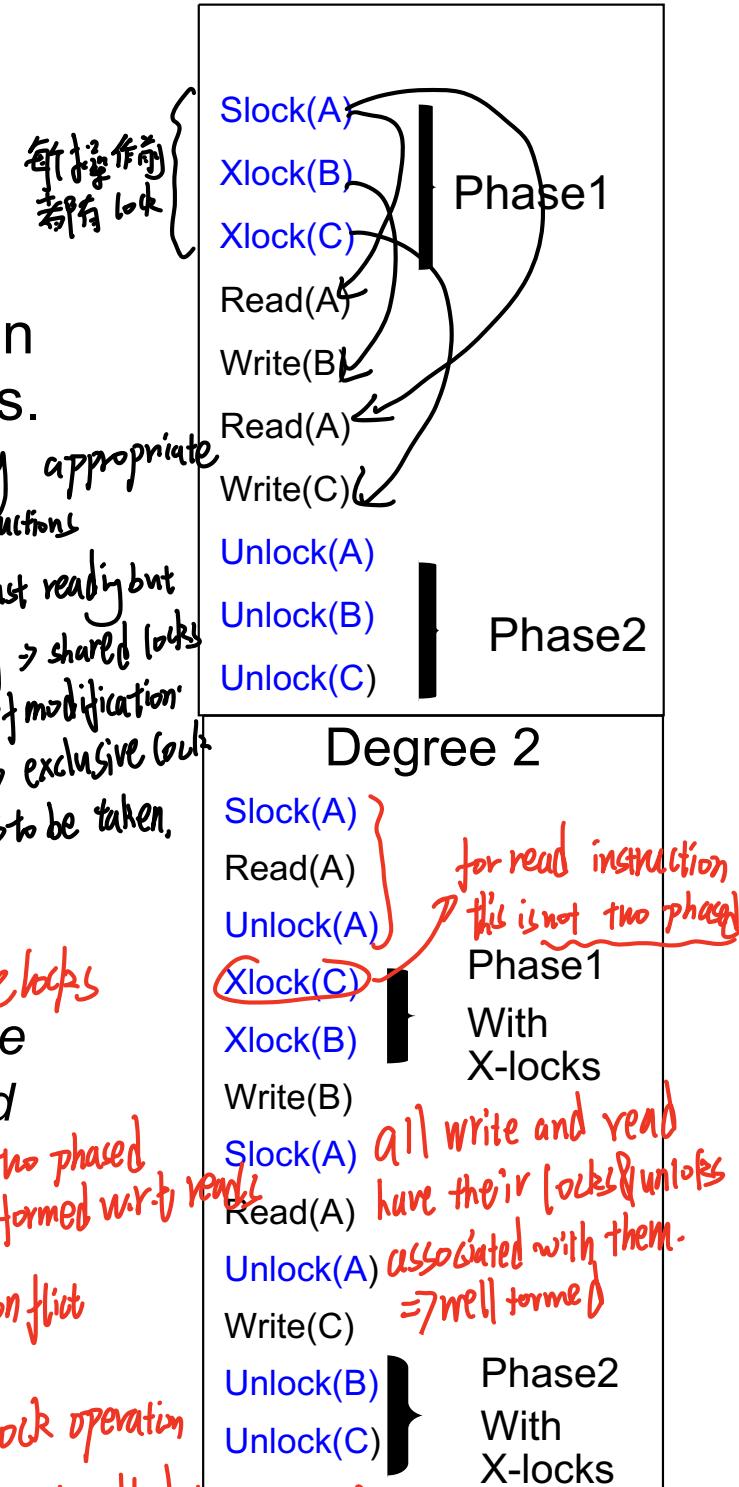
It is sensitive to the following conflicts:

write->write; write ->read;

cannot handle read->write conflict

For reads operation → the unlock operation comes before some other lock operation

For reads → Not two phased ⇒ Non repeatable reads (can not be handled by Degree 2)



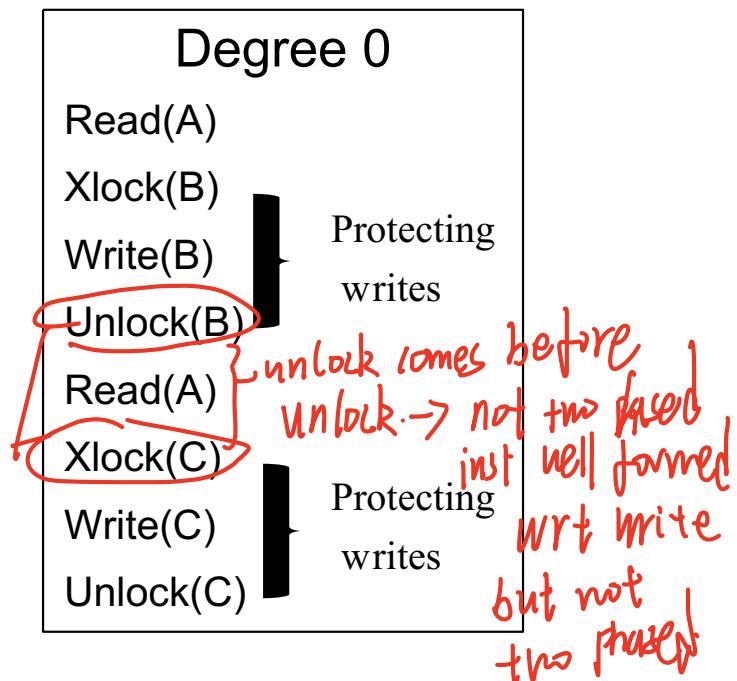
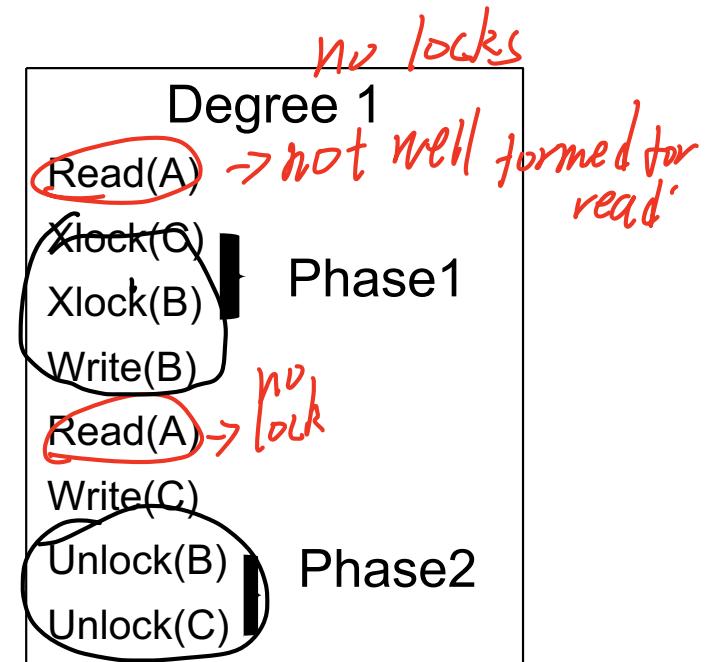
Degree 1: A One degree isolation has no lost updates. *if read(A) operation does not have any locks \Rightarrow degree 1 isolation*
Lock protocol is two phase with respect to exclusive locks and well formed with respect to writes.

It is sensitive the following conflicts:
write->write; *(read->write or write->read)* *都无法处理*

Degree 0 : A Zero degree transaction does not overwrite another transactions dirty data if the other transaction is at least One degree.

Lock protocol is well-formed with respect to writes.

It ignores all conflicts.???



read-write ? $\text{Read}(A) \rightarrow \underline{\text{xlock}(A)} \rightarrow \underline{\text{rite}}$
 write-read ? $\underline{\text{xlock}(A)} \rightarrow \underline{\text{read}(A)} \rightarrow \underline{\text{R}}$
 unlock(A)
 write-write ? $\underline{\text{xlock}(A)} = \underline{\text{xlock}(A)}$
 write.
 $\underline{\text{unlock}}$

