



# COMP90050 ADBS

授课老师: Rita





# 1 – Indexing

字典目录（帮助快速定位）





## DBMS

DBMS admin 管理 index

- o DBMS admin generally creates indices to allow almost direct access to individual items (index = indice)

- o DBMS must support:

- insert/ delete/ modify record 增删改
- read a particular record (specified using record id) (读特定 record)
- scan all records (possibly with some conditions on the records to be retrieved), or scan a range of records





primary key 索引

加速查詢

## Indexing

- Indexing mechanisms used to speed up access to desired data in a similar way to look up a phone book or dictionary

{ ID } -> { name } { ID, name }

- Search Key** - attribute or set of attributes used to look up records/ rows in a system like an ID of a person (The ID will take us to get the entire record)

- An **index file** consists of records (called index entries) of the form search-key, pointer to where data is

index file → 

age 5.
-----------

 nrecords << original data

- Index files are typically much smaller than the original data files and many parts of it are already in main memory (main memory is faster than disk)

↓ size 小、速度快





Index 可以加速存取 → index 能直接帮助指定一个找位置  
Indexing makes **Disk Access** faster through

how {

- ① records with a **specified value** in the attribute accessed with minimal disk accesses (e.g., Student ID = 101, the index file tells us which data block in disk to go)  
*index file 有着一些 specified value 以便找到存储位置。*
  - ② records with an attribute value falling in a **specified range** of values can be retrieved with a single seek and then consecutive sequential reads (e.g., ID:100- 200, find where is the first record (single seek of starting block), then read sequentially)  
*index file 先做位置 → 然后顺序读取*
- Note:** not all DB stores sequential data sequentially, so this method may not useful for some DB  
*(有些数据库是跳跃存取的, specified range 不好使)*





## Criteria to choose index – always ~~tradeoff~~

- 插入时间  
Insertion time to index is also important
- 删除时间  
Deletion time is important as well
- 重排重组 不能变动太大  
No big index rearrangement after insertion and deletion
- 空间开销  
Space overhead needs to be considered for the index itself
- 没有单一的索引技术是最好的。相反，每种技术最适合特定的应用程序。  
No single indexing technique is the best. Rather, each technique is best suited to particular applications.





## Types of indices based on search keys

- **Ordered indices** – Search keys are stored in some order
    - Clustering index / **primary index**
    - Non-clustering index / **secondary index**
  - **Hash indices** – Search keys are distributed
  - **B+tree**
  - **Bitmap index**
  - **Quadtree (k-d tree)** → 向具体是什么
  - **R-tree**
- 考得多 → 很多情况下适合 database，search key 是什么？
- 疑问：index原理、优点上，什么情况下不适合？什么情况下表现更好。





data is ordered based on salary  $\Rightarrow$  index is also ordered based on the same key.  $\Rightarrow$  clustered index

## Clustering Index / Primary Index

in a sequentially ordered file, the index whose search key specifies the sequential order of the file

(The search key of a primary index is usually but not necessarily the primary key)

eg: 有序文件 ID name favorite subj Salary 升序大

ID	name	favorite subj	Salary
10101	Srinivasan	Comp. Sci.	65000
12121	Wu	Finance	90000
15151	Mozart	Music	40000
22222	Einstein	Physics	95000
32343	El Said	History	60000

W1  
W2  
W3  
Z1  
Z2  
Z3  
Z4

这位数据在 disk 上的位置  $\Rightarrow$  disk access time





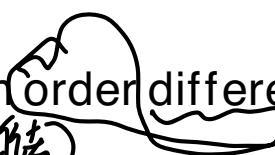
Index在1D上：要选择去 retrieve every row  $\Rightarrow$  query optimizer会选择 non-index query path  $\Rightarrow$  using index is over head  
COMP90050 ADBS

Cannot provide any benefit  $\otimes$  but the cost of index is minimum  $\Rightarrow$  just stay in main memory.

| Index subject:

## Non-clustering Index / Secondary Index

data is not ordered based on that column  $\Rightarrow$  it would be non-clustering index.

an index whose search key specifies an order different from the sequential order of the file  
  
(不按川序存儲)

(Secondary indices improve the performance of queries that use keys other than the search key of the clustering index.)



ID	name	dep
1	Jane	SC
2	John	Biology

ID	position	Salary
1	Lecturer	75000
1	PA	40000
2	SL	82000

find the name of instructors with salary  $> 70000$

① No index

② Bitree index on Salary

③ 二叉搜索树到大根堆的转化  $\Rightarrow$  堆



Answer: 若每个人的 Salary 都大于 80000  $\Rightarrow$  performance are the same

若每个人 Salary > 0  $\Rightarrow$  查的是通过 Salary > 0 的  $\Rightarrow$  -样的 performance

Conclusion 如果是那种需要 retrieve every row of table 的查询  $\Rightarrow$  有无 index 是一样的.  
or the row satisfy the condition  
 $\Rightarrow$  we have to retrieve the whole table  
then do the join.

这个 index 是在 join 前的  $\Rightarrow$

join 后无法使用，先 index  $\Rightarrow$  再 join  
index

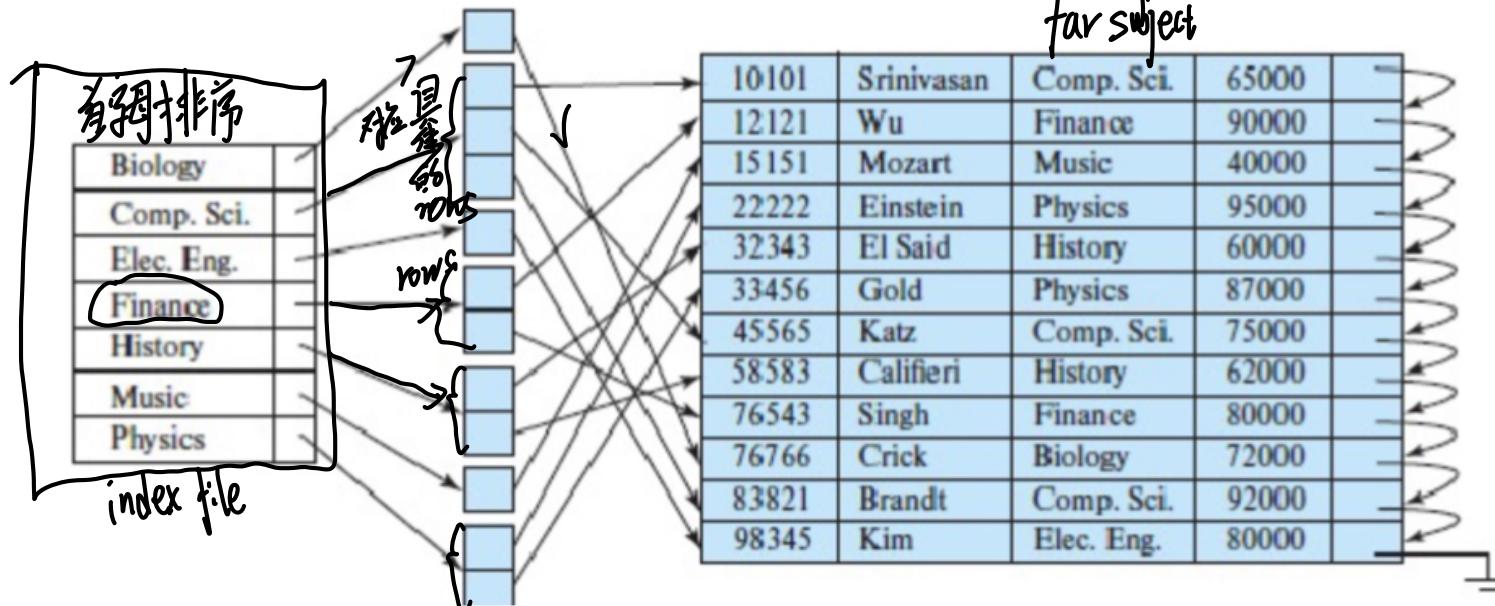
$\Rightarrow$  it will not reduce  
any number of  
pages

non



不是按“行”排序  $\Rightarrow$  而是找  $\downarrow$  search key  $\rightarrow$  找 search key  
的“列”排序

## Non-clustering Index / (Secondary Index) - e.g.





## Hash indices

search keys are distributed hopefully uniformly across “buckets” using a “function”

- A **hash index** organizes the search keys, with their associated record pointers, into a hash file structure. Order is not important. (与前两种不同，不看任何次序) 完全按 hash function 来。
- Given a key, the aim is to find the related record on file in one shot which is important.
- An ideal hash function is **uniform**, i.e., each bucket is assigned the same number of search-key values from the set of all possible values. 每桶均匀分布原则
- Ideal hash function is **random**, so each bucket will have the same number of records assigned to it irrespective of the actual distribution of search-key values in the file. (随机数据分布无关)
- Typical hash functions perform computation on the internal binary representation of the search-key. 在 attribute 上做二进制计算





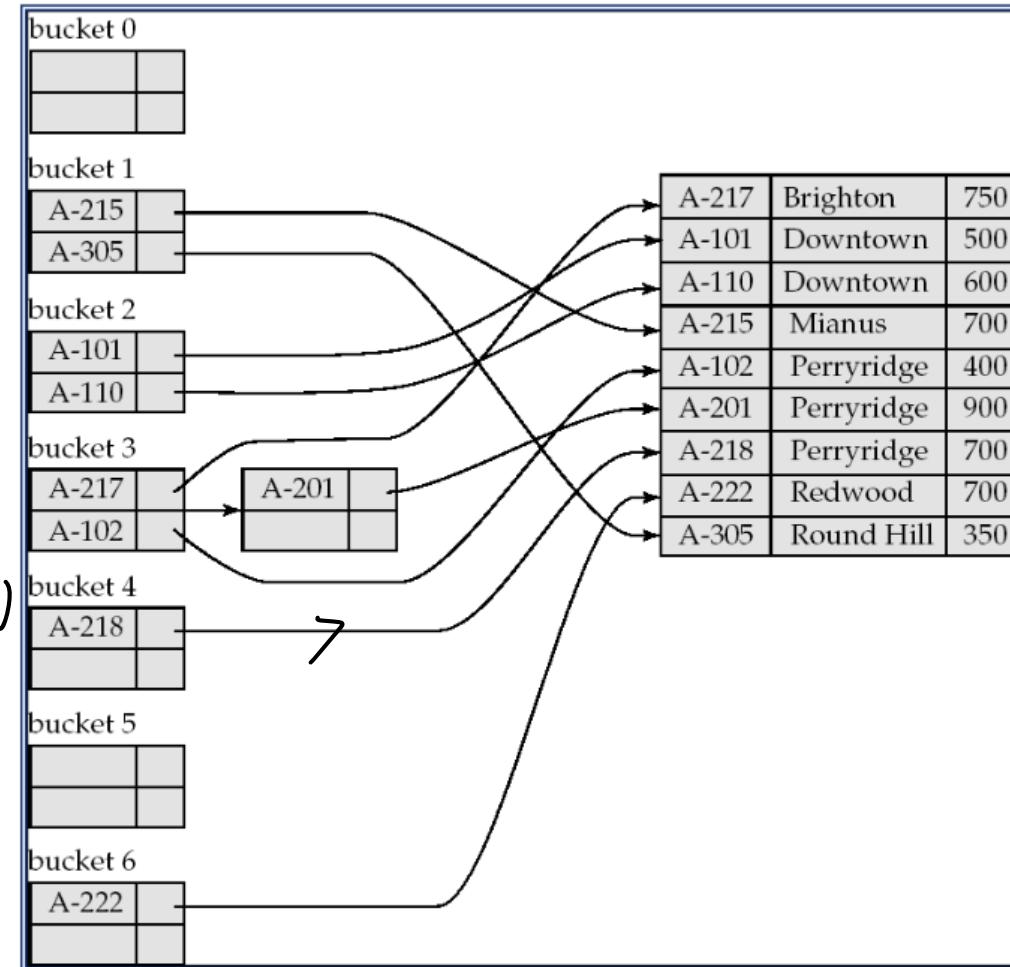
## Hash index – e.g.

hash function:  $\text{hash}(x) = y$

bucket (桶): 0 - 6 共 7 桶，第 6 桶

$\text{Hash}(A-215) = 1 \Rightarrow \text{bucket } 1 \Rightarrow \text{Row}$

$\text{Hash}(A-218) = 4 \Rightarrow \text{bucket } 4 \Rightarrow \text{Row}$





## Bitmap index

- Records in a relation are assumed to be numbered sequentially from, say 0
- Applicable on attributes that take on **a relatively small number of distinct Values**
  - e.g. gender, country, state, ...
  - e.g. income-level (income broken up into a small number of levels such as 0-9999, 10000- 19999, 20000- 50000, 50000- infinity)
- A bitmap is simply an array of bits bit map
- In its simplest form, a bitmap index on an attribute has a bitmap for each value of the attribute
  - Bitmap has as many bits as records
  - In a bitmap for value v, the bit for a record is 1 if the record has the value v for the attribute, and is 0 otherwise
  - Used for business analysis, where rather than individual records say how much of one type exists is the query/ important

(数据行的表达方式  $\Rightarrow$  提供分析基础) male 和 L1 的表达一致  $\Rightarrow$  可能得出相关关系吗?





## Bitmap index - e.g.

要从表上提取适当的数  
据

table  
4 attributes 5 row  
record number

record number	name	gender	address	income_level
0	John	m	Perryridge	L1
1	Diana	f	Brooklyn	L2
2	Mary	f	Jonestown	L1
3	Peter	m	Brooklyn	L4
4	Kathy	f	Perryridge	L3

Bitmaps for gender

m	1 0 0 1 0
f	0 1 1 0 1

Bitmaps for income level

L1	1 0 1 0 0
L2	0 1 0 0 0
L3	0 0 0 0 1
L4	0 0 0 1 0
L5	0 0 0 0 0

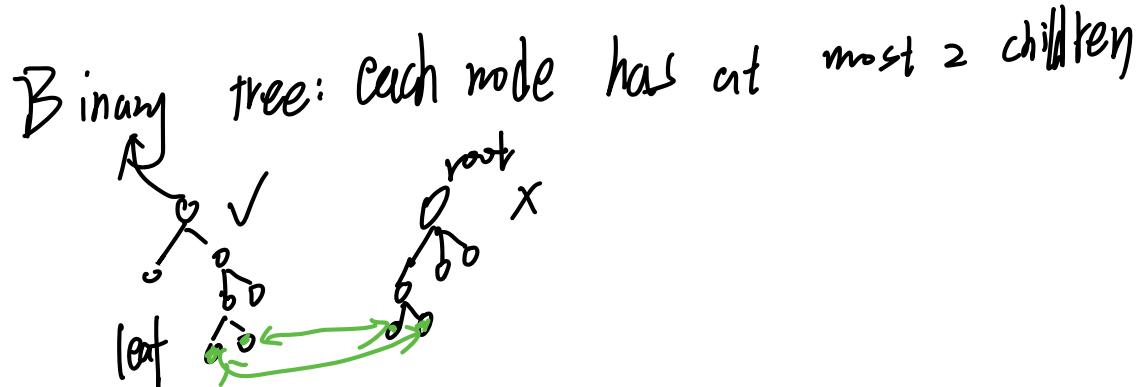
bitmap  
No = 5 (-共5位)

5位的15位bit

gender  
1 0 0 1 0  
女 女

全table没有出现0



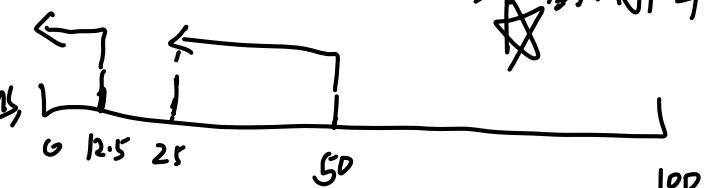


## B+tree

Why?

- Keeping files in order for fast search ultimately degrades as file grows, since many overflow blocks get created.
- So binary search on ordered files cannot be done. : insert & delete 都会破坏数据
- Periodic reorganization of entire file is required to achieve this. : reorganized, 增加以后变动的东西太多

With B+tree /  
We do not need to reorganize  
the entire file in the  
face of insertions and deletions,  
instead we just need small  
local changes.

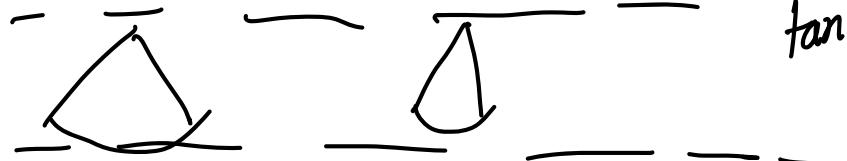


binary search 的前提条件: order  $\Rightarrow$  查找过程无法进行否则





fan out : for binary tree is small, while for B+tree,  
fan out is large.



## B+tree Definition

- It is similar to a binary tree in concept but with a fan out that is defined through a number  $n$  fanout
- All paths from root to leaf are of the same length (depth)

- ① ■ Each node that is not a root or a leaf has between  $\lceil n/2 \rceil$  and  $n$  children 中间结点有  $\lceil n/2 \rceil, n$  个 children  
 $\lceil n/2 \rceil$  和  $n$  个 children  $\lceil n/2 \rceil$  中间结点有  $\lceil n/2 \rceil, n$  个 children
- ② ■ A leaf node has between  $\lceil (n-1)/2 \rceil$  and  $n-1$  values 叶子结点有  $\lceil (n-1)/2 \rceil, n-1$  个值
- ③ ■ Special cases:  
● If the root is not a leaf, it has at least 2 children.  $\rightarrow$  根不是叶子：至少 2 孩子  
● If the root is a leaf (that is, there are no other nodes in the tree), it can have between 0 and  $(n-1)$  值。 $\rightarrow$  根是叶子：有 0 到  $n-1$  个值。





## B+tree Typical Node

- Typical node



e.g.,  
ID的值

- $K_i$  are the search-key values
- $P_i$  are pointers to children (for non-leaf nodes) or pointers to records or buckets of records (for leaf nodes)

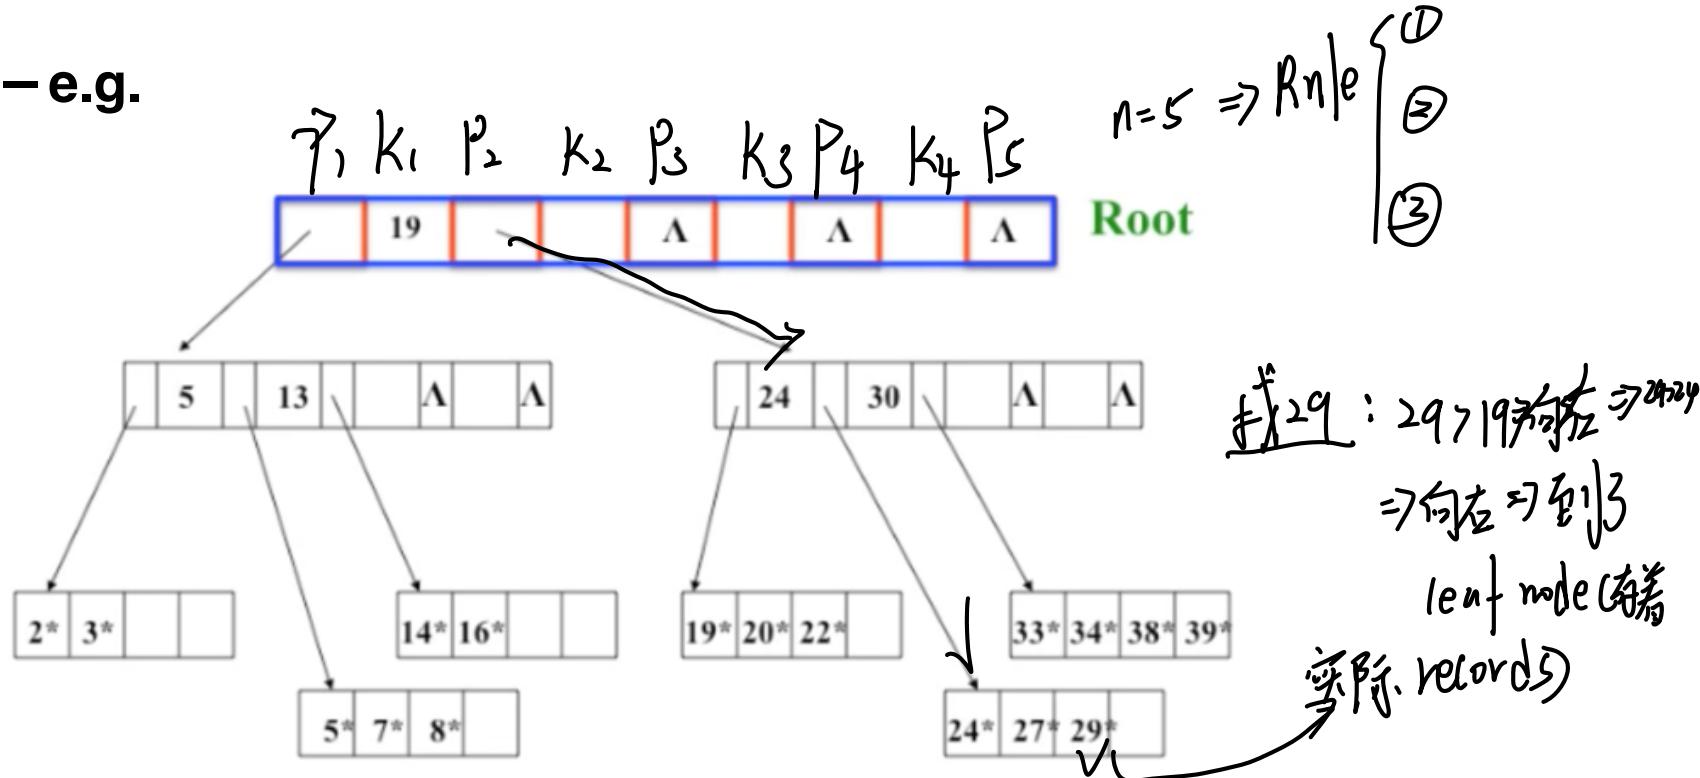
- The search-keys in a node are ordered

$$K_1 < K_2 < K_3 < \dots < K_{n-1}$$

leaf node 在 disk

NOTE: Most of the higher level nodes of a B+ tree would be in main memory already!  
(leaf node will be in disk)



**B+tree - e.g.**



## Run a query in B+tree

- Finding all records with a search-key value of  $k$ .

1.  $N = \text{root initially}$  从root开始往node上的值向右找， otherwise follow

2. Repeat

1. Examine  $N$  for the smallest search-key value  $> k$ .

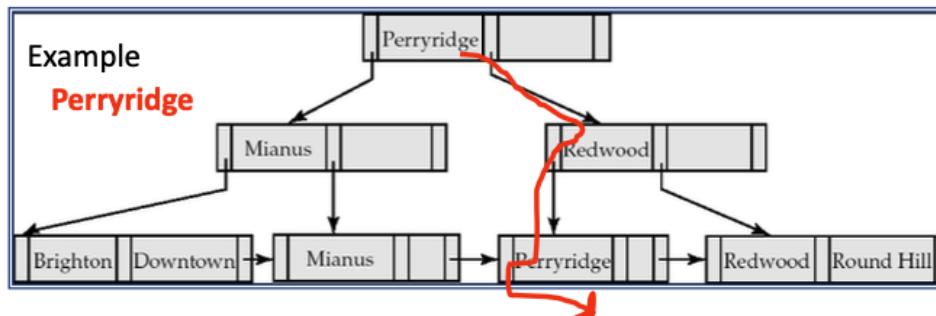
2. If such a value exists, assume it is  $K_i$ . Then set  $N = P_i$

3. Otherwise  $k \geq K_{n-1}$ . Set  $N = P_n$ . Follow pointer.

Until  $N$  is a leaf node

3. If for some  $i$ , key  $K_i = k$  follow pointer  $P_i$  to the desired record or bucket.

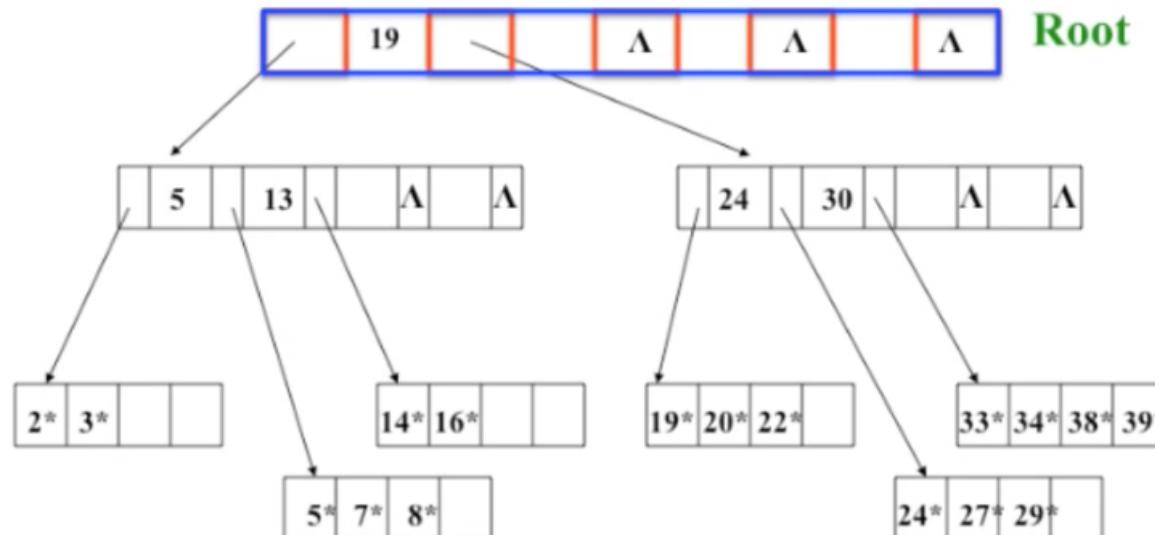
4. Else no record with search-key value  $k$  exists.





## Run a query in B+tree – e.g., find 28/29

找 28 19向右  $\rightarrow$  24 向右  $\Rightarrow$  没有  $\square$  = 被回  
not found



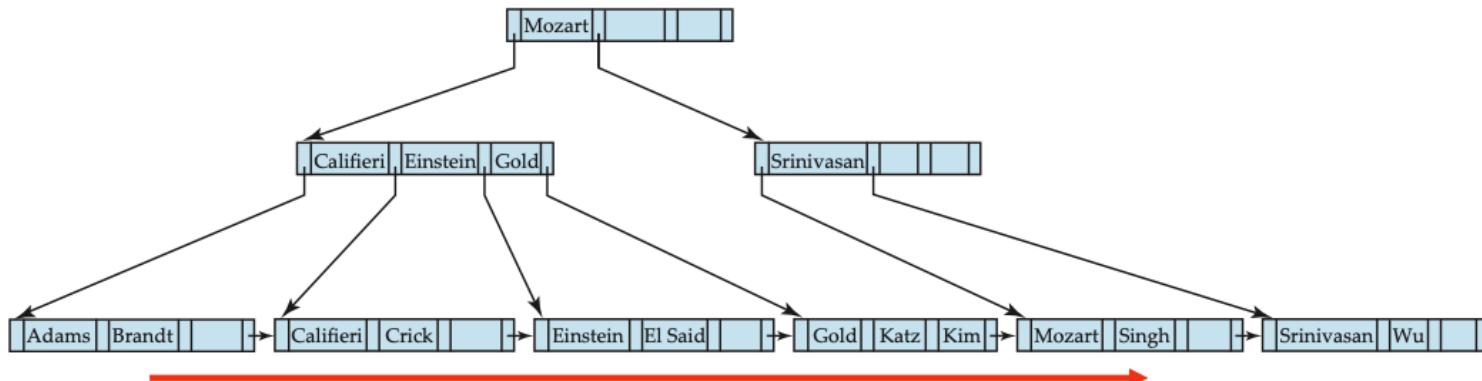


## Run range queries in B+tree (一条扫描)

无论单个range, B+tree都记得住

Range queries find all records with search key values in a given range

e.g., from Brandt to Singh  
(Run range query)



(Brandt → Singh) 扫描





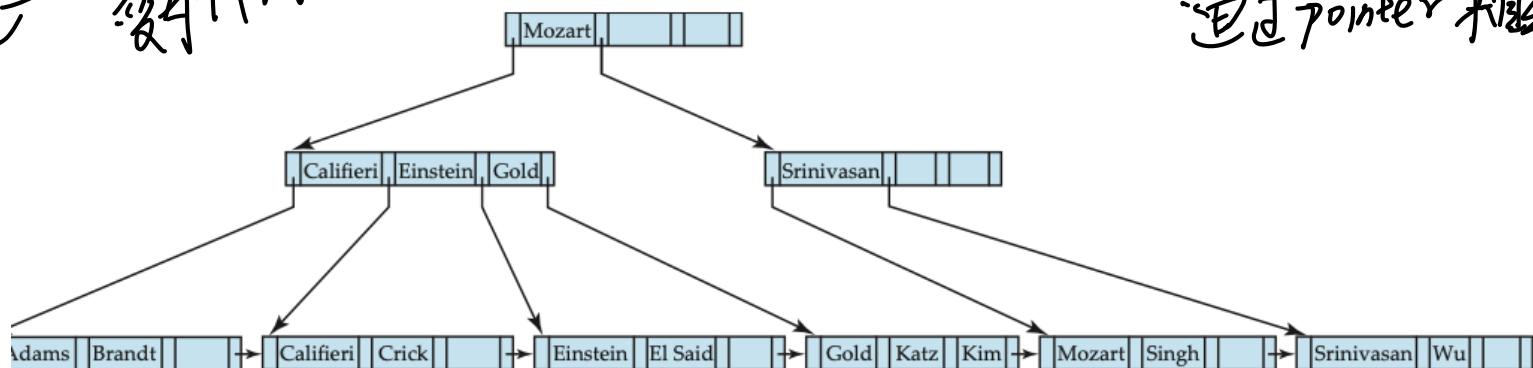
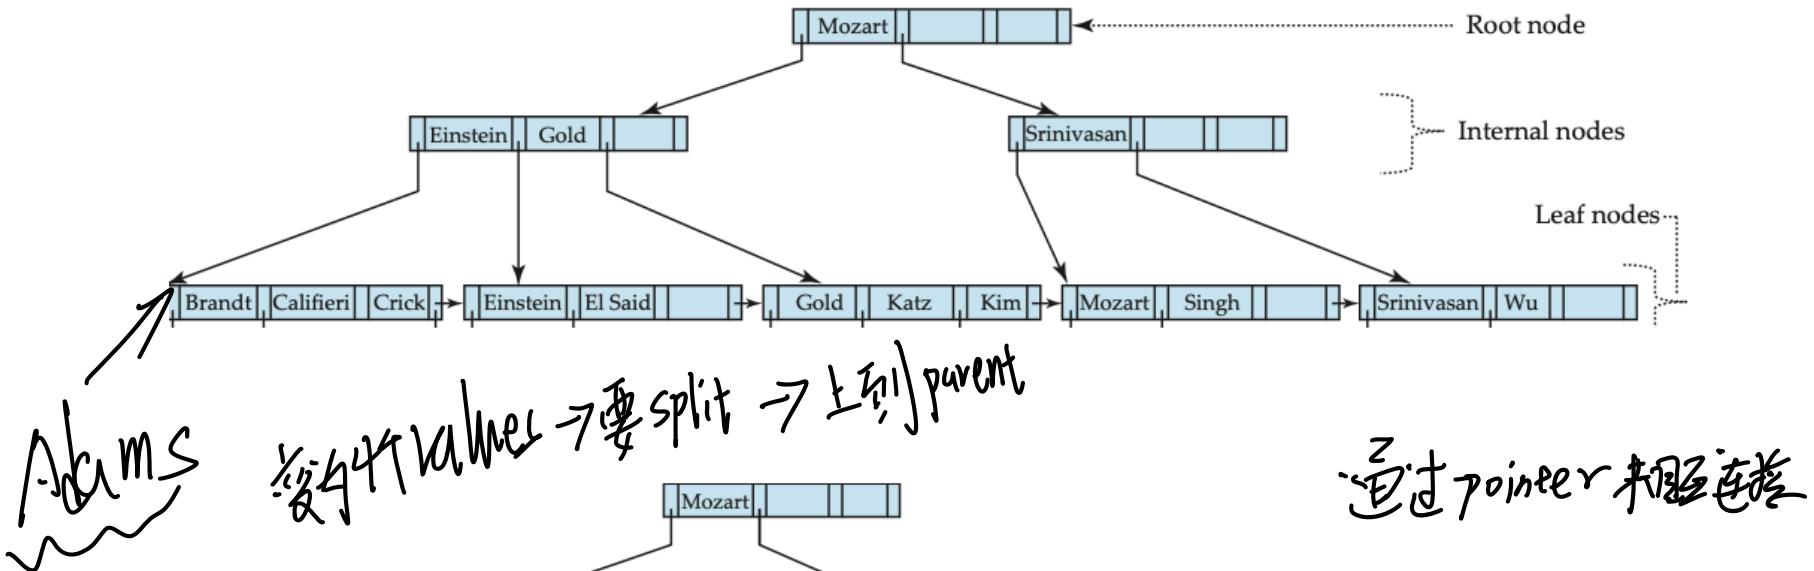
## Basic Rules for this Tree ( $n = 4$ )

- Each node that is not a root or a leaf has between  $\lceil n/2 \rceil$  and  $n$  children
- A leaf node has between  $\lceil (n-1)/2 \rceil$  and  $n-1$  values
- Special cases:
  - If the root is not a leaf, it has at least 2 children.
  - If the root is a leaf (that is, there are no other nodes in the tree), it can have between 0 and  $(n-1)$  values.





## B+tree - Insertion (Type 1 – Parent node has room)

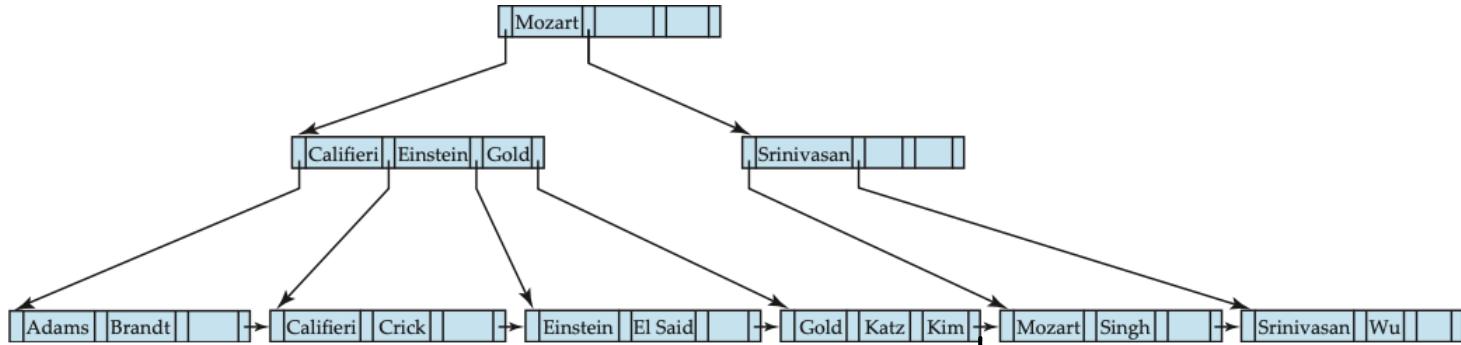


B+-Tree before and after insertion of “Adams”

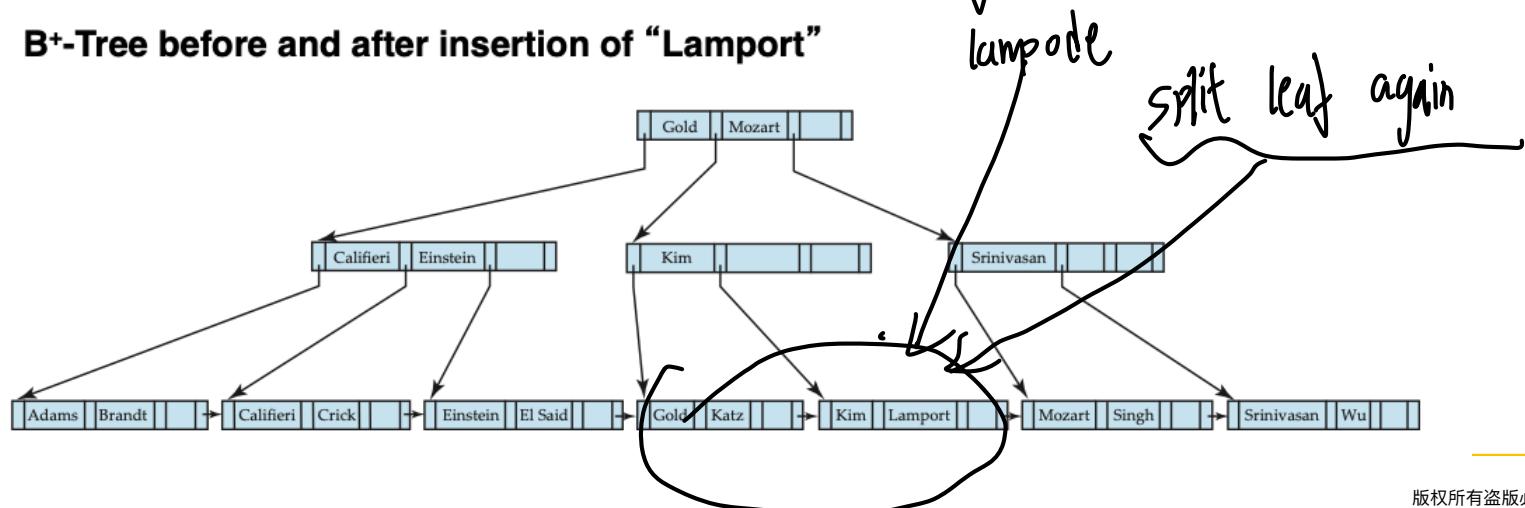




## B+tree - Insertion (Type 2 – Parent node does not have enough room)

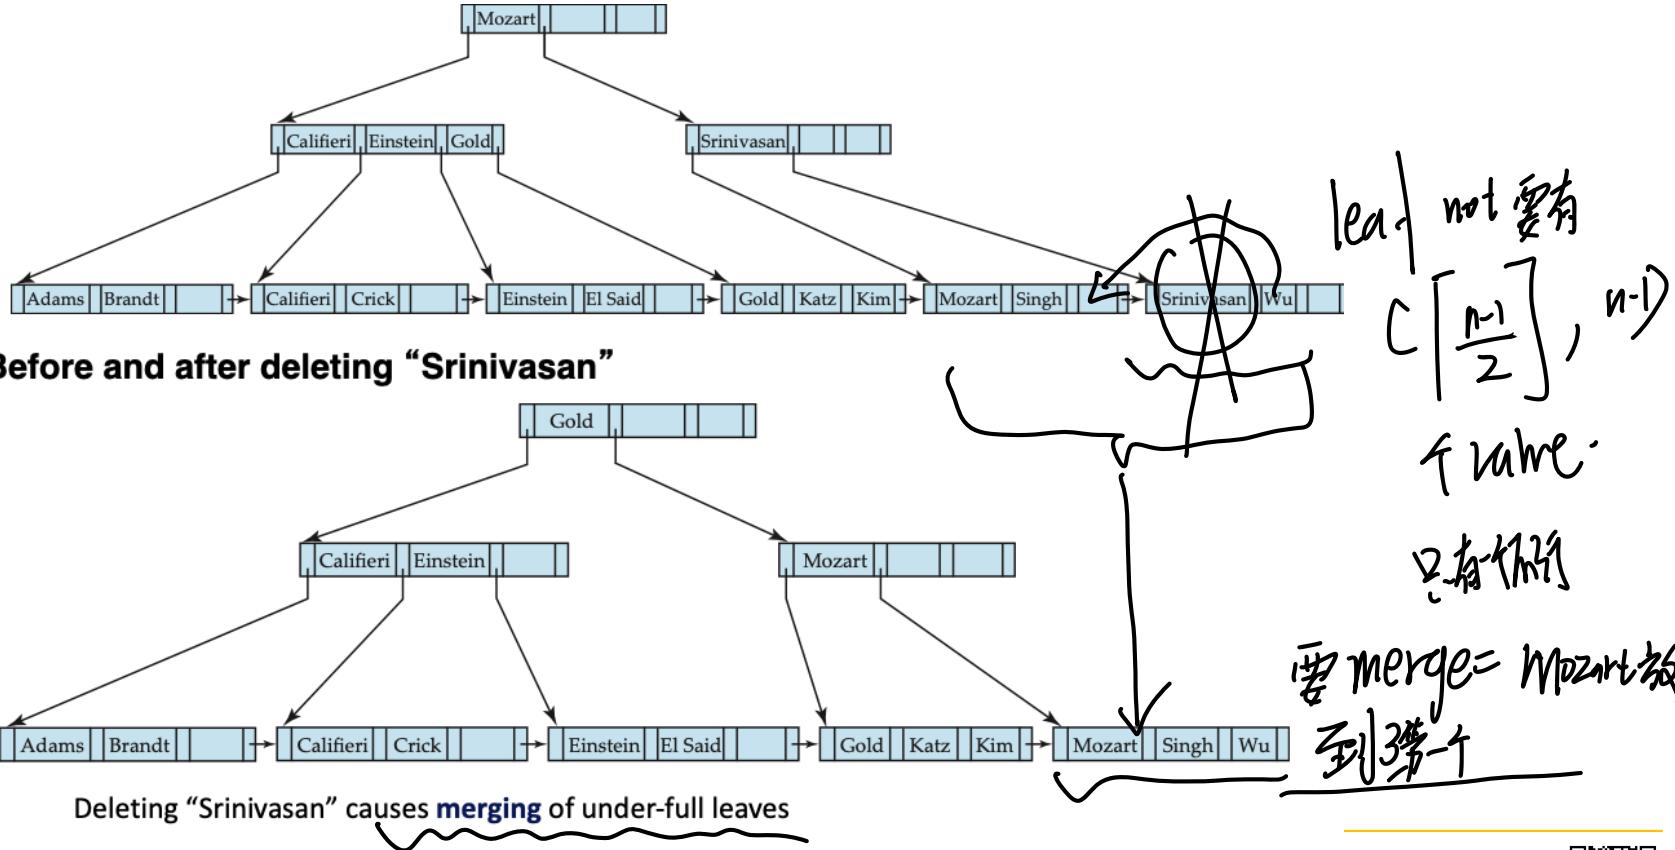


B+Tree before and after insertion of “Lamport”



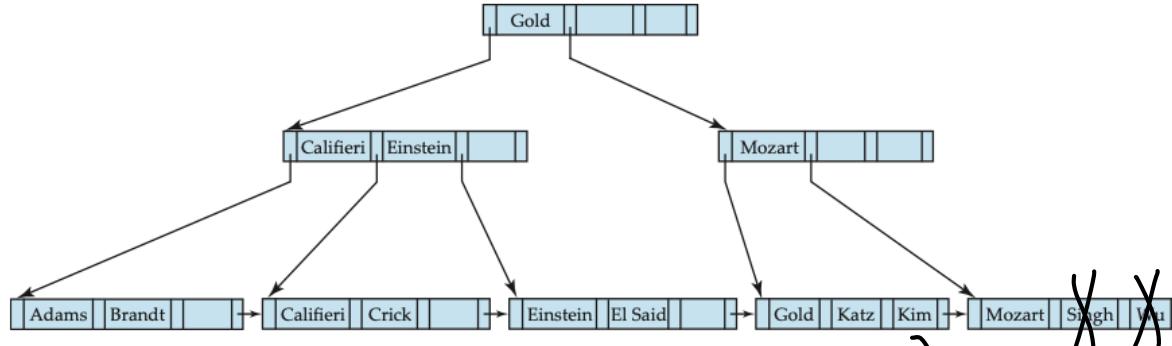


## B+tree - Deletion (Type 1 - merge)

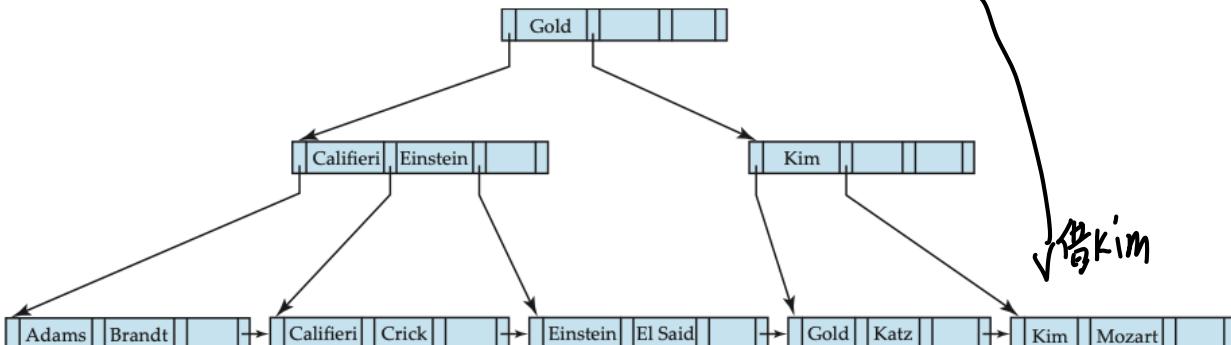




## B+tree - Deletion (Type 2 - borrow)



**Before and after deleting “Singh” and “Wu”**



**Leaf containing Singh and Wu became underfull, and borrowed a value  
Kim from its left sibling**

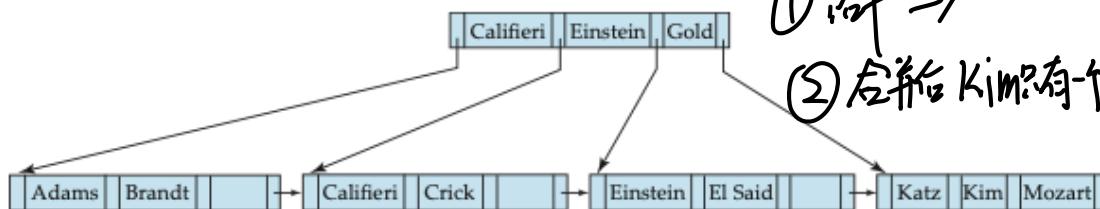
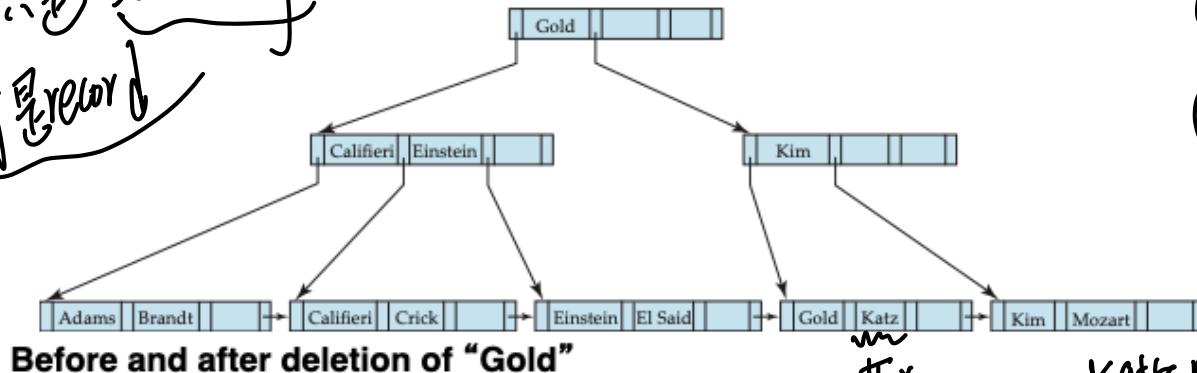
**Search-key value in the parent changes as a result**





## B+tree - Deletion (Type 3 – one level is removed)

node上是 search key  
叶上是 record



- Node with Gold and Katz became underfull, and was merged with its sibling
- Parent node becomes underfull, and is merged with its sibling
- Value separating two nodes (at the parent) is pulled down when merging
- Root node then has only one child, and is deleted

$n=4$

- ① leaf node 2-3个 value
- ② pig node 2-4个 children
- ③ root 至少有2个 children.

① 去掉 Gold => Katz merge with left  
② 合并后 Kim 只有一个 child 3 => Kim 和左边的 merge:  
{ Califieri, Einstein, Kim }  
Kim merged with left





## B+tree – File organization

- Leaf nodes in a B<sup>+</sup>-tree file organization store records, instead of pointers to children
- Helps keep data records clustered (ordered) even when there are insertions/deletions/updates
- Insertion and deletion of records are handled in the same way as insertion and deletion of entries in a B<sup>+</sup>-tree index.



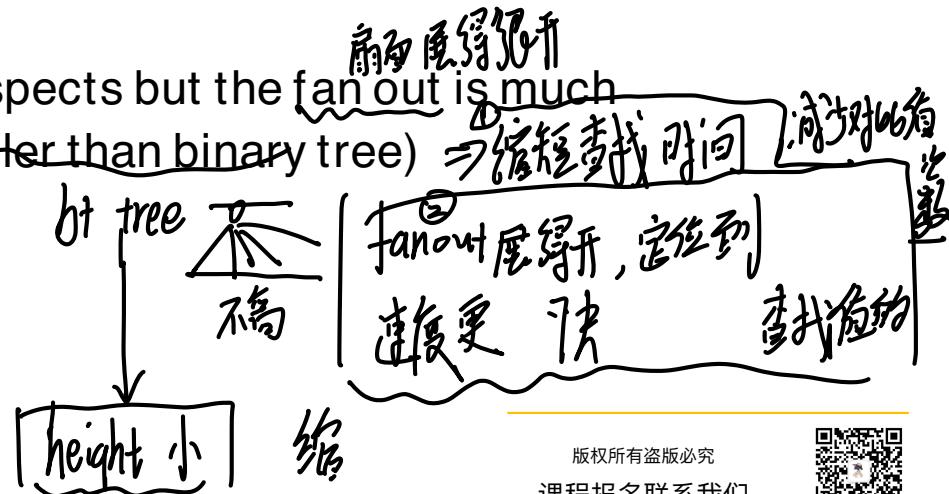


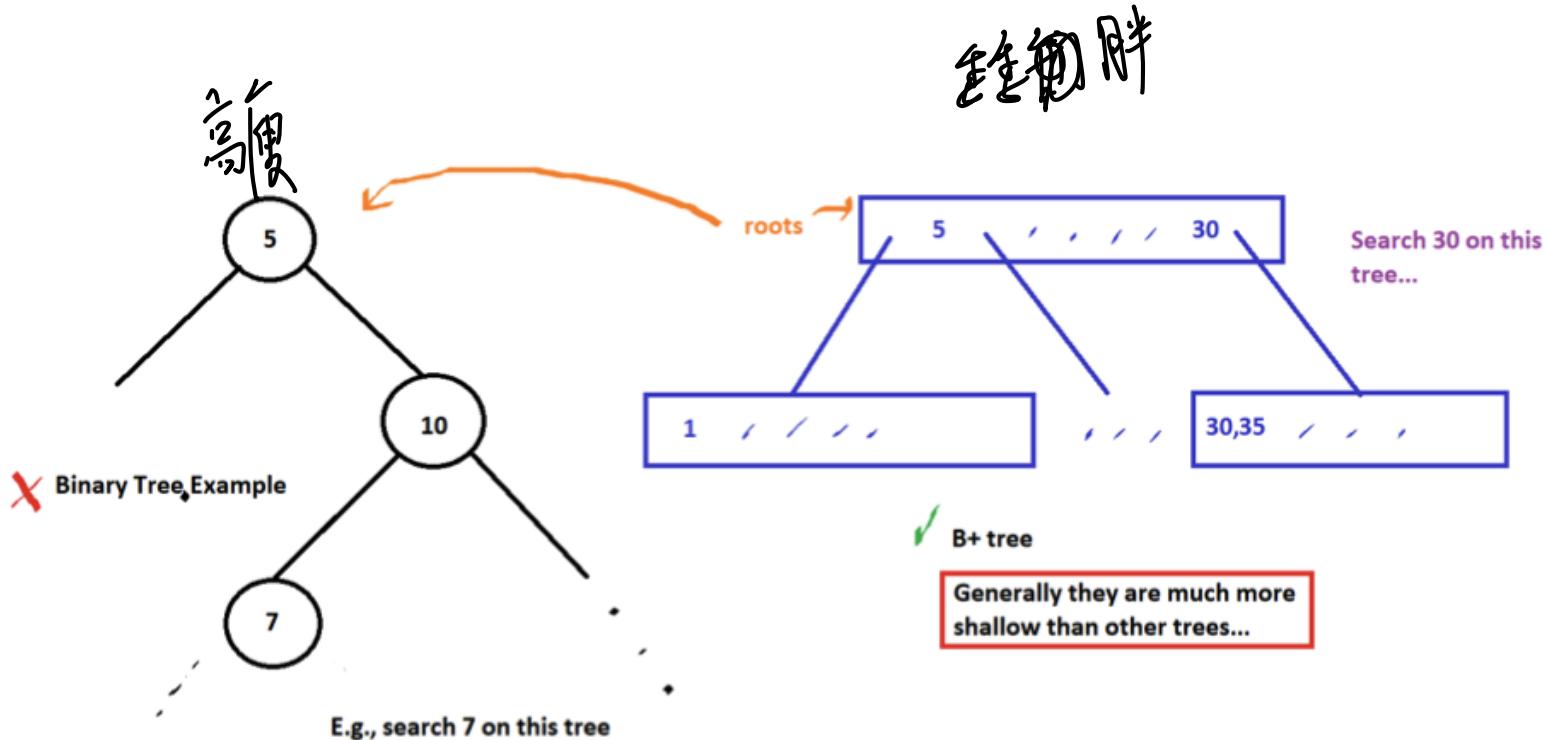
## B+tree

### Advantages

- Automatically reorganizes itself with small, local changes in the face of insertions and deletions. 无需对整个文件进行重新组织.
- Reorganization of entire file is not required to maintain performance.

※ Similar to Binary tree in many aspects but the fan out is much higher (height of B+ tree is smaller than binary tree) 减少查找时间.







## B+tree

### Disdvantages

Extra insertion and deletion over head and space over head (reorganisation)

Advantages of B+- trees outweigh disadvantages for DBMSs  
(B+ trees are used extensively)

查找速度快





多维  
计算复杂

## Considering location data (spatial data)...

Unlike things we can access by names, ids, there is a lot of data that exists, and increasingly that requires special indexing

For example, **spatial data requires more complex computations** for accessing data, e.g., intersections of objects in space

There is no trivial way to sort items which is a key issue,  
e.g., range query on a simple set of items      5km内饭馆.  
e.g., a Nearest Neighbour query      寻找最近的饭馆





index 索引      Q-tree      对数据点

## Quadtree (for location data - points)

Each **node** of a quadtree is associated with a rectangular region of space; the top node is associated with the entire target space.

Each **division** happens with respect to a rule based on data type.

Each non-leaf **nodes** divides its region into four equal sized quadrants

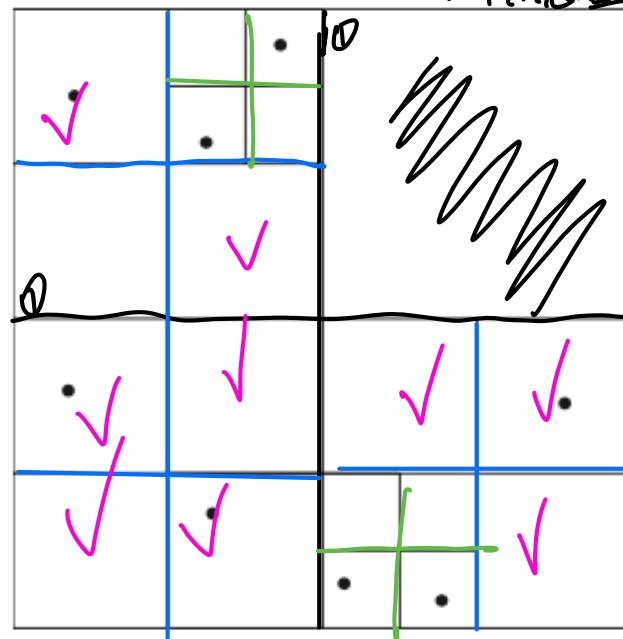
Thus each such node has four child nodes corresponding to the four quadrants and **division continues recursively until a stopping condition**





## Quadtree – e.g.,

Example: Leaf nodes have between zero and some fixed maximum number of points (set to 1 in example below)



rule

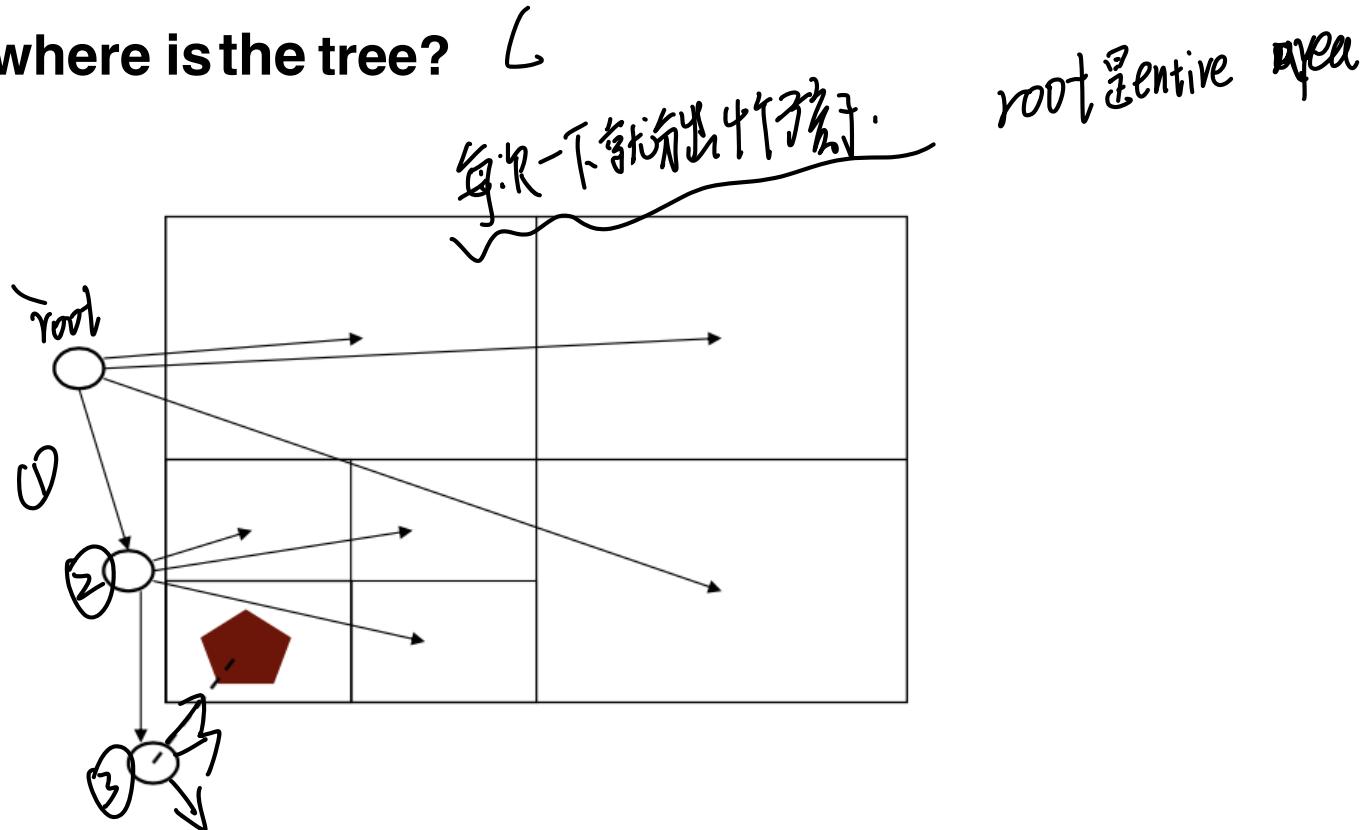
点 0~1 ✓ 不够

如果都符合规则 再切  
✓ 是符合的  
剩下的再切  
3. 之前符合



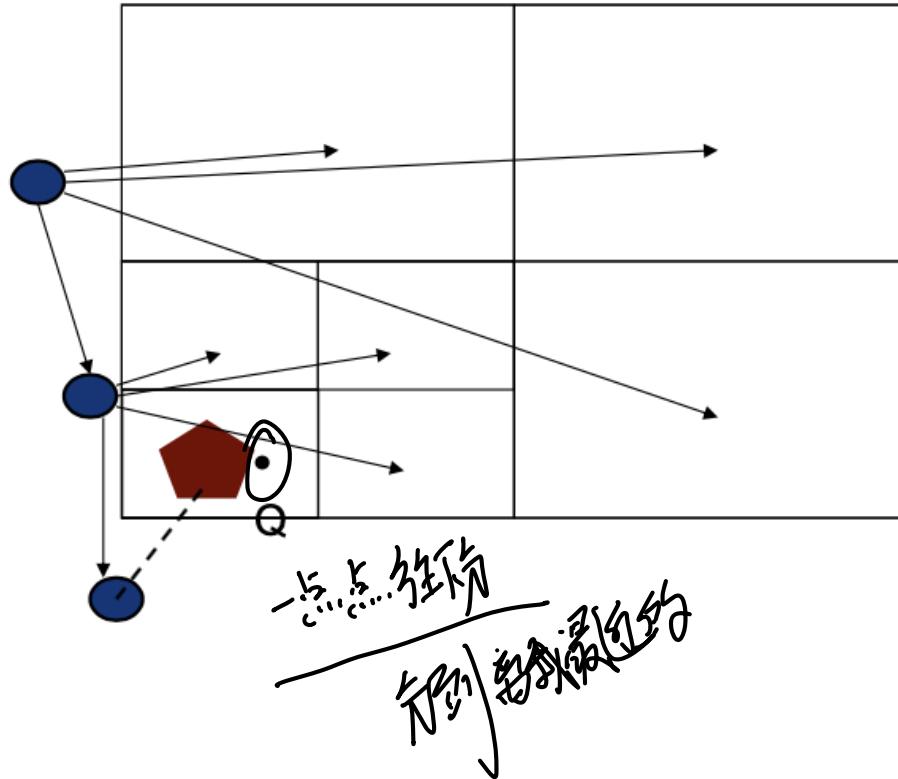


## Quadtree – where is the tree?





## Quadtree – e.g., run NN query



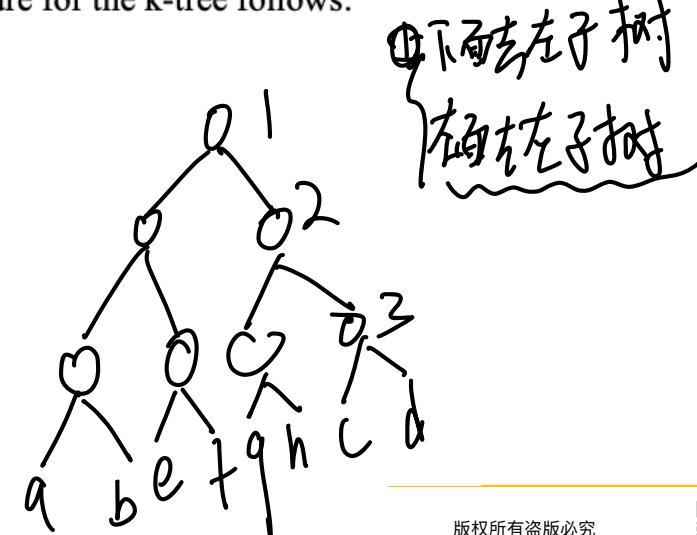
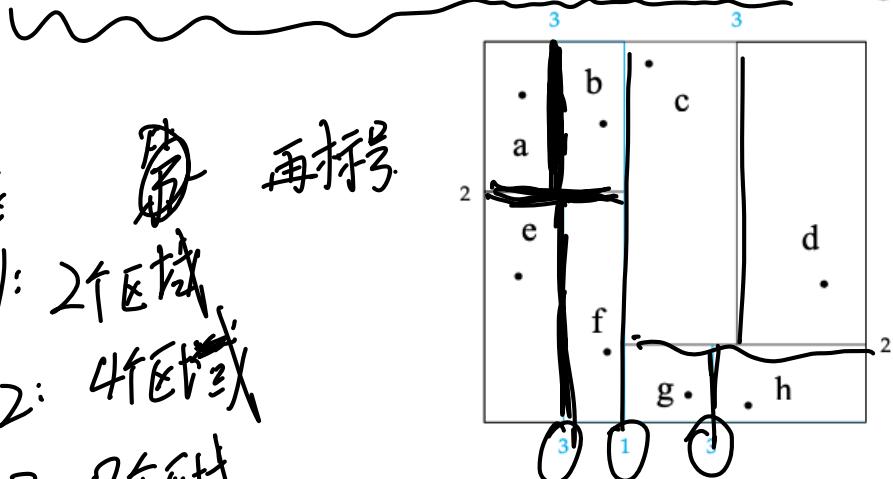


## 2022 S1 (k-d tree)

### Question 5: [4 Marks]

Given the k-d tree below with point data, where black labeled dots represent spatial coordinate data, and the rectangular area is divided into regions with the division order given with numbers: draw the associated k-d tree as a tree structure with leaves labeled as the data labels given below. Assume left subarea of a division goes to a left subtree, and lower subarea of a division also goes to a left subtree. The figure for the k-tree follows:

分区  
1: 2个区域  
2: 4个区域  
3: 8个区域  
再标注





## R-tree (for location data – rectangles & polygons)

**R-trees** are an N-dimensional extension of B<sup>+</sup>-trees, useful for indexing sets of rectangles and other polygons.

Supported in many modern database systems, along with variants like R<sup>+</sup>-trees and R\*-trees.

**Basic idea:** generalize the notion of a one-dimensional interval associated with each B+ -tree node to an N-dimensional interval, that is, an N-dimensional rectangle.

Will consider only the two-dimensional case ( $N = 2$ )

- generalization for  $N > 2$  is straightforward, although R-trees work well only for relatively small N

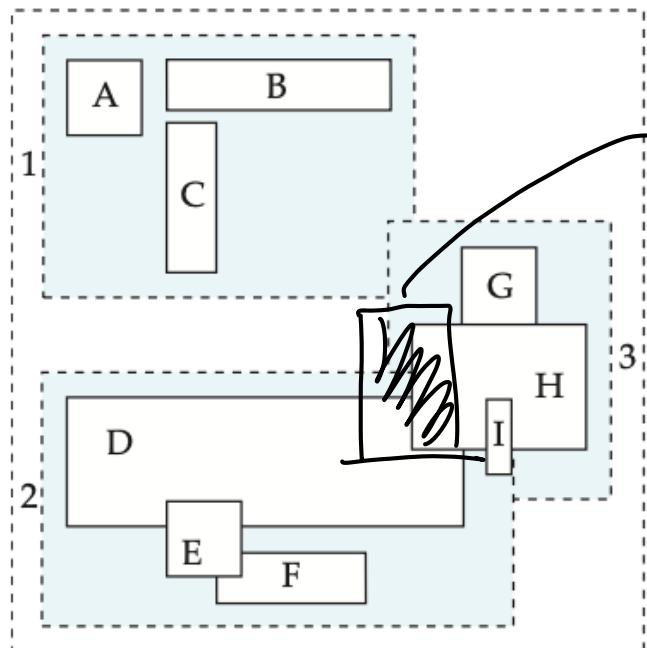
*Bounding boxes of children of a node are allowed to overlap*

Note: A **bounding box** of a node is a minimum sized rectangle that contains all the rectangles/polygons associated with the node

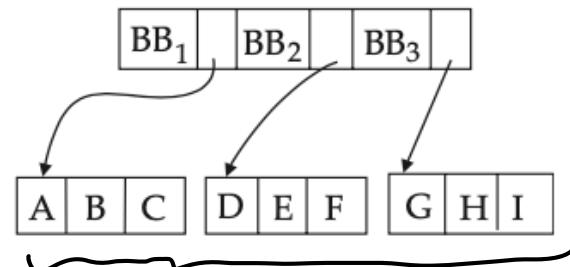




## R-tree – e.g.,



→ 16 号的 query 与 3 bounding box 2 和 bounding box 3  
和 BB 的 D、BB 的 H 相同  $\Rightarrow$  return R.H  
∴ 3 个 ~~黑色~~ 区



leaf node

leaf node  
B+树节点：point to actual record.





## Search in R-tree

To find data items intersecting a given query point/region, do the following, starting from the root node:

- If the node is a leaf node, output the data items whose keys intersect the given query point/region.
- Else, for each child of the current node whose bounding box intersects the query point/region, recursively search the child.

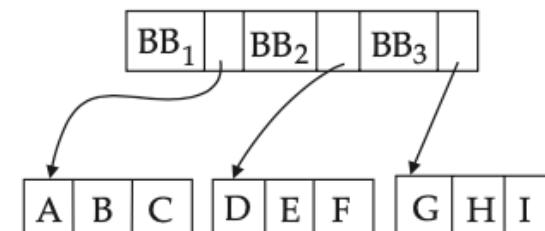
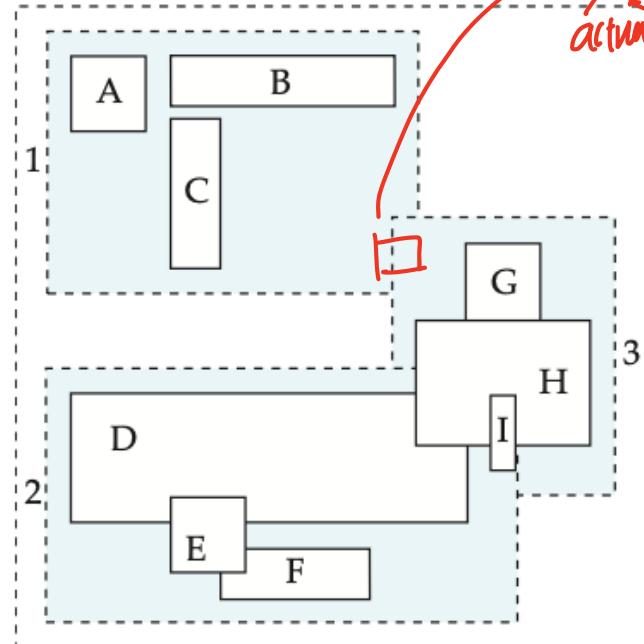
Can be very inefficient in worst case since multiple paths may need to be searched due to overlaps, but works acceptably in practice.

There might be multiple overlapping boundary boxes  $\Rightarrow$  which may lead to explore multiple parts or like going through branches of this tree to find the actual result. 举例看下面:





## Search in R-tree – e.g.,

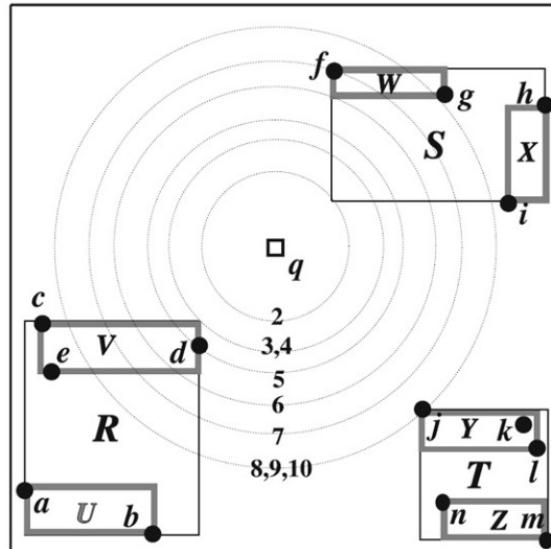
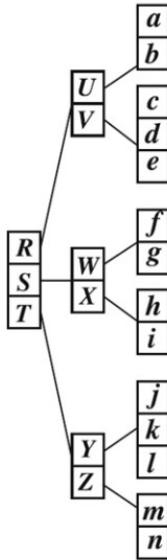


是否 intersect? → 是 OK return  
否跳过，下一个 →





## Continued R-tree



Step	Priority queue	Retrieved NN(s)
1	$\langle S, R, T \rangle$	None
2	$\langle R, W, T, X \rangle$	None
3	$\langle V, W, T, X, U \rangle$	None
4	$\langle d, W, T, X, c, e, U \rangle$	None
		return d
Step 5 finds d as the first NN (using Best First search)		

R-tree 存在问题

举例说明





## 2 – Using Indexes in SQL





## Index Definition in SQL

- Some indexes are automatically created by the DBMS  
*(unique → non-clustered index)* For UNIQUE constraint, DBMS creates a non- clustered index.
  - For PRIMARY KEY, DBMS creates a clustered index  
*选 -> Table →*
  - You can create indexes on any relation (or view) – view: e.g., join result of 2 tables  
*Table* *合并结果*





## Create an Index

```
create index <index-name> on <relation-name>  
(<attribute-list>)
```

e.g.,

1. Create a clustered index on a table

```
CREATE CLUSTERED INDEX index1 ON table1 (column1);
```

2. Create a non-clustered index with a unique constraints

```
CREATE UNIQUE INDEX index1 ON table1 (column1 DESC,  
column2 ASC, column3 DESC);
```

(A unique index is one in which no two rows are permitted to have  
the same index key value)





## Drop an Index

```
drop index <index-name>
```

Note: Most database systems allow specification of type of index, and clustering





## Specialized Index – filtered index

```
CREATE INDEX index1 ON table1 (column1)
```

```
WHERE Year > '2010';
```

(filtered index)

(A filtered index is an optimized nonclustered index, suited for queries that select a small percentage of rows from a table. It uses a filter predicate to index a portion of the data in the table)





## Specialized Index – spatial index

KD tree  
R tree → boundary box

```
CREATE SPATIAL INDEX index_name ON
table_name(Geometry_type_col_name) WITH ( BOUNDING_BOX = ( 0,
0, 500, 200 ) );
```





## Things we need to know

- There is no point going through all index types for all data types
  - There are hundreds of them
  - Even each type would have many subtypes
    - ▶ E.g., MX-CIF quadtree is one quadtree type among many
  - Same with R-trees, etc
  - Same with many other index types





## Things we need to know

- Given a data set, when uploading to the DBMS
  - Find the potential query types → frequently (频繁) (数据类型)
  - Research what indices that particular DBMS would have for that data type
  - Research for what queries you would better do on what index
  - Create index if you have large data
  - Monitor performance
  - Tune or create other indices
  - Your DMBS will have a version of the “create index” SQL statement





# Quiz





## 1. Which one of the following statements is incorrect?

- A. Hash index can be used for fast access to individual records with id type attribute.
- B. Bitmap index is not commonly used in data analysis.
- C. Quadtree index is suitable for nearest neighbour queries on multi-dimensional data.
- D. A B+tree index is a balanced index

*B+tree is balanced*  
*the length of the paths from root node to all leaf nodes are the same*





2. A restaurant database stores the name, cuisine type, GPS location, phone number as well as average expert rating of restaurants in a table. Customers can access the database via their web browsers and the database access website contains a map for querying as well. The main query type that is issued is a NN query by customers where given a point location on the map, the system shows the nearest restaurant. Which one of the following is the best index to create on this table for fast access?

- A R-tree index on GPS location
- B Bitmap index on cuisine type
- C Hash index on phone numbers
- D Hash index on GPS location





### 3. Which of the following statements is true for transaction processing?

- A. Running transactions concurrently but not in isolation with each other is not desirable.
- B. Transactions are expected to be run obeying Chemical properties such as properties of acids and metals.
- C. Durability is the least important ACID property.
- D. SQL queries are the main atomic units of execution in relational DBMSs





4. If I have a hash index given as follows, what should we expect to happen during query processing using this index for accessing the associated table?

The hash index is created on the "id" attribute of a table containing info about more than **a billion** social network users. Each "id" value is unique and generated per user at the time of creation of their social network account. This index is used to access the info associated with a user, i.e., given their id, e.g. by the system administrators.

For example, administrators access a user's info during a service phone call from the customer taking the id value and entering to the system to retrieve the associated customer row from the customer info table.

The DBMS uses a simple hash function which converts this big integer "id" number which has 20 decimal digits to binary form and takes the middle 4 bits out of the binary representation and uses it as the bucket number to access a bucket in the associated hash table.

In this case we expect that:

- A. The hash function is simple but elegant and thus saves time in data access and cannot be better.  
*无法处理 (无法) (查得慢) (效率太低)*
- B. As this is a very simple hash function it will lead to slower than expected data access.
- C. The hash index finds the data location on the disk as quickly as it could be as there is an even distribution of ids to hash table buckets.
- D. None of the above is relevant as we do not know which four bits are read per "id" instance at this time, i.e., 0000 or 1111 or etc.