



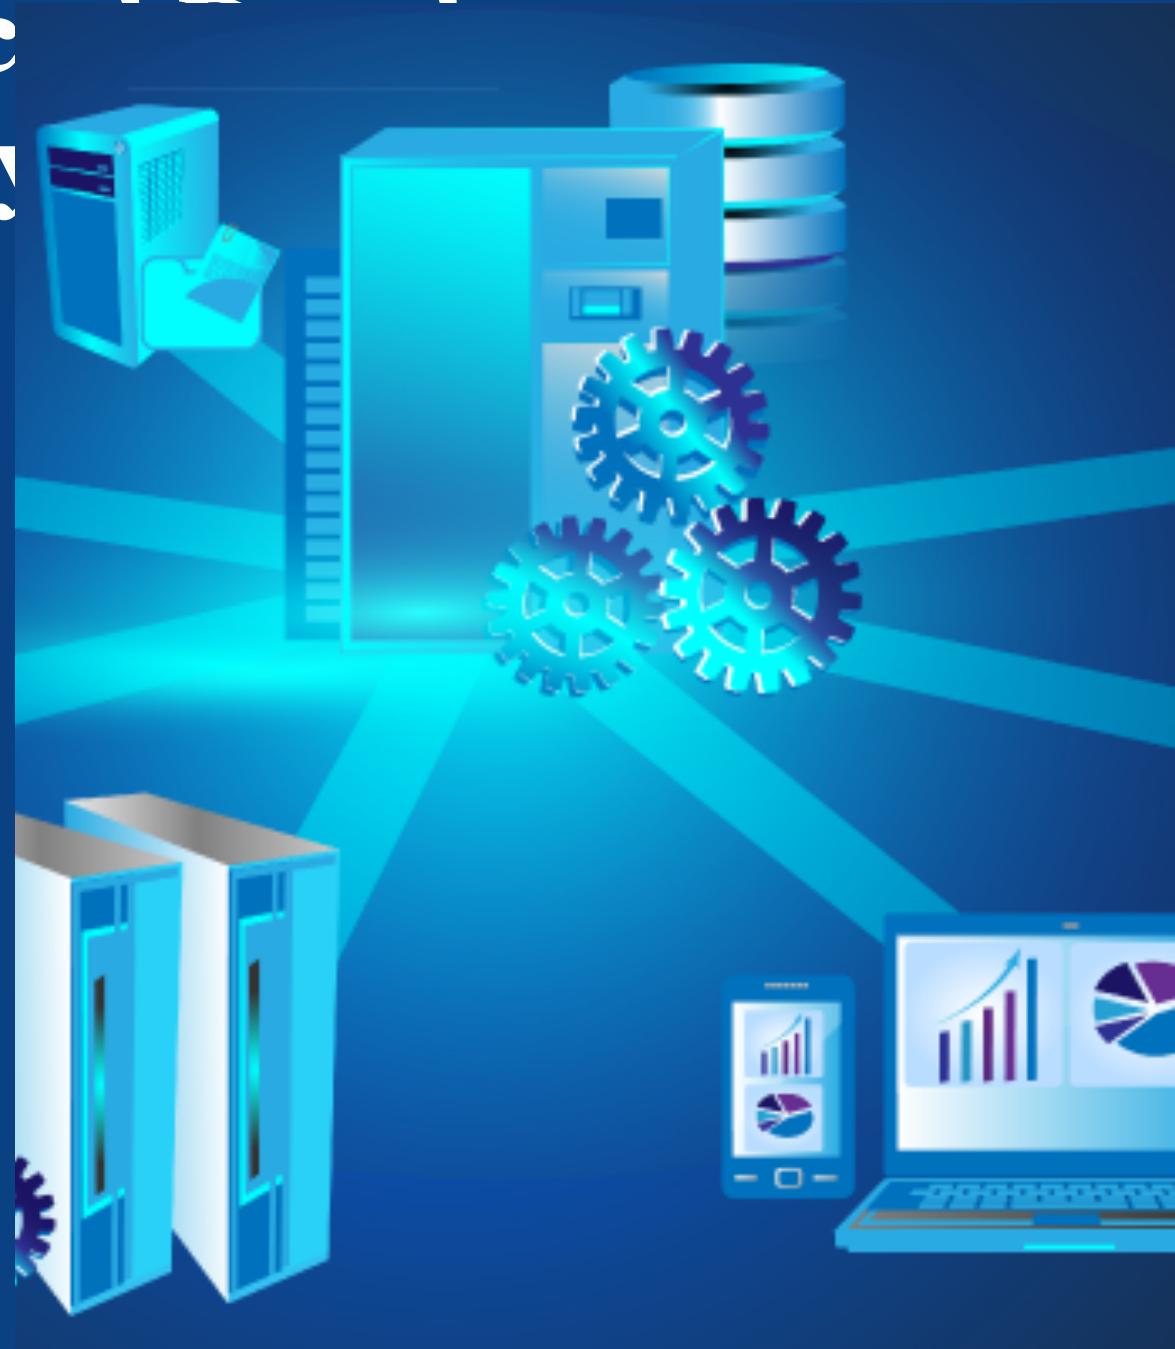
# COMP90050:

## Advanced Systems

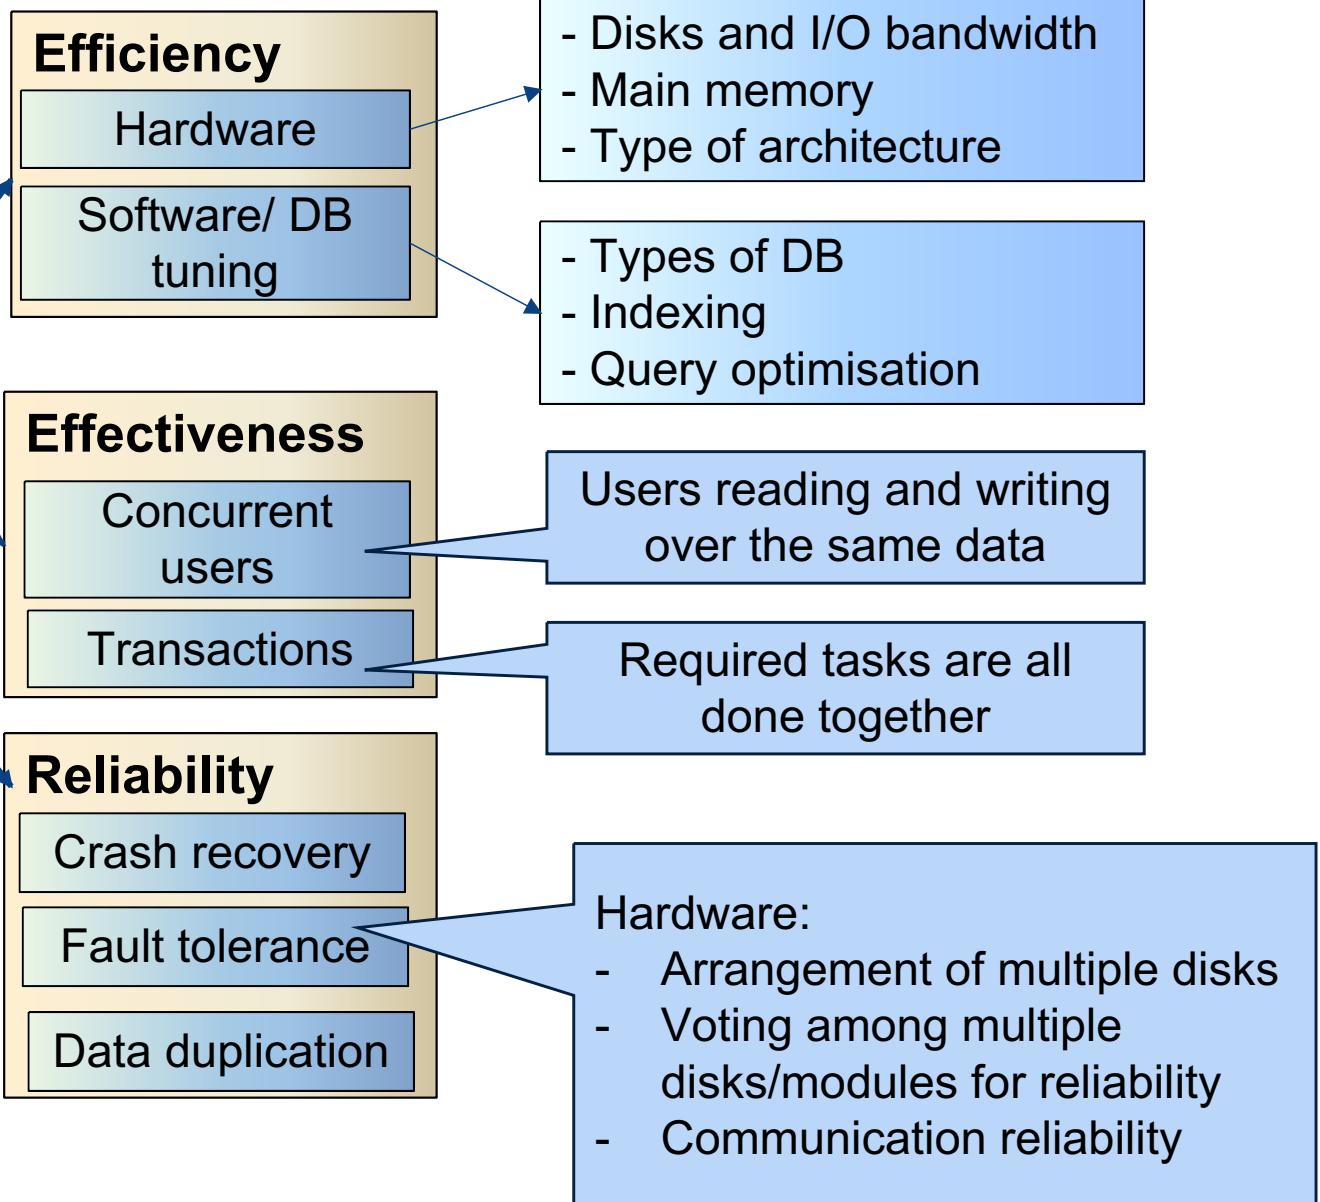
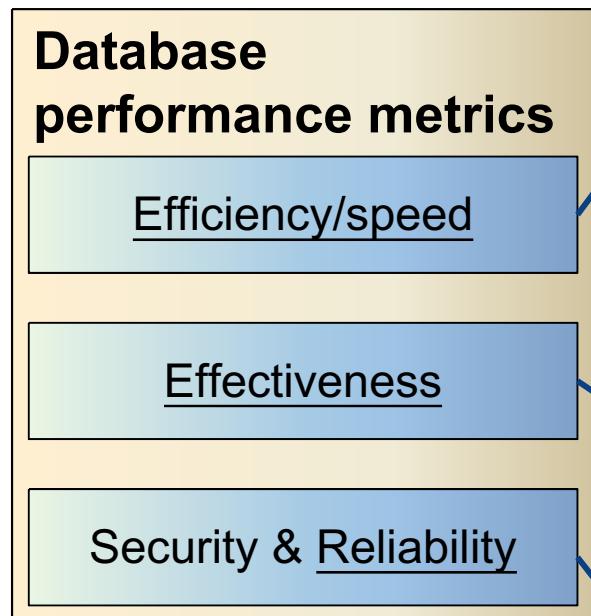
**Lecturer: Farhana Choudhury (PhD)**

**Fault tolerance**

**Week 8**



# Core Concepts of Database management system



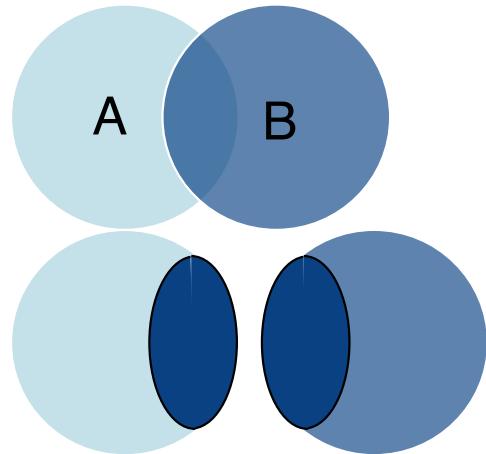


# Fault Tolerance

The property that enables a system to continue operating properly in the event of the failure of some of its components.



# Some statistics first!



#  $P(A)$  = probability of an event A is happening in a **certain period**.

#  $P(A \text{ and } B)$  = probability both A and B happening in that period

=  $P(A) * P(B)$  assuming A and B are independent events.

#  $P(A \text{ or } B) = P(A) + P(B) - P(A \text{ and } B)$

=  $P(A) + P(B) - \underbrace{P(A) * P(B)}_{\text{Assuming A and B are independent}}$

$P(A) + P(B)$  [if  $P(A)$  and  $P(B)$  are very small]

# Mean time to event A is,  $MT(A) = 1/P(A)$

(The time elapsed before A occurs)

# If events A, B, have mean time  $MT(A)$ ,  $MT(B)$ , then the mean time

to the first event,  $MT(A \text{ or } B) = 1/P(A \text{ or } B)$

The time that elapsed before one of these events occurs



If there are  $n$  events, each with the same probability  $p$ , then

Probability that one of the events occur =  $p + p + \dots$  ( $n$  times)

$$= n * p \text{ [assuming } p \text{ is small]}$$

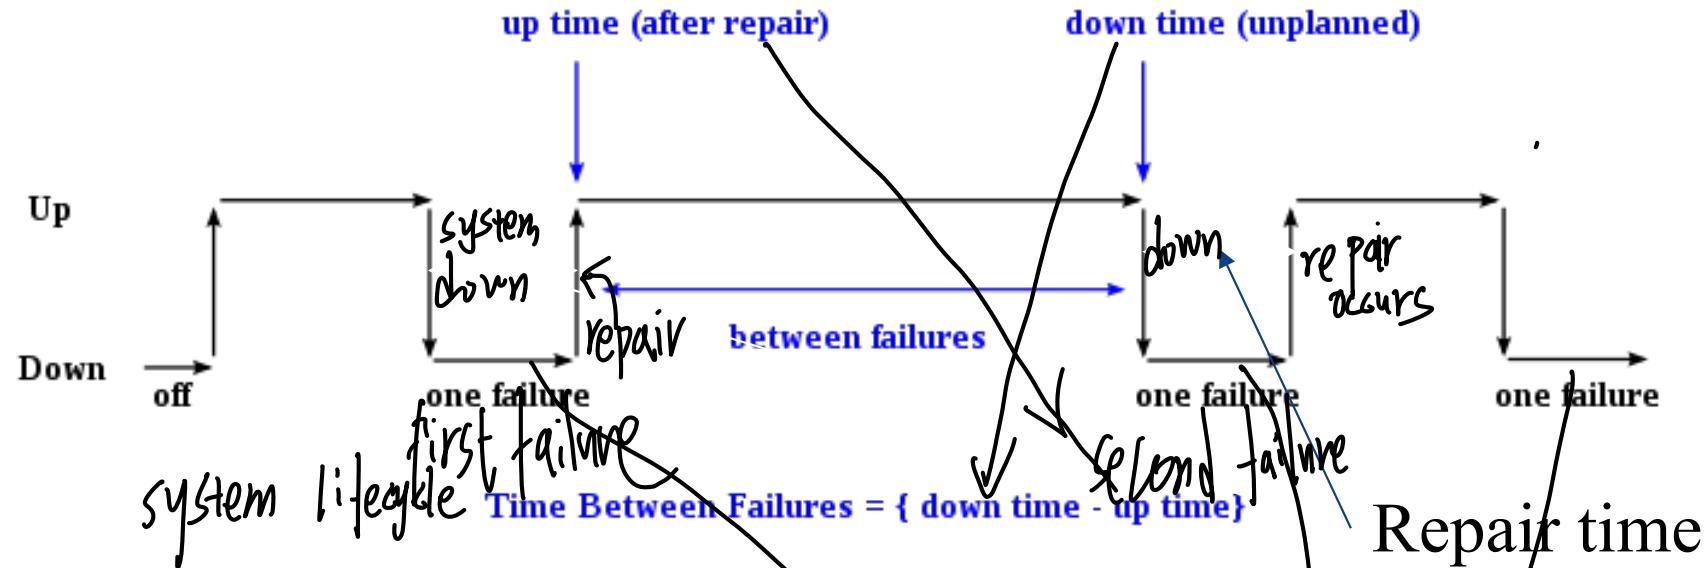
2. Mean time to one of the events (i.e., mean time to the first event)

$$= 1/(n * p)$$

$$= (1/p) * (1/n)$$

$$= m * (1/n) = m/n$$

# A system's lifecycle



Module availability : measures the ratio of service accomplishment to elapsed time

the **time** elapsing before a failure is experienced

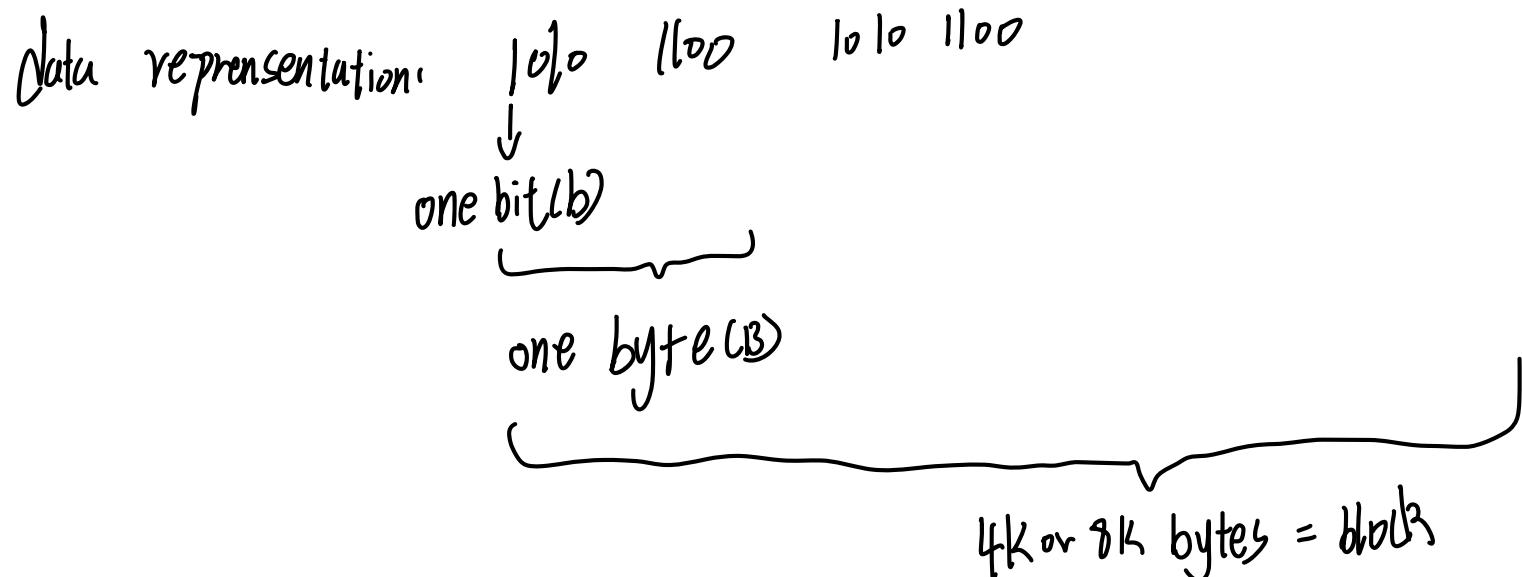
$$= \frac{\text{Mean time to failure}}{\text{Mean time to failure} + \text{mean time to repair}}$$



# Fault tolerance by RAID

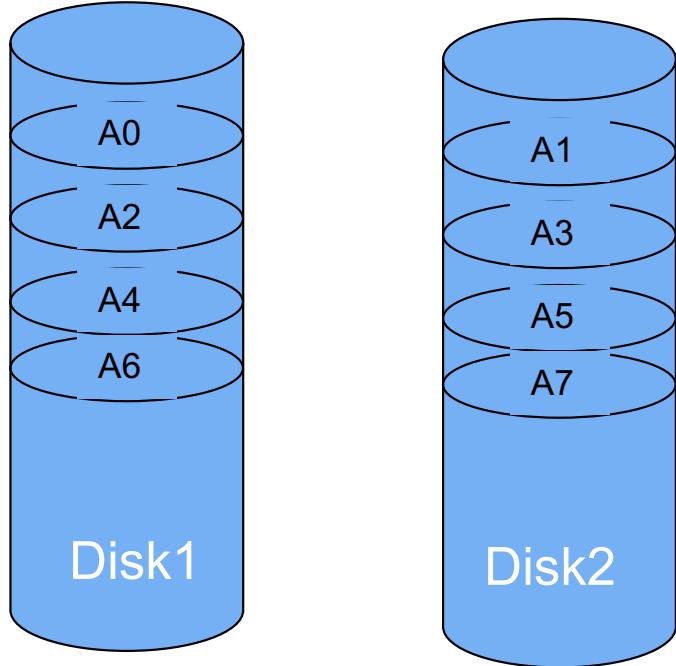
RAID: Data will be stored in more than one disk even one of the disks fail, the system will be still able to operate, and reading and writing operation can still be done there ↪

Redundant Array of Independent Disks – different ways to combine multiple disks as a unit for fault tolerance or performance improvement, or both of a database system



都是 fault tolerance 的方法

# RAID 0 (Block level Striping)



$\rightarrow$  1 block

A0, A1, A2, ... are contiguous blocks of data of a file

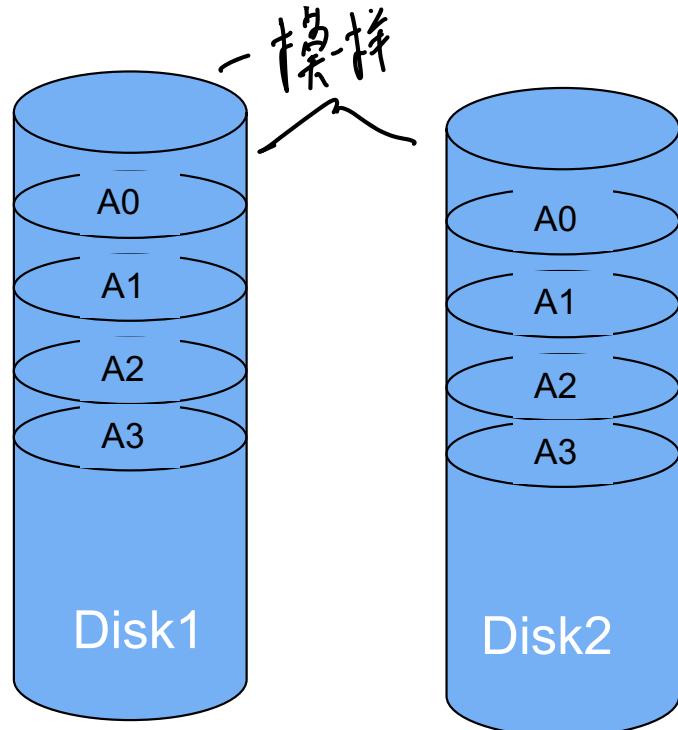
- Provides balanced I/O of disk drives – throughput ~doubles (disks can send data parallelly)
- Any disk failure will be catastrophic and MTTF reduces by a factor of 2

优点: Higher throughput at the cost of 缺点: increased vulnerability to failures

$\rightarrow$  if 1 fail, we cannot get full file

- A means Block (4K or 8K bytes of storage)
- MTTF = Mean Time To Failure

# RAID 1 (mirroring)



2gb info  
only contain 1gb  
real info

most commonly used

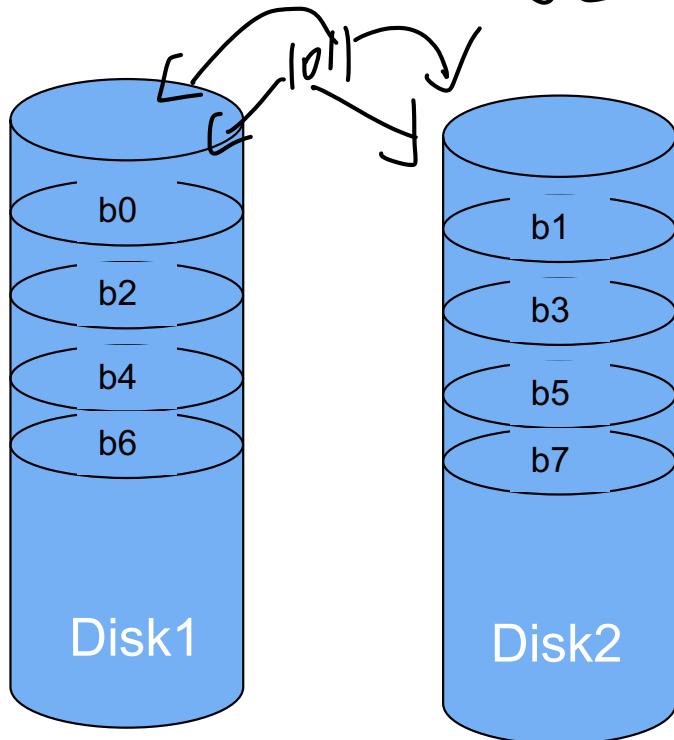
- Provides higher read throughput but lower write throughput (~~half of the total speed~~ — i.e. single disk speed)
  - Half storage utilization.**
  - MTTF increases substantially (quadratic improvement — i.e.  $MTTF^2$ !)
- reading can be done in parallel whenever anything needed to be written into Disk1, it also needs to be written into Disk2.

Continues to operate as long as 1 disk is functional

Calculation of MTTF values – in tutorials

- A means Block (4K or 8K bytes of storage)
- MTTF = Mean Time To Failure

## RAID 2 (bit level striping)



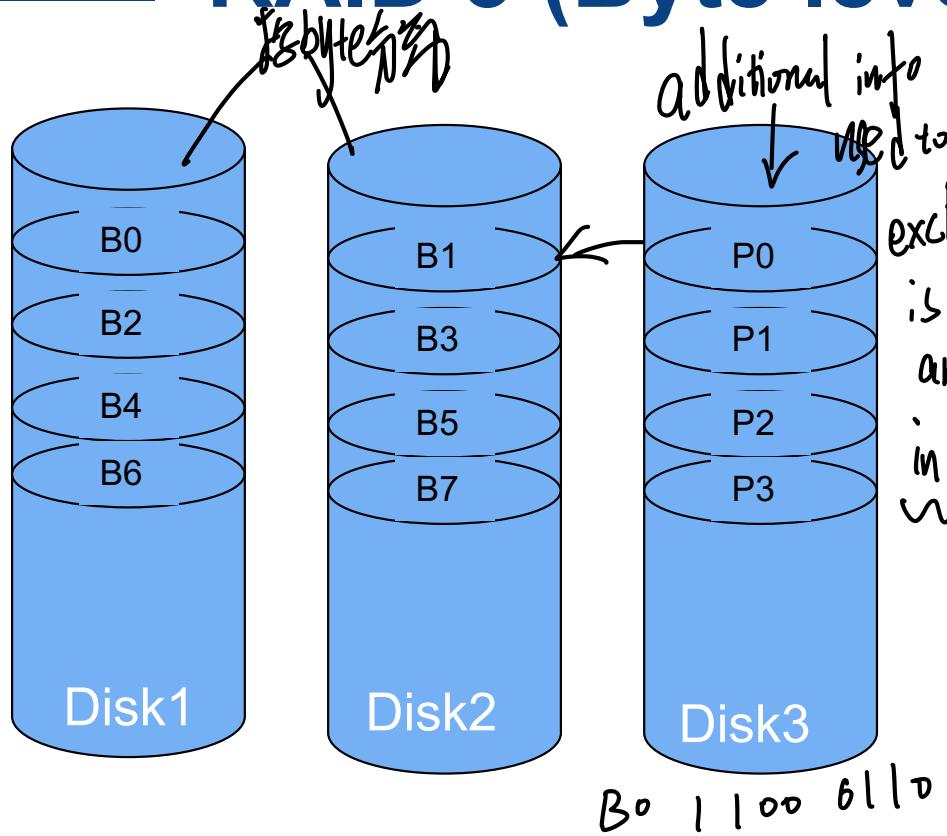
- Striping takes place at bit level
- Provides higher transfer rate (double the single disk)
- MTTF reduced by half as in RAID 0

~~• rarely used~~

*(striping at bit level) is difficult to maintain and read.*

- b means bit
- MTTF = Mean Time To Failure

# RAID 3 (Byte level striping)



- B means Byte
- P is parity
- MTTF = Mean Time To Failure
- Parity (or check bits) are used for error detection

$B_0 \ 1 \ 1 \ 0 \ 0 \ 0 \ 1 \ 1 \ 0$   
 $B_1 \ 0 \ 1 \ 1 \ 1 \ 1 \ 0 \ 0 \ 0$   
 parity  $1 \ 0 \ 1 \ 1 \ 1 \ 0 \ 1 \ 0$

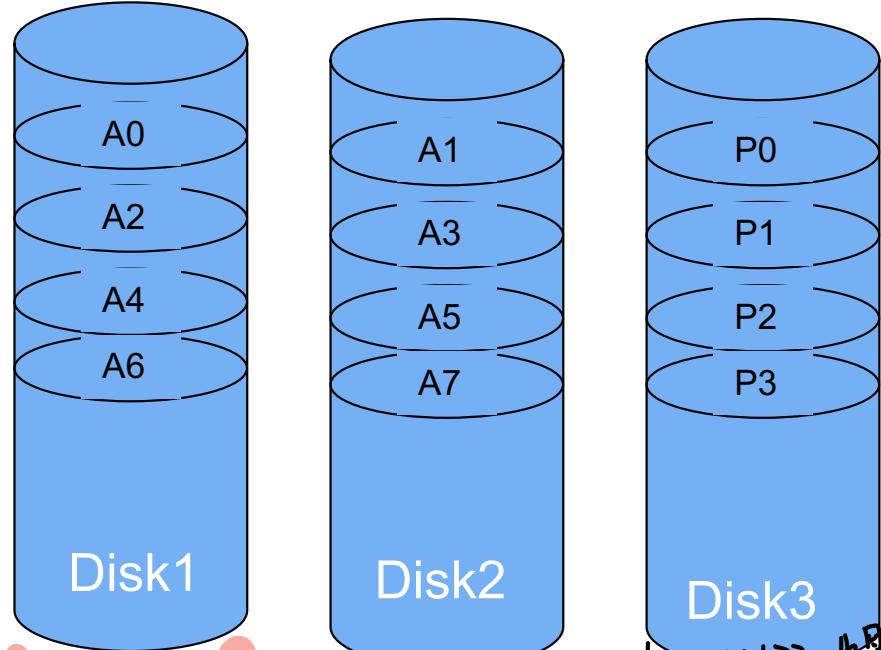
- additional info combining disk1 and disk2 and storing this new info at phis Disk3  
 B0, B1, B2, B3, .. are bytes of exclusive data of file  
 is calculated. Striping takes place at byte level and stored.  
 • Rarely used | byte = 8 bits  
 in disk3  
 Provides higher transfer rate as in RAID 0
- P0 is parity for bytes B0 and B1
  - $P_i = B_{2i} \oplus B_{2i+1}$ , here  $\oplus$  is exclusive-or operator

MTTF increases substantially ( $1/3$  of RAID1 =  $MTTF^2/3$ ), as 1 disk failure can be recovered from the data of the other 2 disks

exclusive value :  $\begin{cases} 0 & \rightarrow 1 \\ 1 & \rightarrow 0 \\ 0 & \end{cases}$

# RAID 4 (Block level level striping)

byte level  
f



- Striping takes place at block level
- Dedicated disk for parity blocks
- Provides higher throughput but, writing very slow writes. Disk3 has more ~~slow~~ writes as Parity needs to be updated for every data write.
- MTTF increases substantially (same as RAID3)

$$P_i = A_{2i} \oplus A_{2i+1}, \text{ here } \oplus \text{ is an exclusive-or operator}$$

~~Why writing throughput is slow? (for RAID3)~~

~~Calculate and update Parity, if anything gets updated in data, the corresponding parity has to be calculated as well.~~

- A means Block (4K or 8K bytes of storage)
- P is parity
- MTTF = Mean Time To Failure

~~So disk 3 gets more operations than any of those 2 disks.~~

~~So the whole write operation gets bottlenecked.~~

*slow writing speed has been addressed by RAID5.*

*by the write speed of parity disk 2*

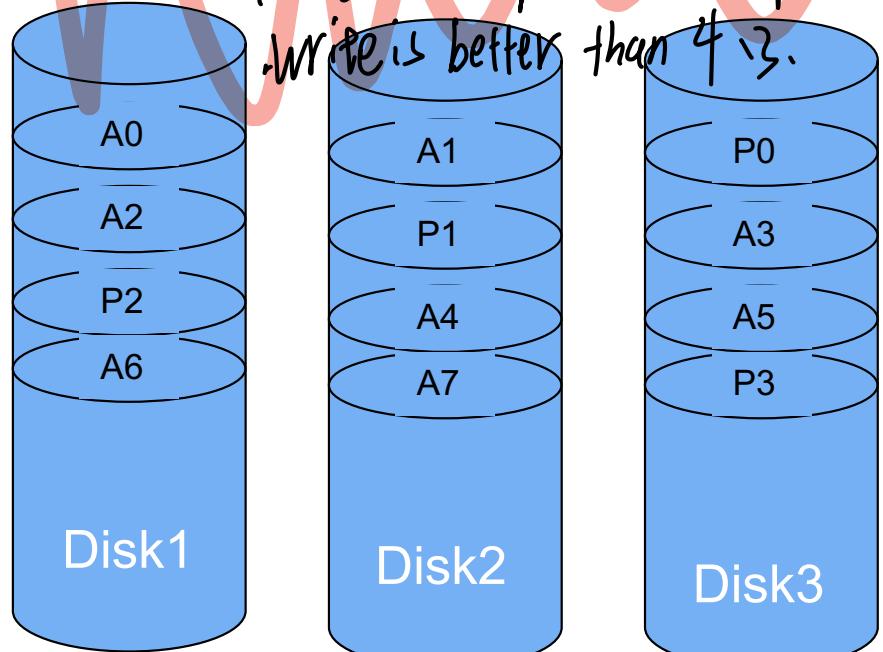
→ 163.4 MB/s

## RAID 5 with 3 disks (Block level striping)

*parity is distributed across disks no specific disk for storing parity.*

*hence There is no single disk with high number of writes.*

*Write is better than 4/3.*



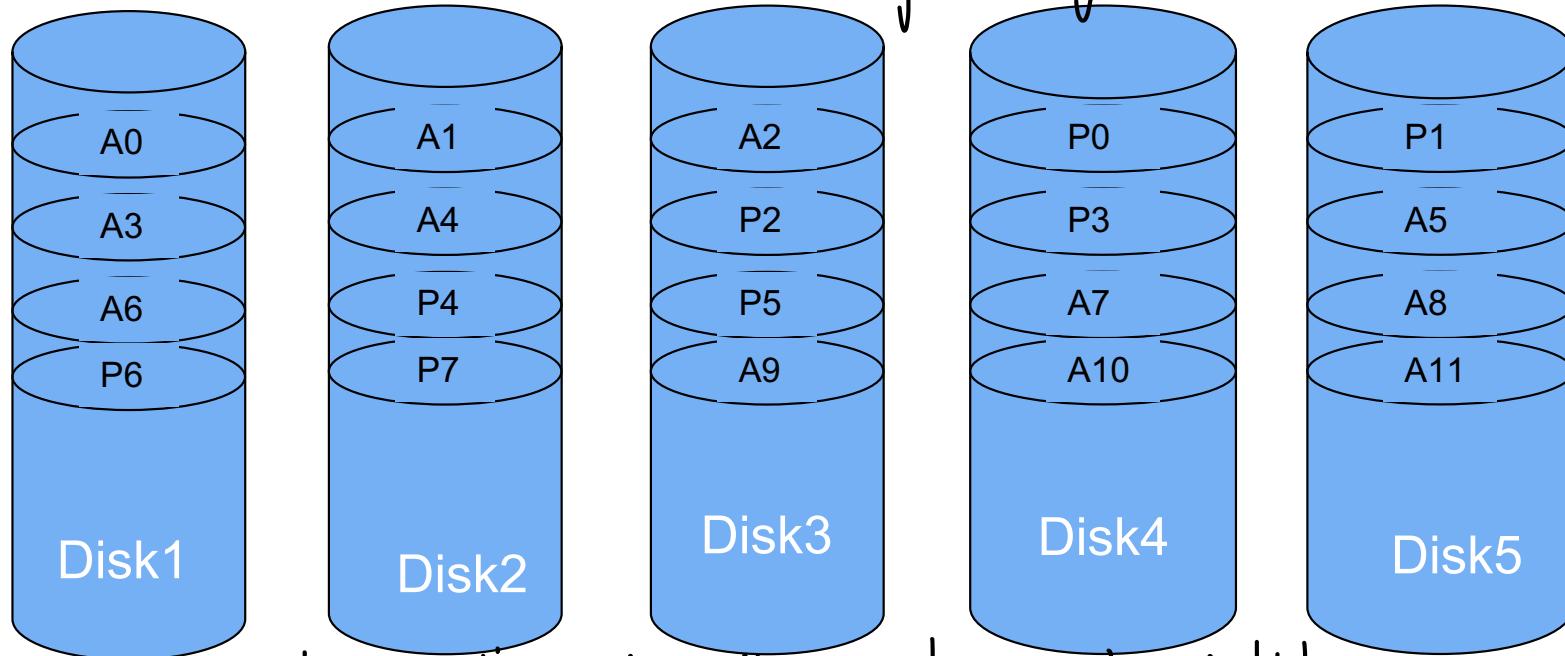
*The # of write operations are on average equal across three disks.*

- A0, A1, A2, A3, .. are contiguous blocks of data of a file
- Striping takes place at block level
- Parity blocks are also striped
  - read from multiple disks all in parallel.
- Provides higher throughput but slower writes but better than RAID 4 as Parity bits are distributed among all disks and the number of write operations on average equal among all 3 disks.
- MTTF increases substantially (same as RAID3)
- $P_i = A_{2i} \oplus A_{2i+1}$ , here  $\oplus$  is an exclusive-or operator

- A means Block (4K or 8K bytes of storage)
- P is parity
- MTTF = Mean Time To Failure

## RAID 6 (Block level level striping)

- A means Block (4K or 8K bytes of storage)
- P is parity (3行数据. 2个存 parity) parity blocks are distributed across the disks. cause higher # of blocks



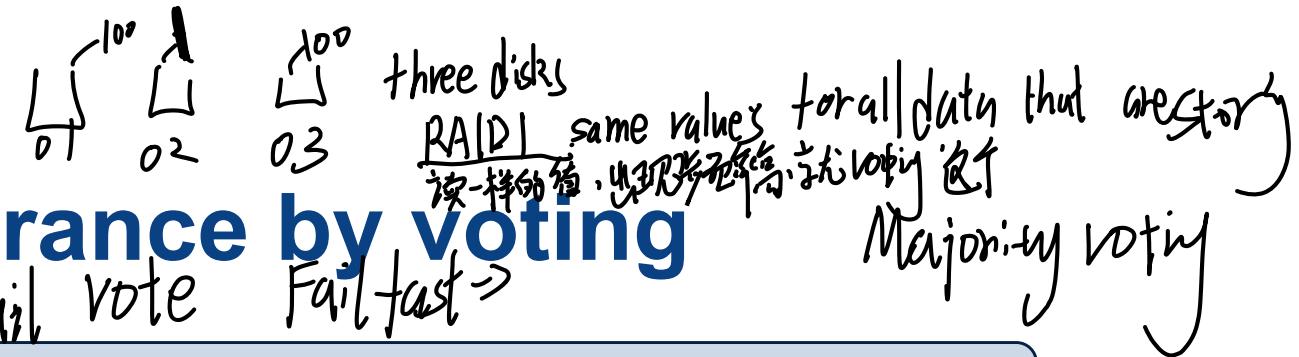
All disks have similar # of writes, reading throughput is higher

- Similar to RAID 5 except two parity blocks used.

if we have higher requirement for fault tolerance  $\rightarrow$  use this (more than one disk failure needs to be recovered) and system PRO (two disk failures can be safe to recover the data.)

CON ① Parity calculation and recovery process for RAID 6 is much more complex

Still needs to keep running

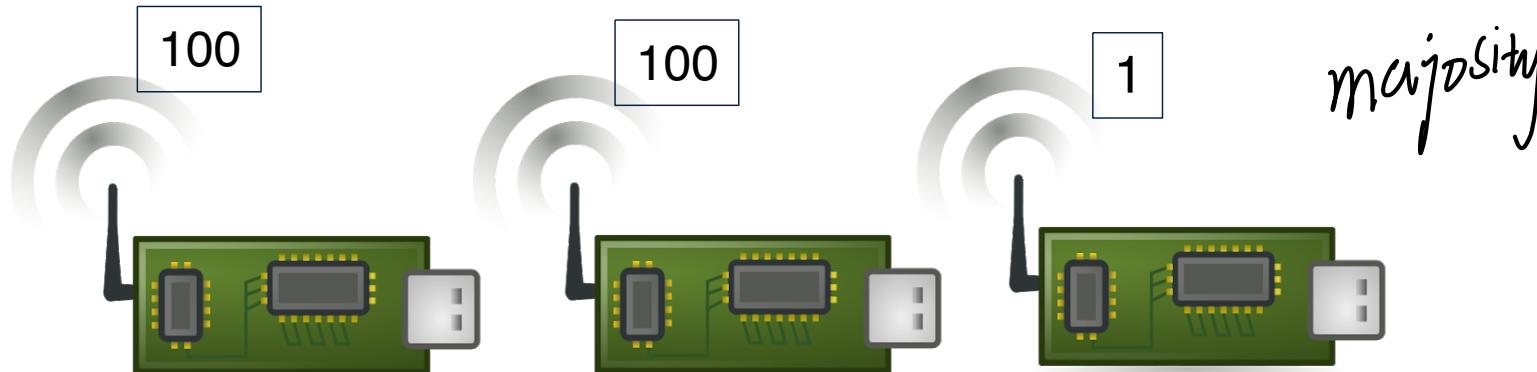


## Fault Tolerance by voting

Fail vote Fail fast  $\Rightarrow$

Majority voting

Use more than one module, voting for higher reliability



**Failvote** - Stops if there are no majority agreement

**Failfast (voting)** - Similar to failvote except the system senses which modules are available and uses the majority of the available modules.

- A 10 module Failfast system continues to operate until the failure of 9 modules whereas Failvote stops when 5 modules fail.
- Failfast system has better availability than failvoting (since failvote stops when there is no majority agreement).

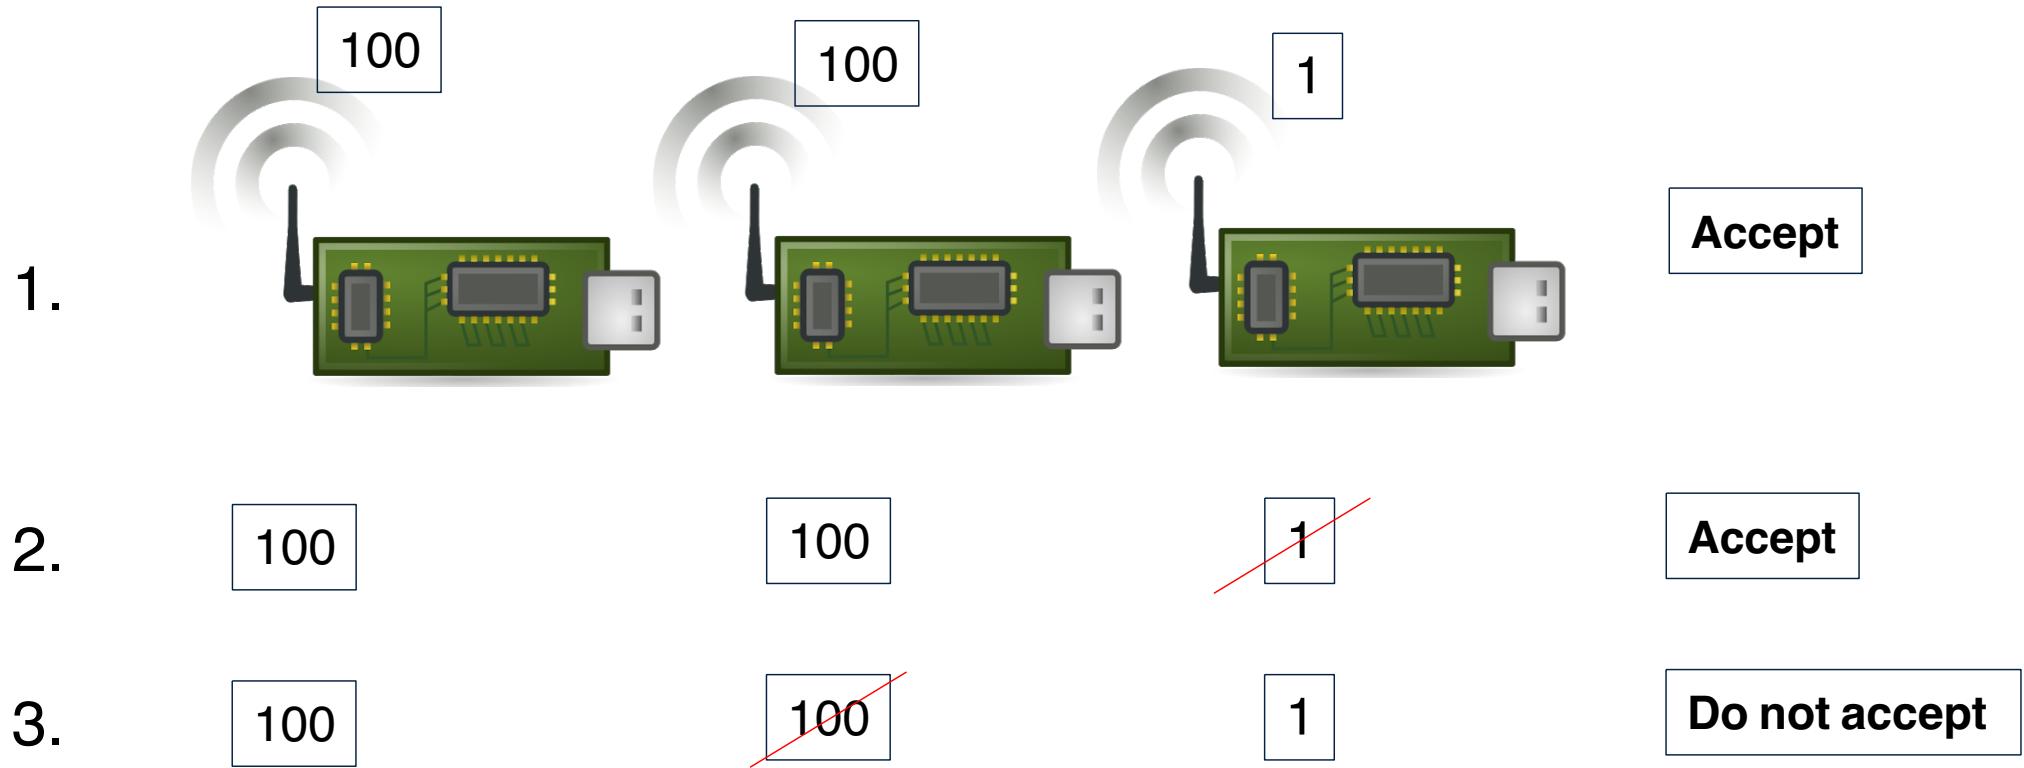
fail rate 5个有3个一致 ✓  
4个有3个一致 ✓

fail fast 系统能知道到哪个 module 正常运行，只对正常运行的计算进行 majority voting

5个仅有两个正常运行  $\Rightarrow$  这2个进行 majority voting 2个必须一致 ✓  
available

# Fault Tolerance ...

Failvote needs majority agreement to accept an action (eg, Read/write)



If we start with 10 devices, the system works as long as 6 of them are working. Action is accepted when 6 or more agreeing on the decision. The moment 5<sup>th</sup> one fails, system stops as there cannot be 6 devices agreeing



## Fault Tolerance ...

In Failfast, we are only concerned of majority among the working ones. We are assuming that we can tell which ones are working. Hence we can continue to operate until 2 working ones and if both agree we can proceed with the action. But if they differ the system stops.

- 0 devices are faulty, we have 10 working and we need at least 6 to agree
- 1 device is faulty, we have 9 working and we need at least 5 to agree
- 2 devices are faulty, we have 8 working and we need at least 5 to agree
- 3 devices are faulty, we have 7 working and we need at least 4 to agree
- 4 devices are faulty, we have 6 working and we need at least 4 to agree
- 5 devices are faulty, we have 5 working and we need at least 3 to agree
- 6 devices are faulty, we have 4 working and we need at least 3 to agree
- 7 devices are faulty, we have 3 working and we need at least 2 to agree
- 8 devices are faulty, we have 2 working and we need both to agree
- 9 devices are faulty, we have 1 working and we have to stop as nothing to compare!

# Fault Tolerance for disks

**Supermodule** – Naturally, a system with multiple hard disk drives is expected to function with only one working disk (use voting when multiple disks are working/available, but still work even when only one is available)

 ready from three disks if only one is available  
but I still want to use that disk to write or read.

Supermodule: still keep functioning with one working device or one working disk.  
3 devices  $\rightarrow$  fail vote  $\rightarrow$  2 devices to agree (otherwise system fails)  $\Rightarrow$  3个中有2个坏  
 $\left\{ \begin{array}{l} \text{3坏} \\ \text{2坏} \end{array} \right\} \Rightarrow$  系统失败.

if fail fast  $\Rightarrow$  3个坏  $\rightarrow$  系统失败  
rotating  $\left\{ \begin{array}{l} \text{3坏} \\ \text{2坏} \end{array} \right\}$  - 1 by majority  
 $\Rightarrow$  no majority.  
 $\Rightarrow$  the system will fail.

e.g. 2 devices  $\rightarrow$  system fails for 1 failure. C if PCI device fail) =  $P$ )  
 $\Pr(\text{such a system to fail}) = P + P$        $MTTF = \frac{1}{2P}$  (inverse)  $= \frac{1}{2} \times \frac{1}{P}$   $= \frac{1}{2} \text{ years}$   
 $\Rightarrow$  1 device's MTTF  
 $m = 10 \text{ years}$

## Availability of failvote systems

If there are n events, mean time to the first event =  $m/n$

Consider a system with modules each with MTTF of 10 years

Failvoting with 2 devices:  $\rightarrow$  system down for 1 device failure

$MTTF = 10/2 = 5 \text{ years}$  (system fails with 1 device failure)

Failvoting with 3 devices:  $\rightarrow$  system down for 2 devices failure.

$MTTF = 10/3 \text{ for the first failure} + 10/2 \text{ for 2nd failure} = 8.3 \text{ years.}$

Lower availability for higher reliability (multiple modules agreeing on a value means that value is more likely to be accurate/reliable)

1st MTTF + 2nd MTTF = MTTF of full system.  
 one module  $\leftarrow \frac{m}{n}$        $= \frac{10}{3} + \frac{10}{2} = 8.3$

But, but, but.... cannot we have both??



MTTF

$n$ : # of available devices. (3 for 1st, 2 for 2nd)

Current

## Fault tolerance with repair

$MTTR \leq 2 \text{ hours}$

$MTTF \ll MTTR$  待修时间不用很长.  $MTTR$  故障时间

With repair of modules: the faulty equipment is repaired with an average time of MTTR (mean time to repair) as soon as a fault is detected  
(Sometimes MTTR is just time needed to replace)

Typical Values for recent disks:

MTTR = Few hours (assuming we stock spare disks) to 1 Day

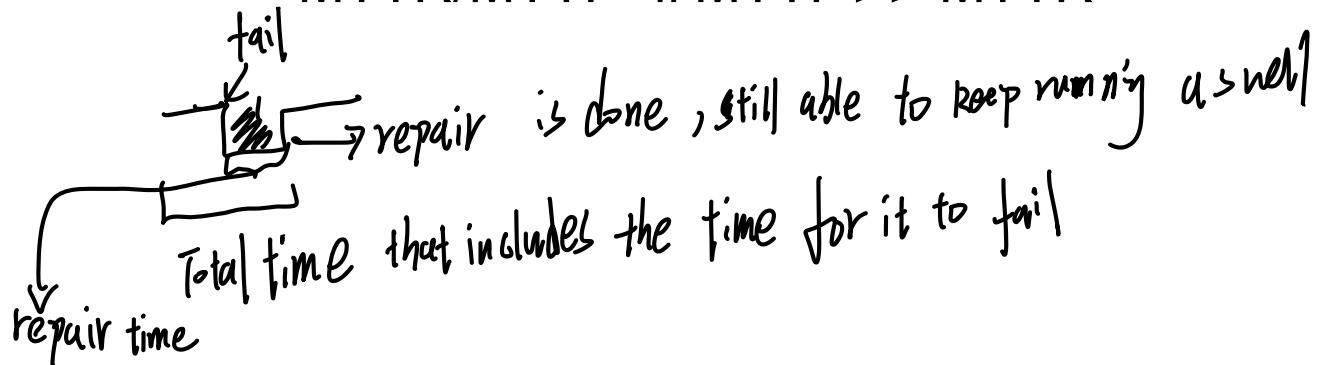
MTTF = 750000 hours (~ 86 years) [hard fault]

Probability of a particular module is not available

$$= MTTR / (MTTF + MTTR)$$

repair is done as soon as fault is detected

$$\approx MTTR / MTTF \text{ if } MTTF \gg MTTR$$



# Fault tolerance of a supermodule with repair

Probability that (n-1) modules are unavailable,  $P_{n-1} = \frac{1}{MTTF} \times \left(\frac{MTTR}{MTTF}\right)^{n-1}$

Probability that a particular  $i^{th}$  module fails,  $P_f = \frac{1}{MTTF}$

Probability that the system fails with a particular  $i^{th}$  module failing last =

$$(n-1) \text{ are unavailable : } P_{n-1} \times P_f = \frac{1}{MTTF} \times \left(\frac{MTTR}{MTTF}\right)^{n-1}$$

last are fails

Probability that a supermodule fails due to any one of the n modules failing last, when other (n-1) modules are unavailable

3 devices , 3 different ways that the last one fails

1 1 1



prob that supermodel

system fails is

$$n \times \frac{1}{MTTF} \times \left(\frac{MTTR}{MTTF}\right)^{n-1}$$



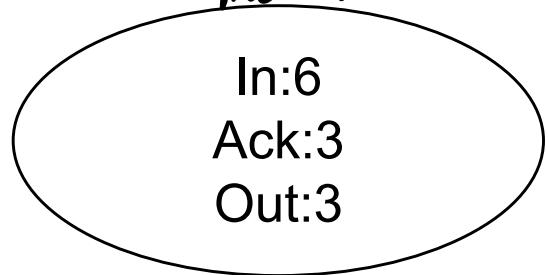
# Communication reliability

ensure the communication  
between them is reliable

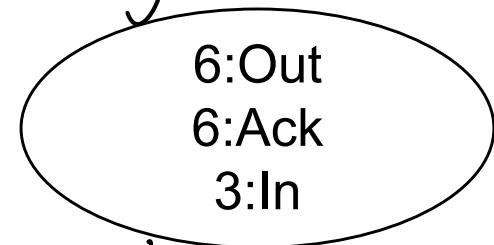
Out = #messages sent

In = #messages received

the data that has been transmitted and received directly is ensured.



node A



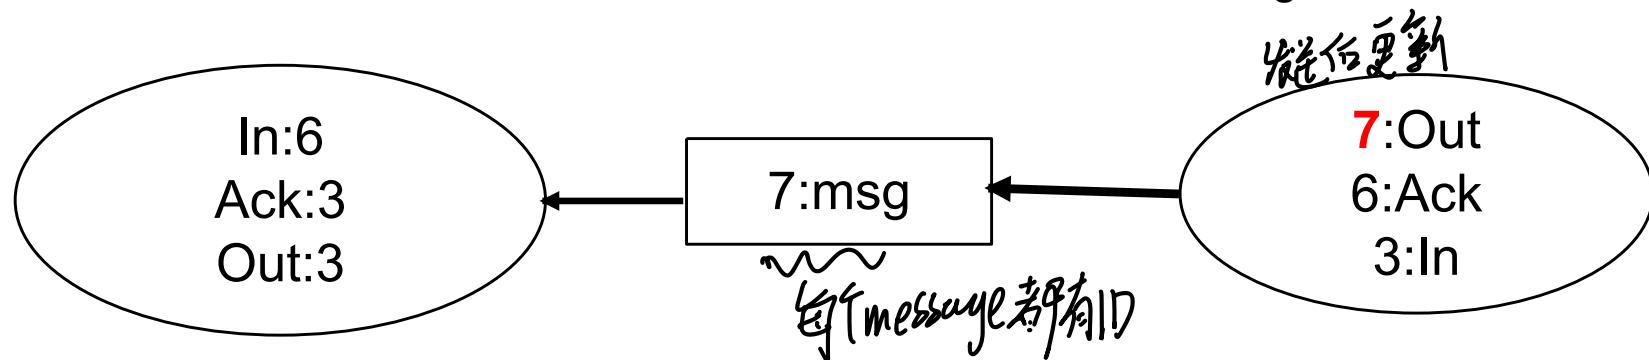
node B

# Communication reliability

Out = #messages sent

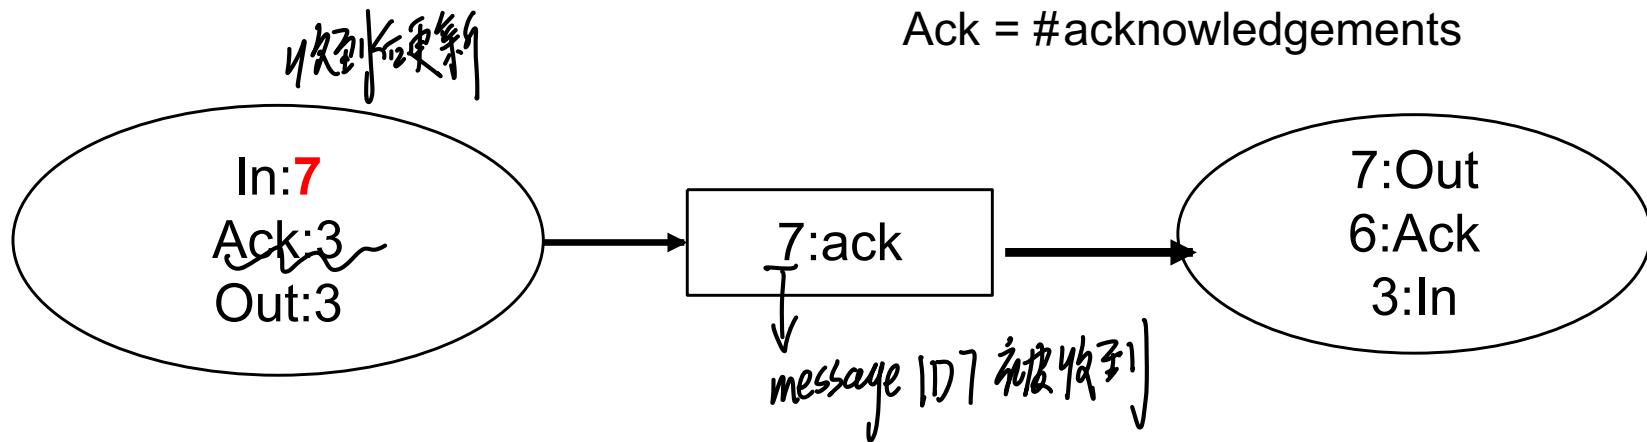
In = #messages received

Ack = #acknowledgements



# Communication reliability

## Reliable message passing



Out = #messages sent

In = #messages received

Ack = #acknowledgements



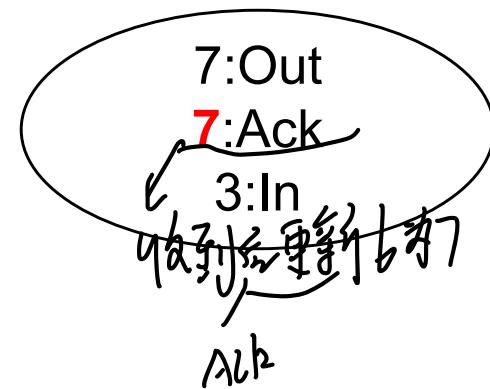
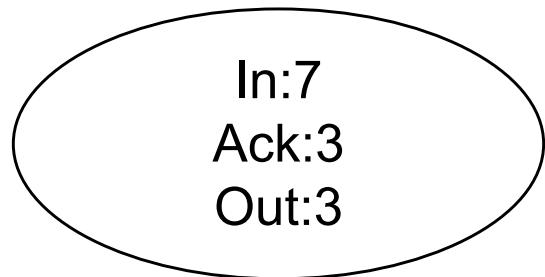
# Communication reliability

Reliable message passing

Out = #messages sent

In = #messages received

Ack = #acknowledgements

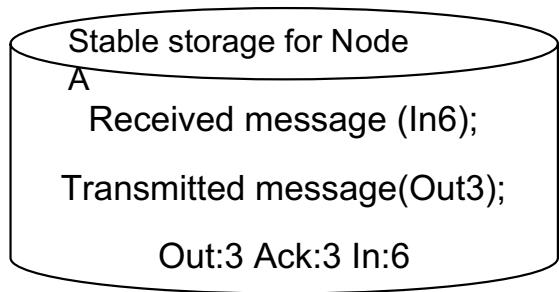


to make sure each message does not get lost  $\rightarrow$  stable communication

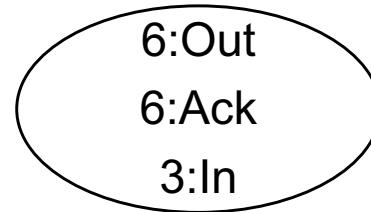
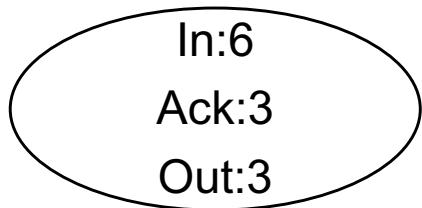
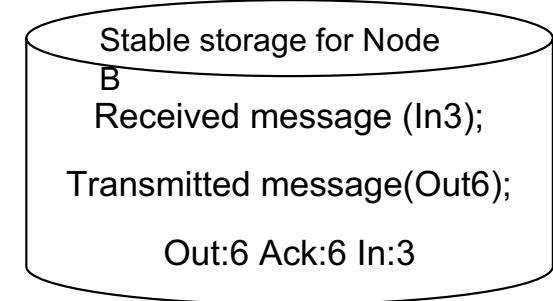


Stable storage  $\rightarrow$  fault tolerance (Data not get lost easily)

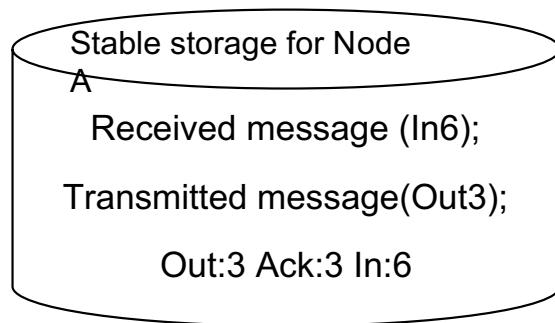
# Communication reliability



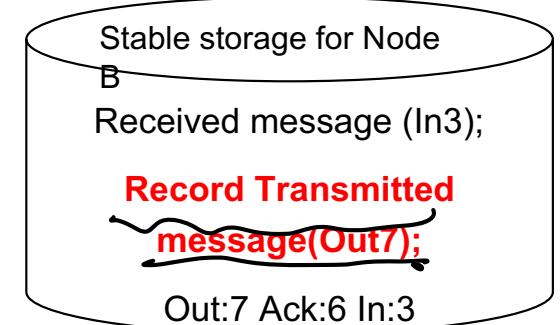
Checkpoint before sending  
a message



# Communication reliability



Checkpoint before sending  
a message

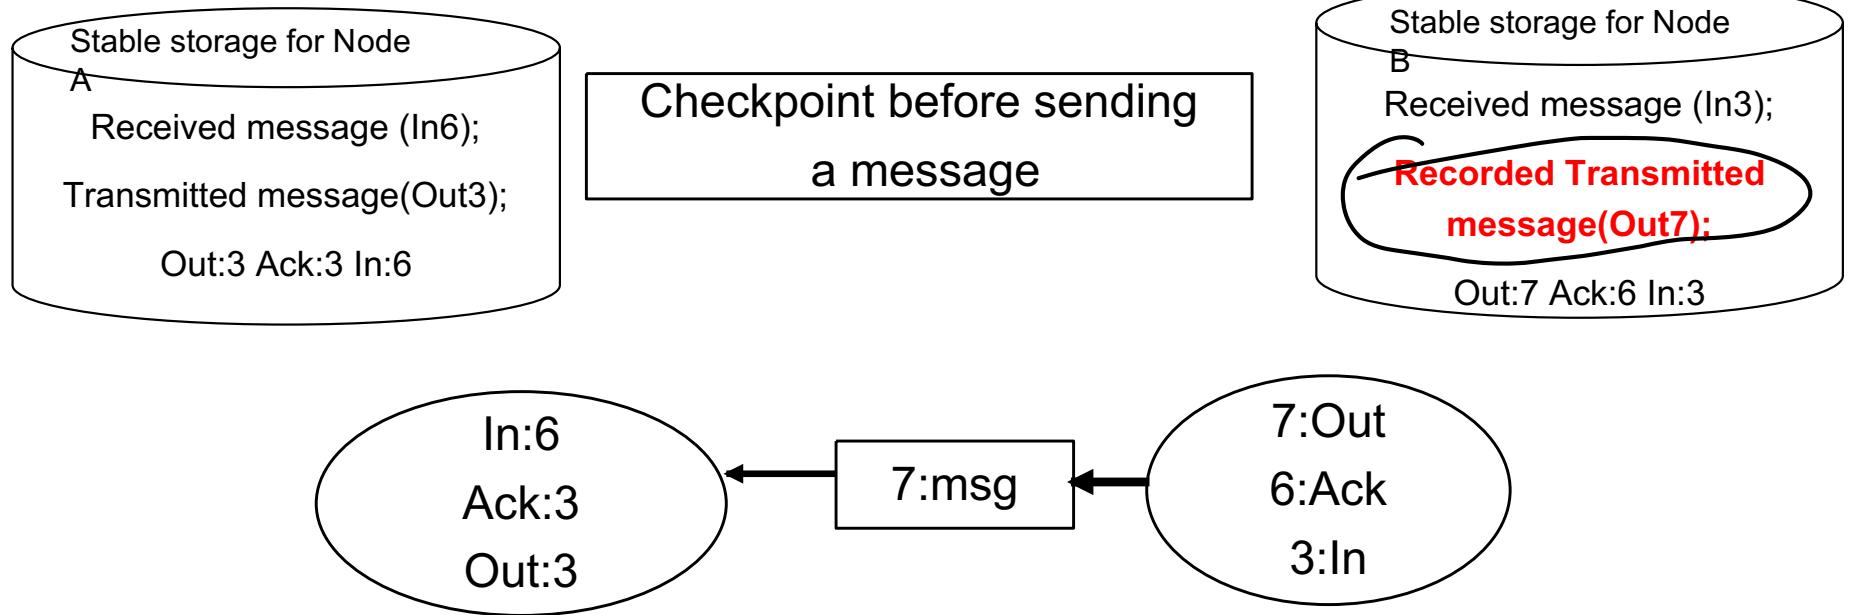


In:6  
Ack:3  
Out:3

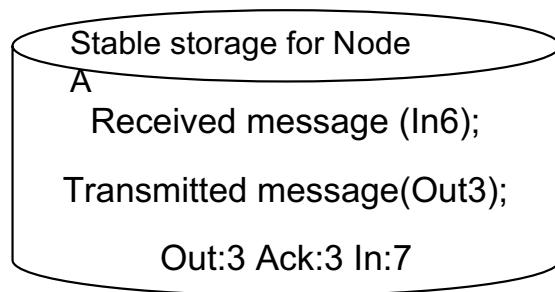
7:Out  
6:Ack  
3:In



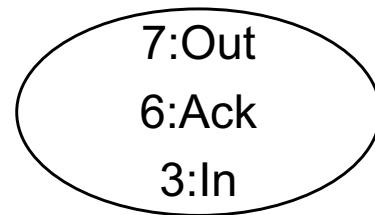
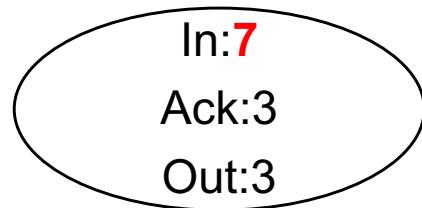
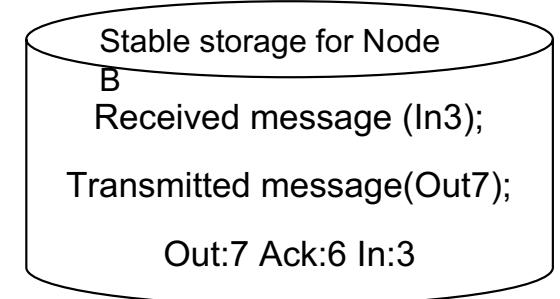
# Communication reliability



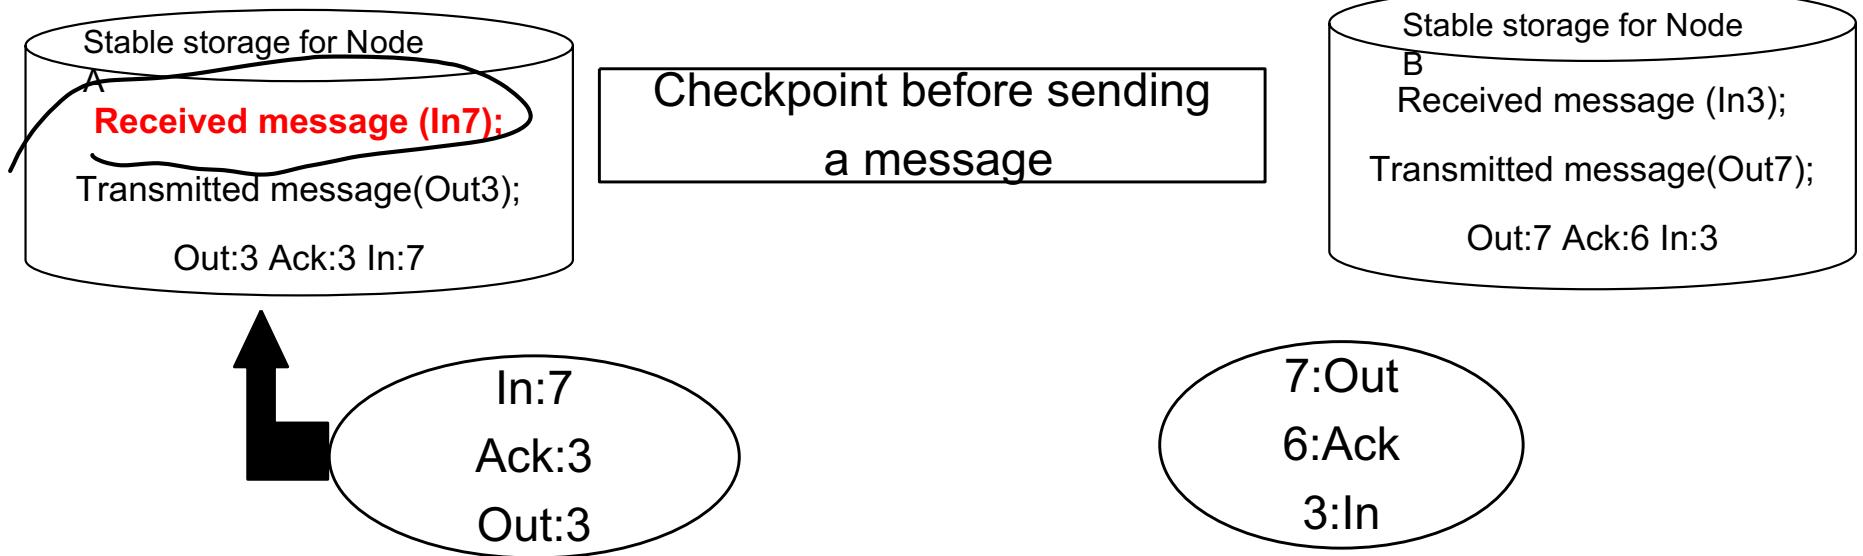
# Communication reliability



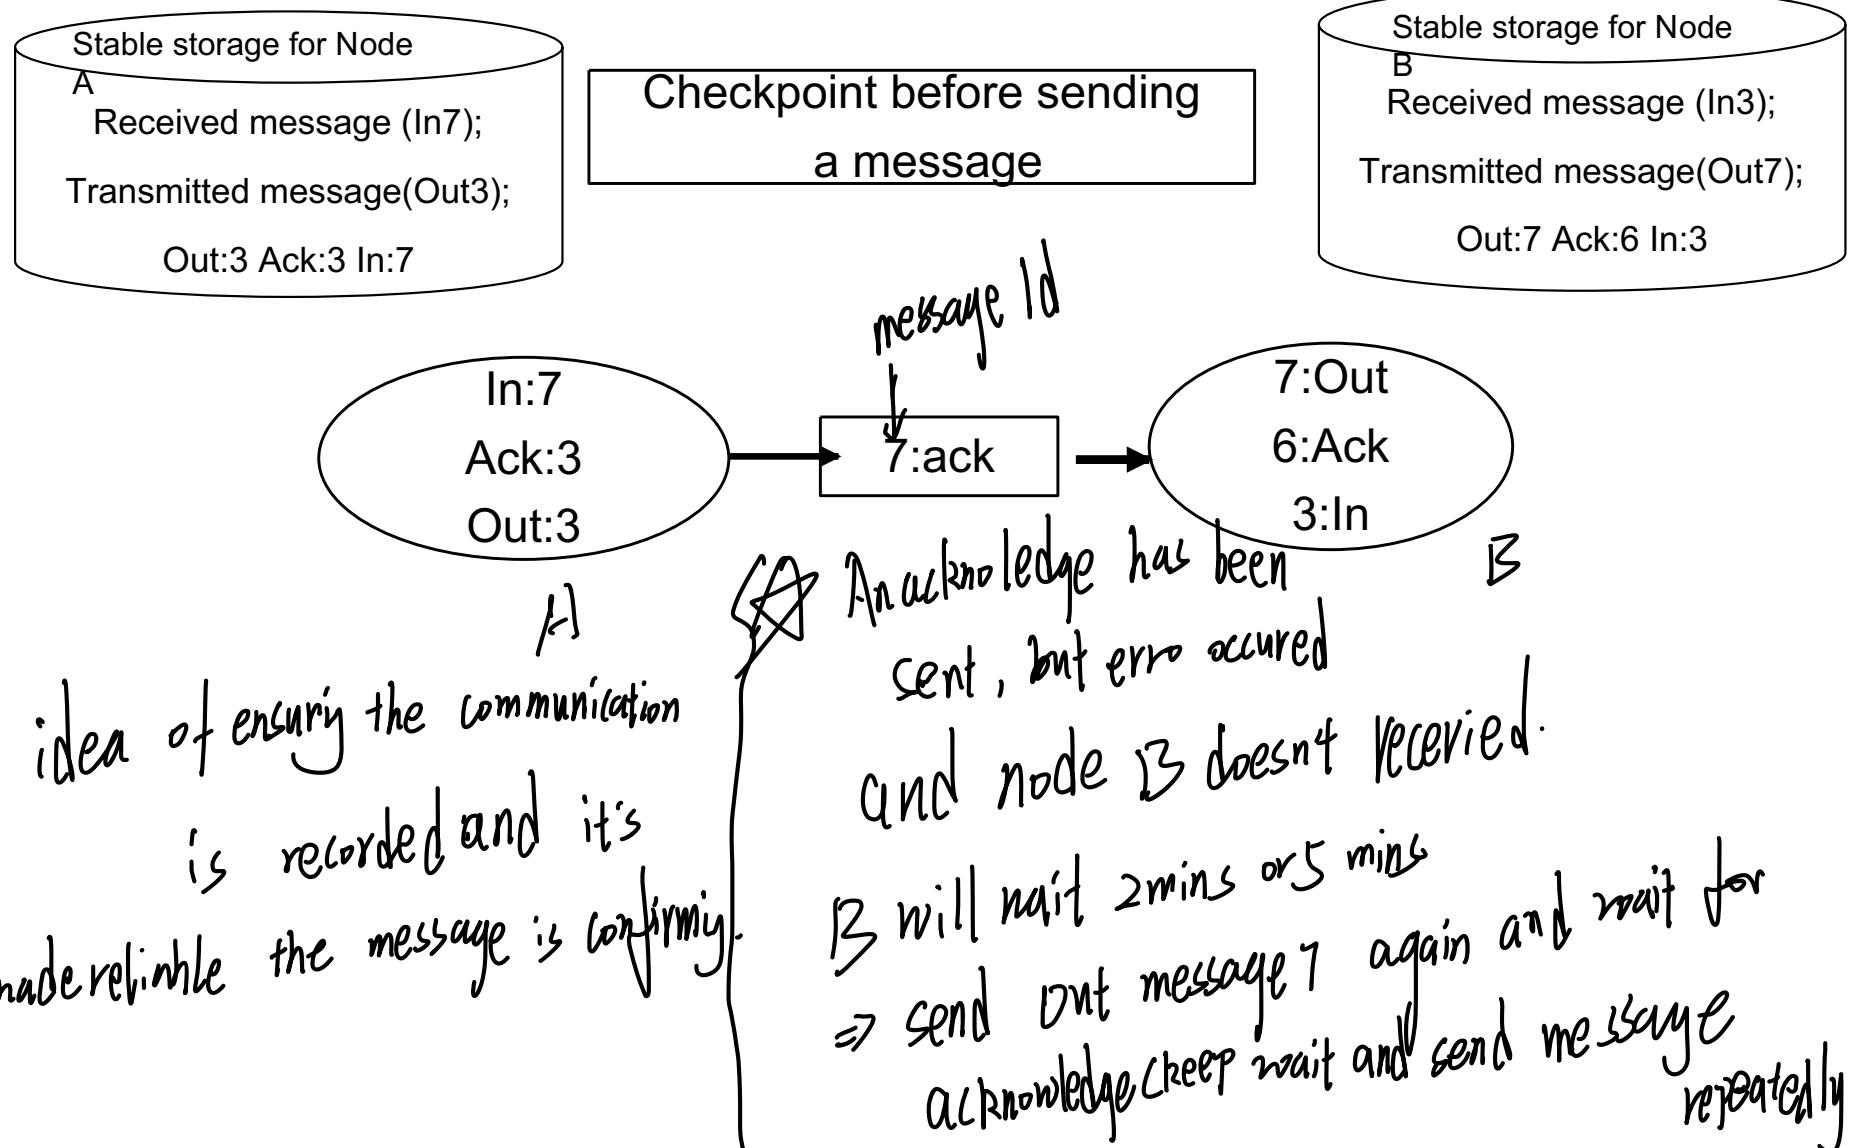
Checkpoint before sending  
a message



# Communication reliability



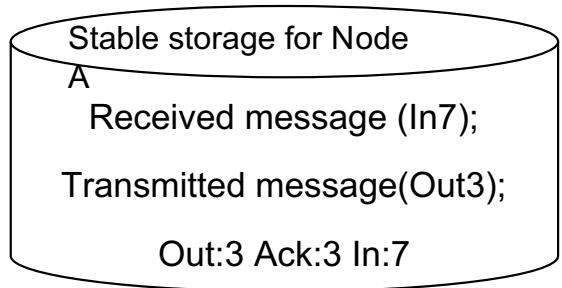
# Communication reliability



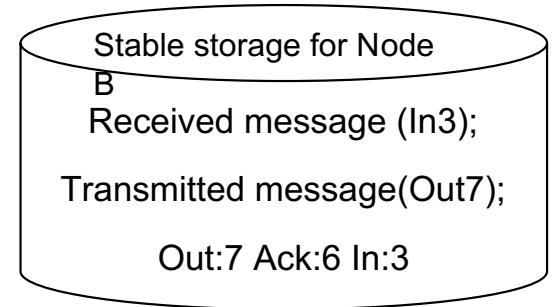


# Communication reliability

As long as it does not receive the acknowledgement.  
Once receive acknowledgement  $\Rightarrow$  message is done completely that  
it has been received by the other  
node correctly.  $\checkmark$



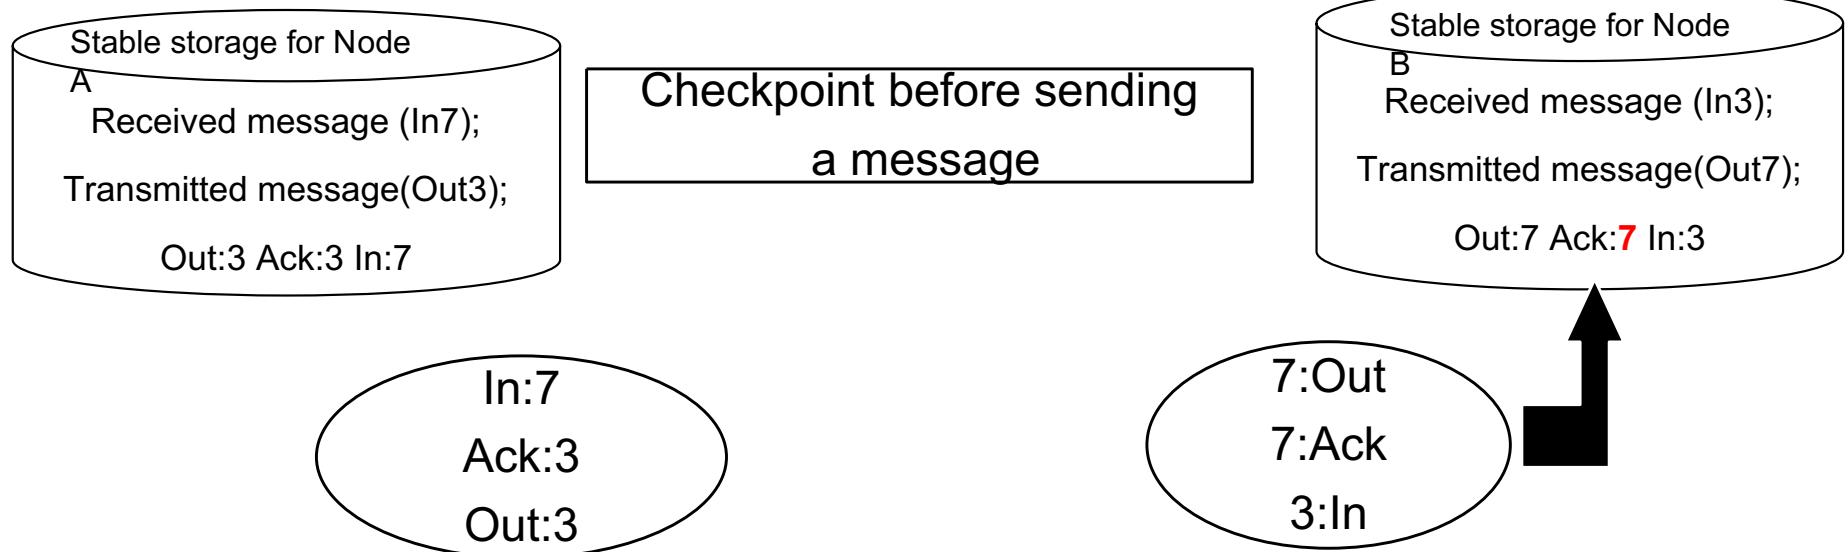
Checkpoint before sending  
a message



In:7  
Ack:3  
Out:3

7:Out  
**7:Ack**  
3:In

# Communication reliability



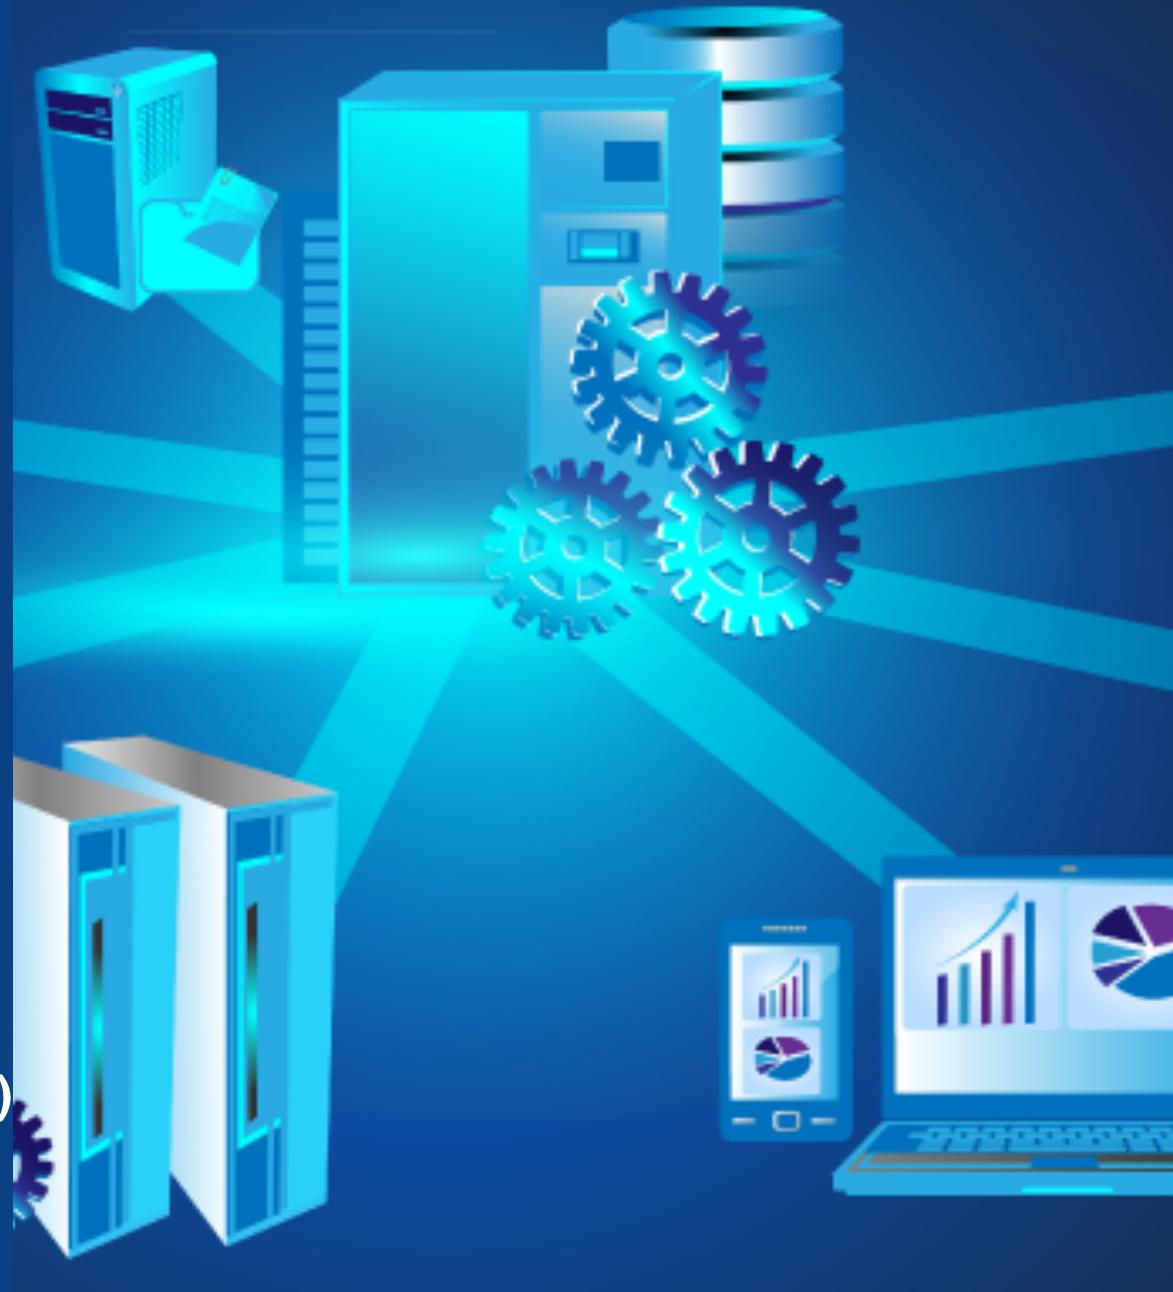


# COMP90050 Advanced Database Systems

## Semester 2, 2024

Lecturer: Farhana Choudhury (PhD)

Live lecture – Week 8

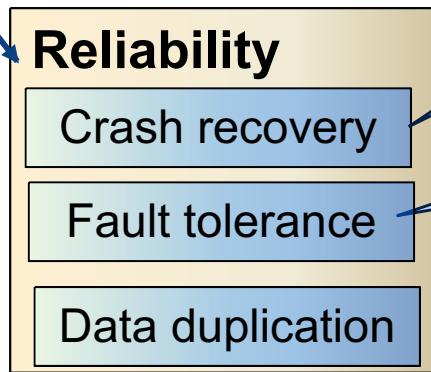
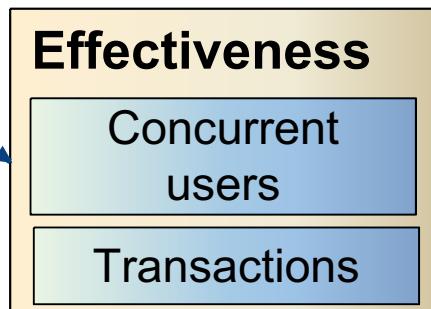
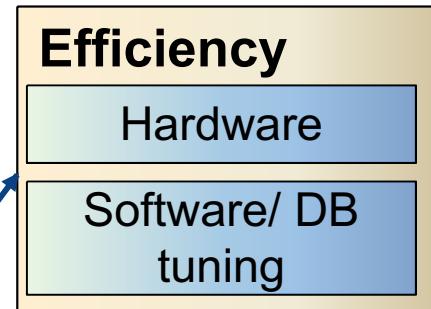
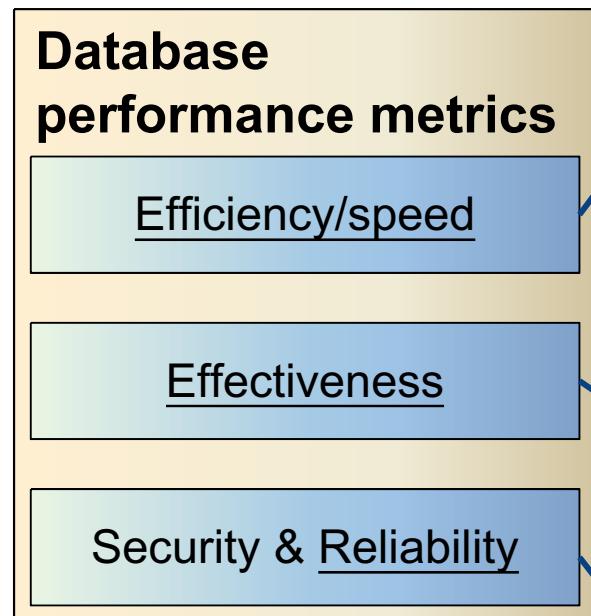




# Presentation details

- Announcement
- Register asap
- Check FAQs

# Core Concepts of Database management system



Reactive

Preemptive

Hardware:

- Arrangement of multiple disks
- Voting among multiple disks/modules for reliability
- Communication reliability

# Fault Tolerance

The property that enables a system to continue operating properly in the event of the failure of some of its components.

We have covered

- Statistics crash course
- Lifecycle of a system
- Different fault tolerance techniques
- Communication reliability



#  $P(A \text{ and } B) = P(A) * P(B)$  assuming A and B are independent events.

#  $P(A \text{ or } B) = P(A) + P(B) - P(A \text{ and } B)$

=  $P(A) + P(B) - P(A) * P(B)$  [ Assuming A and B are independent]

$P(A) + P(B)$  [if  $P(A)$  and  $P(B)$  are very small]

## This is not a math subject!

Understanding the concepts (e.g., the event that both disks fail at the same time is statistically ‘how much’ less likely than a single disk fail)



# Fault tolerance by RAID

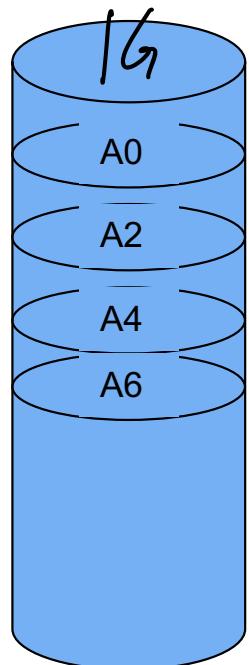
Redundant Array of Independent Disks – different ways to combine multiple disks as a unit for fault tolerance or performance improvement, or both of a database system

## Choosing the suitable RAID level

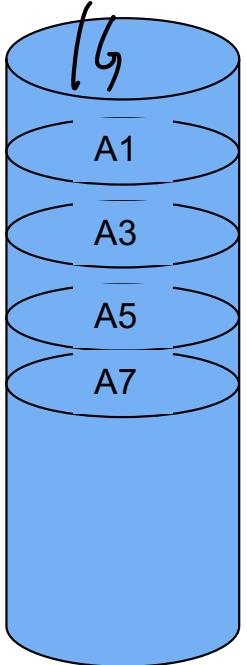
The factors to consider:

- Reliability
- Performance
- Storage utilization
- Price/number of disks
- Additional overhead?

# Example summary: RAID 0 and RAID 2



Block level striping



Bit level striping (rare)

Provides balanced I/O of disk drives

Provides higher throughput (~doubles)

Any disk failure will be catastrophic

MTTF reduces by a factor of 2

Higher throughput at the cost of increased vulnerability to failures

Calculation of MTTF values – in tutorials

Storage utilization?

Number of disks needed?

Overhead difference?

## Comparison:

RAID 0 VS RAID 2

- ① Storage utilization: 对上面两个是平等的，但有时会有 mirror 镜像 (RAID 1). 2 blocks 只有 一半 ~~50%~~ pure info. (100% storage utilization vs 50% storage utilization)
- ② # of disks needed: (RAID 0 VS RAID 2)
- ③ overhead difference? Stripping in bit level ~~lost~~ more than stripping in block level  
In ① scenario! When we write ~~more dividing or partitioning overhead~~ involved. When we read, we need to merge data from two disks more frequently to form meaningful data that can be displayed to users (much more overhead on RAID 0)
- ④ reliability performance: Any disk failure will be catastrophic (if one block failure, then the whole system will ~~not~~ <sup>same</sup> be able to provide any sort of data. (data will become meaningless))
- ⑤ System throughput: we can read from two disks in parallel, higher throughput  $\Rightarrow$  double higher
- ⑥ balanced I/O of disk drives: the same # of write and read operations happen at both disks. (No bottleneck disk that we have to wait for disk write finish to move to the other one)



# RAID – exercise

Which of the following RAID setting can safely recover data even for two disk failures?

RAID 5 with 3 disks

0%

RAID 1 with 3 disks

✓ Three replications  $\geq$  two failure. Still have third disk (can do read and write)

0%

None of the above

0%

Time for a poll - [Pollev.com/farhanachoud585](https://pollev.com/farhanachoud585)



If somethy failed can I still recover data?

RAID is commonly used for fault tolerance system.

use RAID 1 to store research data (people's GPS data)

\* Data: As people move, their phone continuously send GPS to the server. Over sometime that's big amount of data get collected.

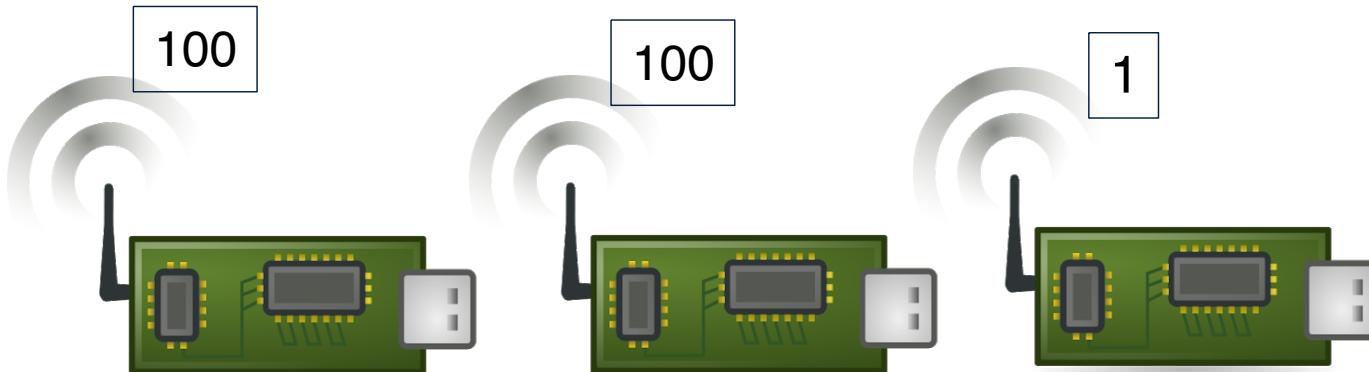
Why?: Data is important, if I lose it I cannot collect them again (So I need mirror to duplicate dataset) - Just guarantee I don't lose it.  
Only for research purpose <sup>only</sup> I don't care anything about read throughput. Not many user or something really needs to read faster.

how much reliability or better throughput?  $\Rightarrow$  RAID levels

# Fault Tolerance by voting

*Check if the data one got is correct*

Use more than one module, voting for higher reliability



**Failvote** - Stops if there are no majority agreement

**Failfast (voting)** - Similar to failvote except the system senses which modules are available and uses the majority of the available modules.

**Supermodule** – A system with multiple modules that use voting when multiple modules are working/available, but still work even when only one is available

*→ if only one device can run, the system can still keep operating.*



# Fault Tolerance by voting

In a Failfast system, which of the following cases we can accept an action?

fail vote  $\geq$  check the total # of devices the do majority voting  
 ✓ for or fail fast,  $\leq$  available (working) devices

Total number of devices	No. of working devices	No. of agreeing devices	Accept?
10	6	4	✓
10	6	3 (3 is not a majority only 6)	X
10	5	3	✓
5	5	3	✓
5	4	2	X
5	2	2	✓
5	1	-	X



# Fault Tolerance by voting

In a Failfast system, which of the following cases we can accept an action?

Total number of devices	No. of working devices	No. of agreeing devices	Accept?
10	6	4	y
10	6	3	n
10	5	3	y
5	5	3	y
5	4	2	n
5	2	2	y
5	1	-	n

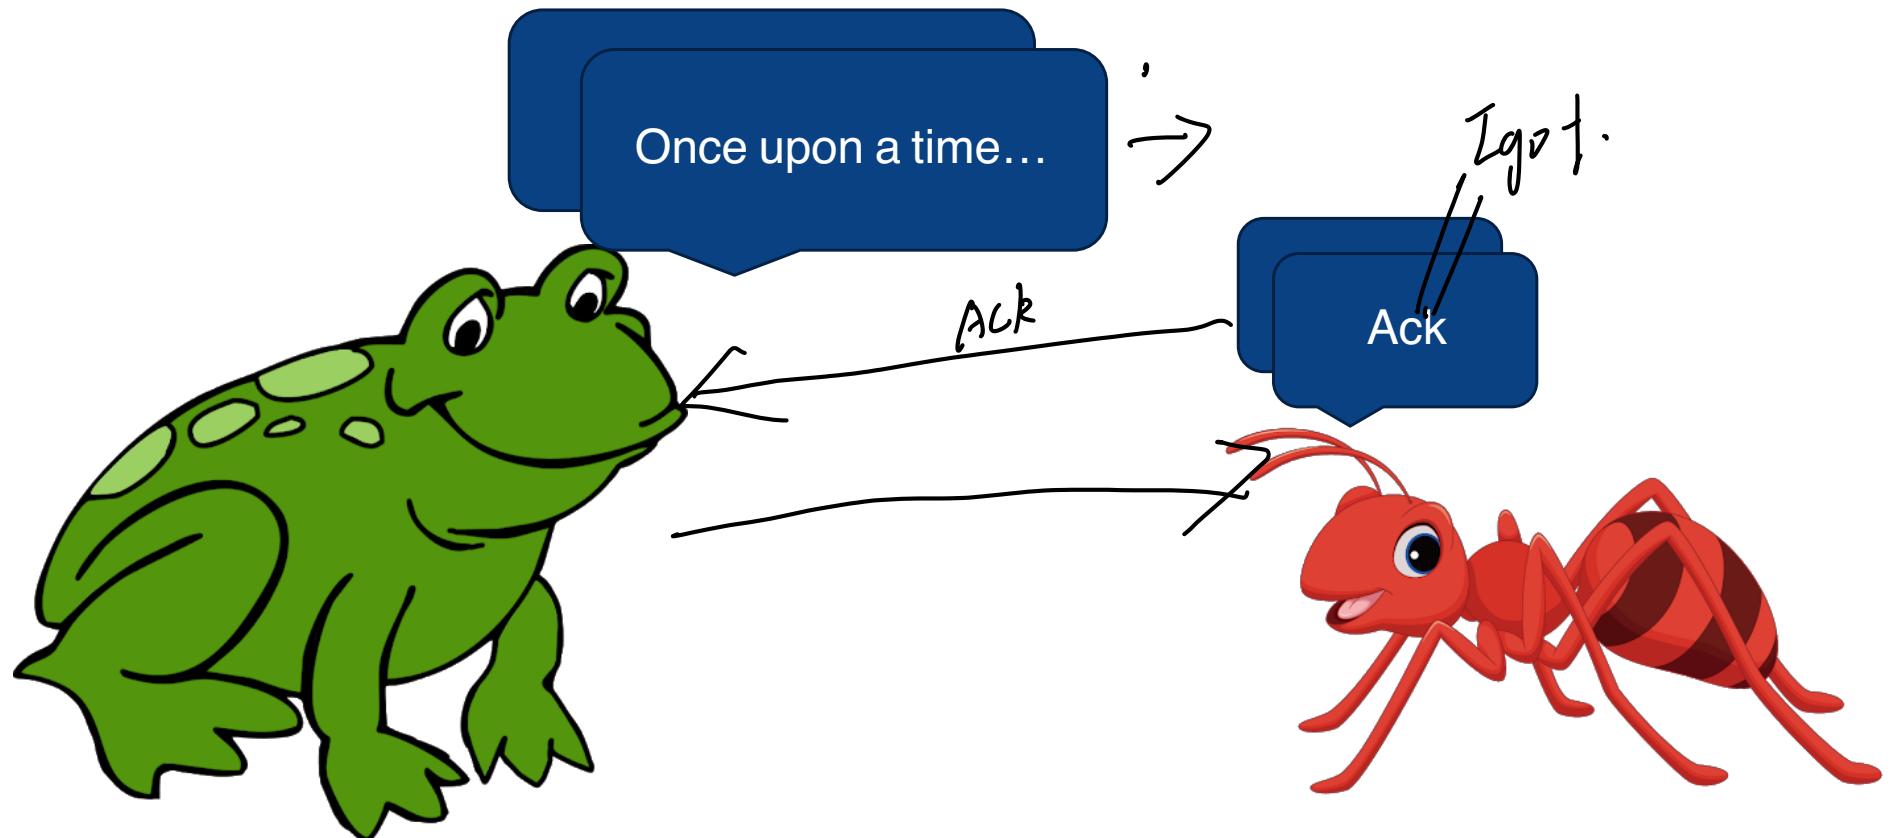
# Communication reliability

- Why important for databases? for database? why we learnt that?

~~As the nodes of distributed database need to communicate with each other,~~  
also : Input in a database does not only from users, but also come from automated or other systems. If input comes  $\Rightarrow$  some acknowledgement or message passing involved  $\Rightarrow$  check it the message or the input has been received correctly.

Additionally : many functionalities need ensure the reliability of communication for consistency: insertion, updates. (communication and messages need to be passed.)

# Communication reliability



# Communication reliability

If acknowledgement not received -



# Core Concepts of Database management system

