



Adaptive Query Processing



Liumingzi Zhu 727740
Xiaolu Zhang 886161

Luyao Chen 1266572
Muhan Guan 1407870





Table of Contents



01

Introduction

02

Related Works

03

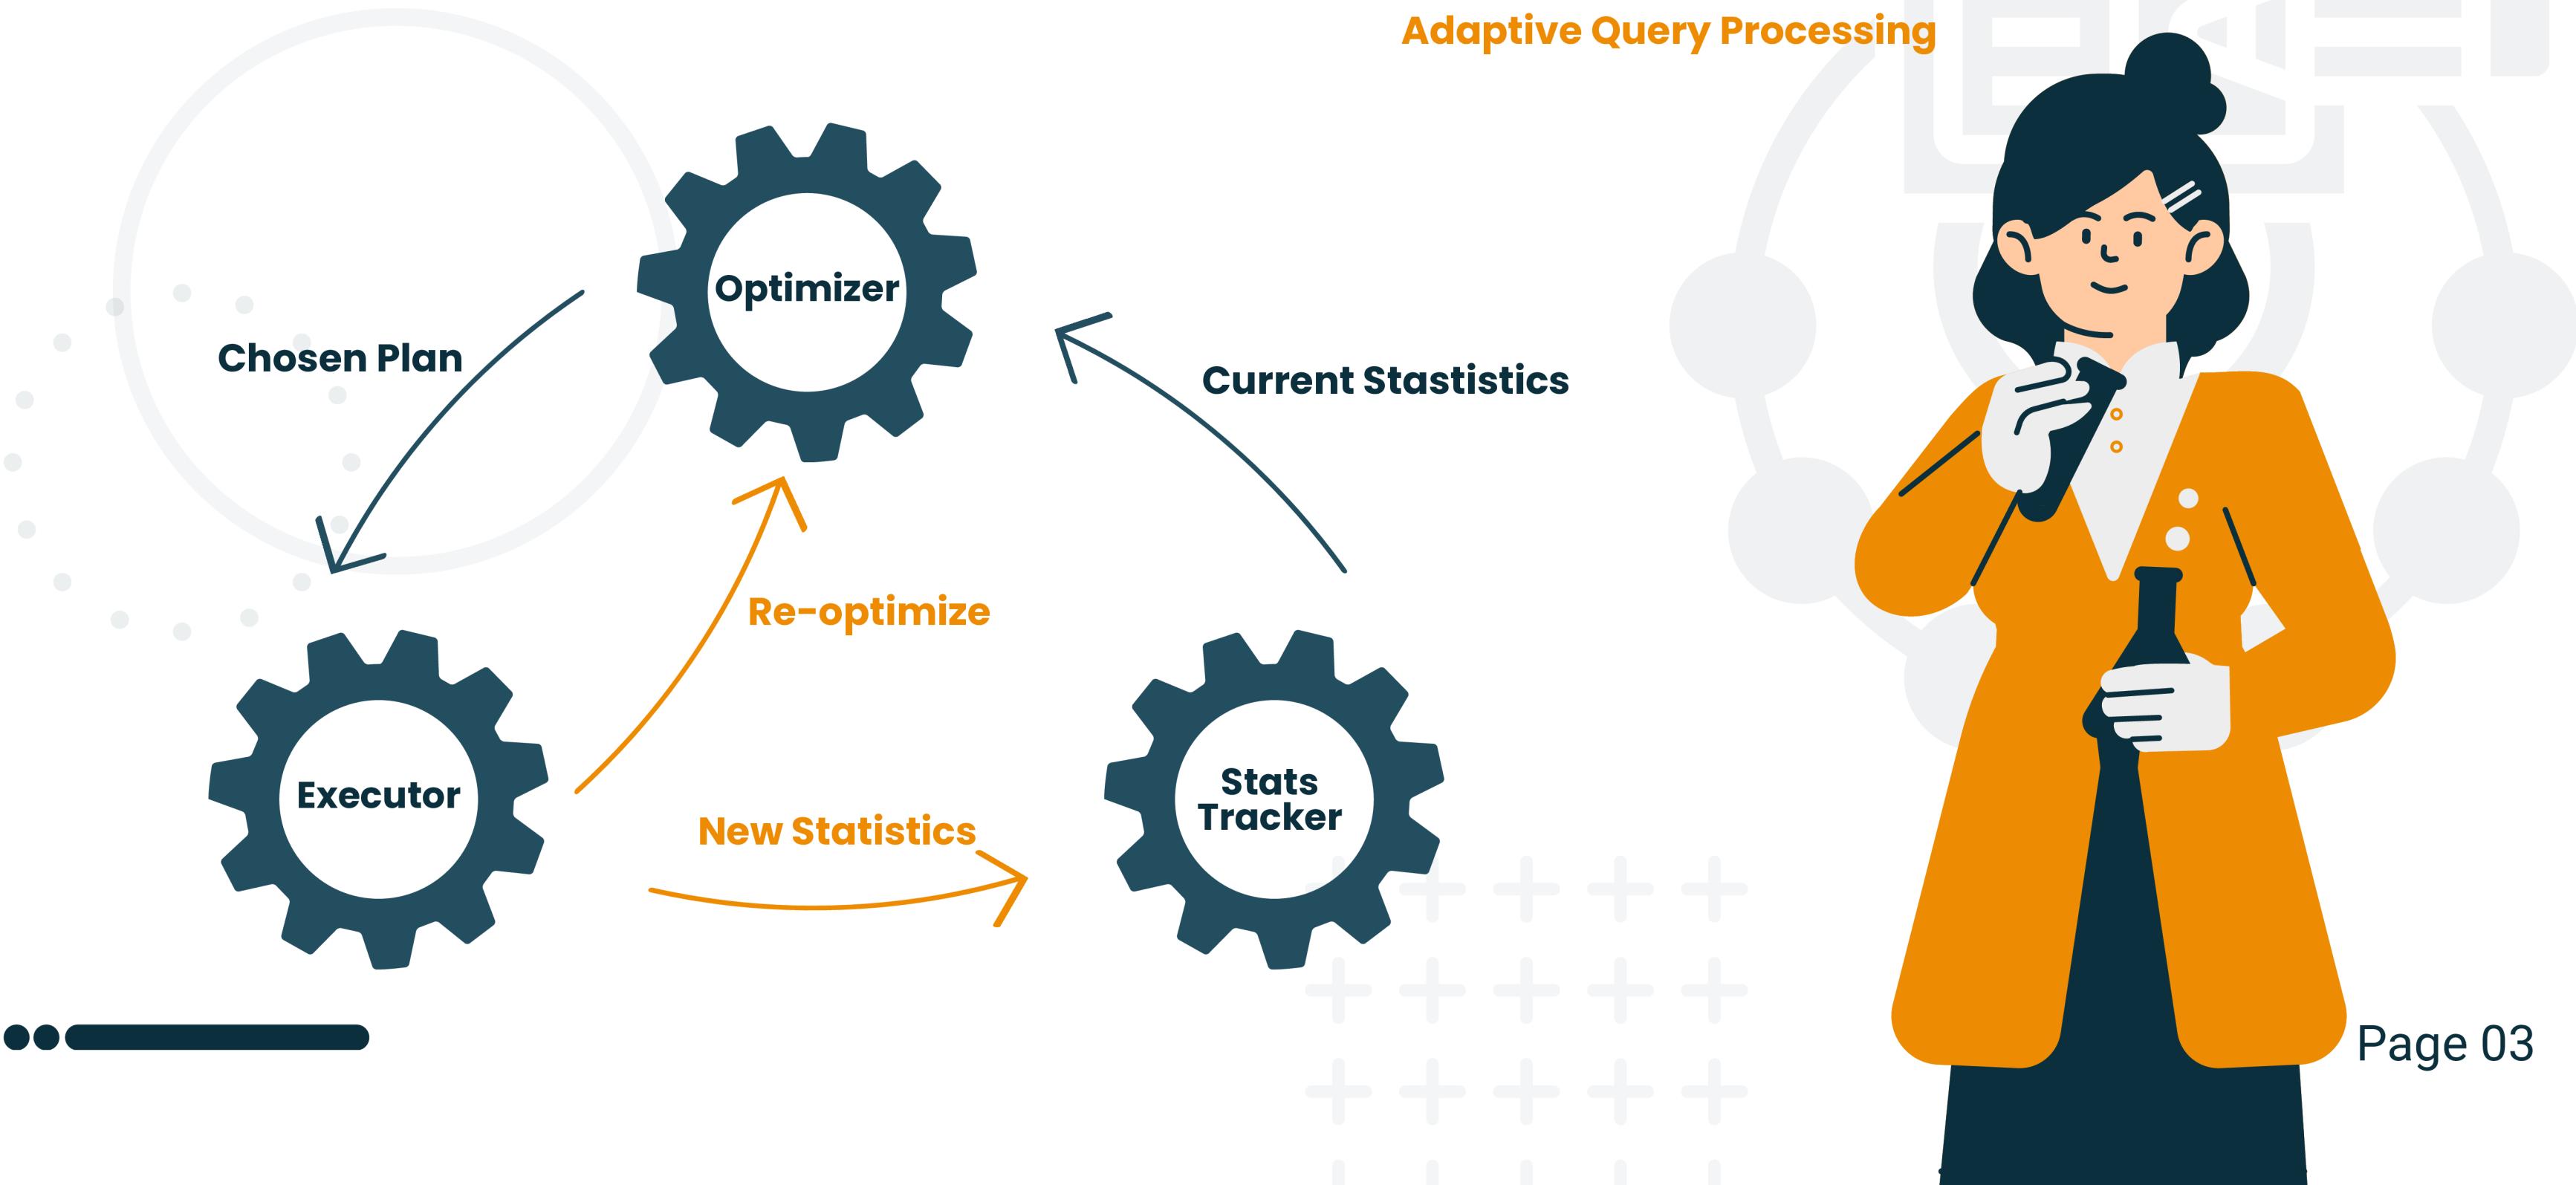
Comparison

04

Conclusion & Future Direction



Introduction





Why AQP

01

Unreliable cardinality estimation

02

Changing data and system resources

03

Error propagation in complex queries

04

Queries in online environment



Related Works

- 
- 
- 01 Plan-based**
 - 02 Routing-based**
 - 03 Continuous Query-based**
 - 04 Reinforcement Learning-based**



Plan-based



Parametric Query Optimization

- Precompute multiple plans for different situations and choose the best plan at run-time
- Examples
 - PQO (1994)
 - PPQO (2008)



Re-Optimization

- Use the optimizer to generate a better plan during query execution
- Examples
 - Re-Opt (1998)
 - POP (2024)
 - QuerySplit (2023)

Competition

- Run multiple plans in parallel and selects the most promising one
- Examples
 - DEC Rdb (1996)





Plan-based



Re-Optimization Frequency



Performance Overhead



Scalability

Competition

Low: one decision per table

High overhead: run multiple plans in parallel

Low scalability

Re-Optimization

Re-Opt: Low
POP: High
QuerySplit: Proactive

Re-Opt: Low
POP: High
QuerySplit: Moderate

Re-Opt: Low
POP: Low
QuerySplit: High

Parametric Query Optimization

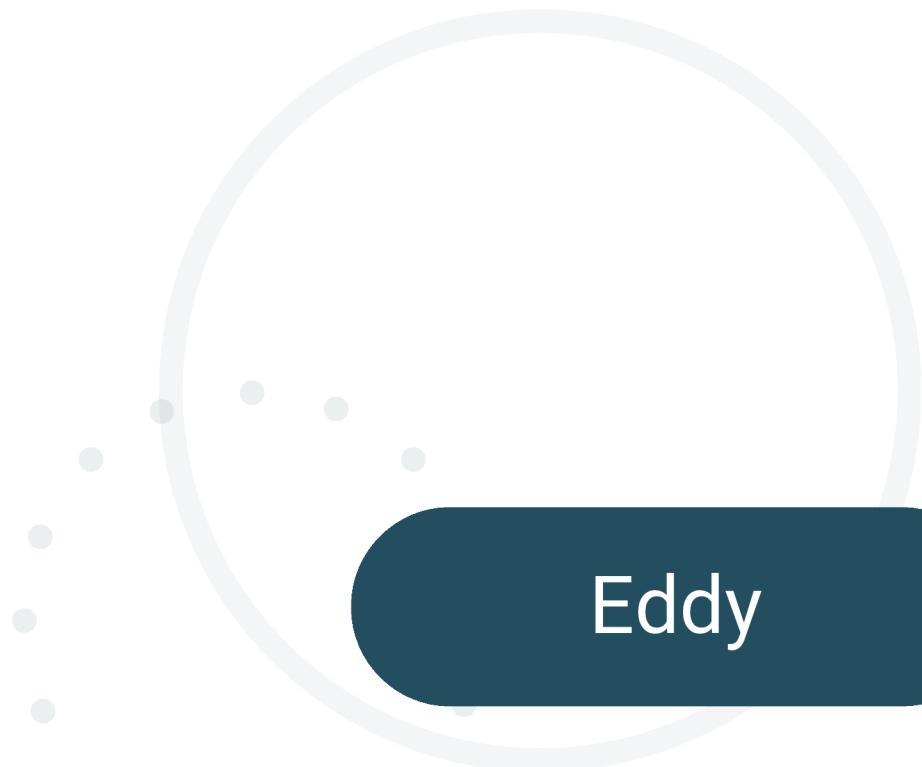
PQO: Depends on condition change
PPQO: Lower than PQO

PQO: Low during execution
PPQO: Lower than PQO

PQO: Low due to high start-up costs
PPQO: Moderate due to plan reuse



Routing-based

A large light gray circle with several smaller gray dots inside, representing a complex routing space.

Eddy



SteMs



STAIRs

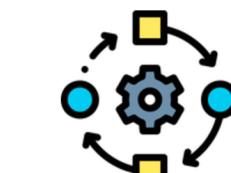
- (+) Dynamic operator ordering
- (-) Sub-optimal routing decisions under **constraints on routing history**

- (+) Improved join performance
- (+) Independence on intermediate state
- (-) Constrained plan choices
- (-) Re-computation of intermediate tuples

- (+) Flexible storage and reuse of intermediate results



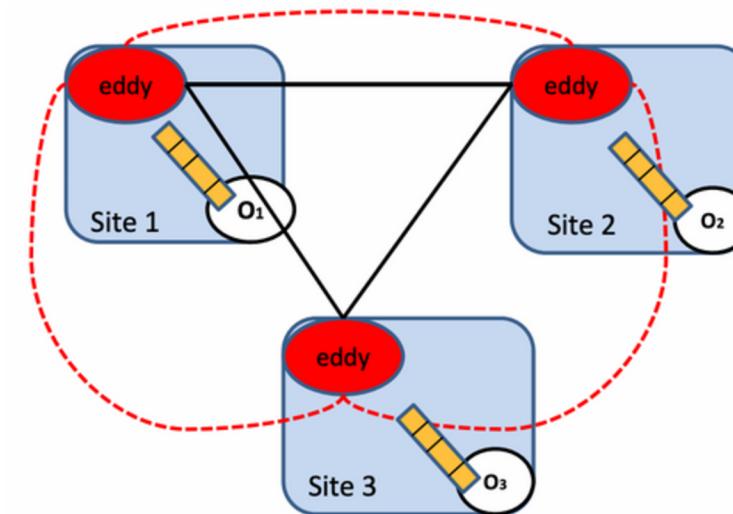
Routing-based

| |  Execution Time |  State Handling |  Adaptivity |
|--------|--|--|--|
| Eddy | Fast | Accumulates state in operators | Moderate |
| SteMs | Slower | No storage of intermediate results | Moderate |
| STAIRs | Faster for natural horizontal partitioning data | Storing long-living intermediate tuples | Higher, allows correction of earlier routing decisions |

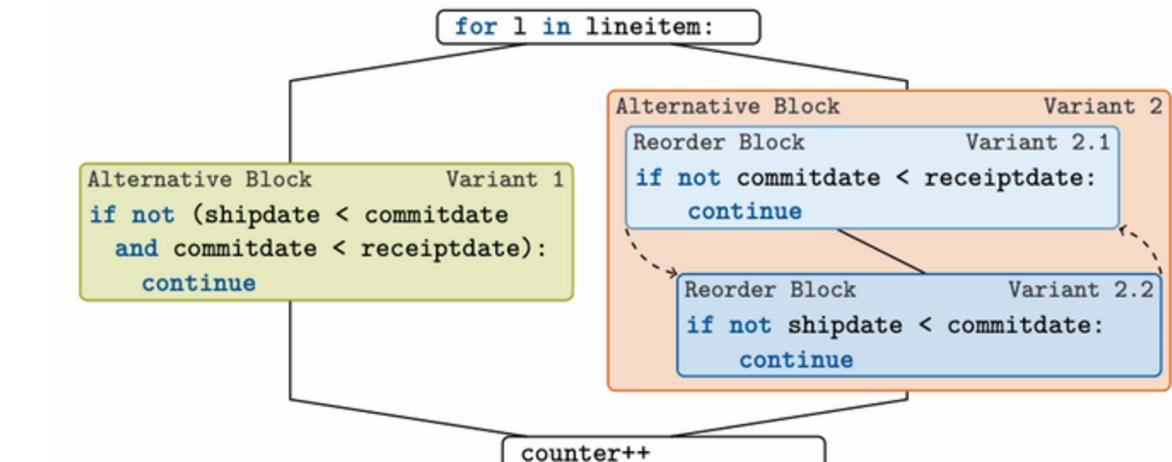


Routing-based

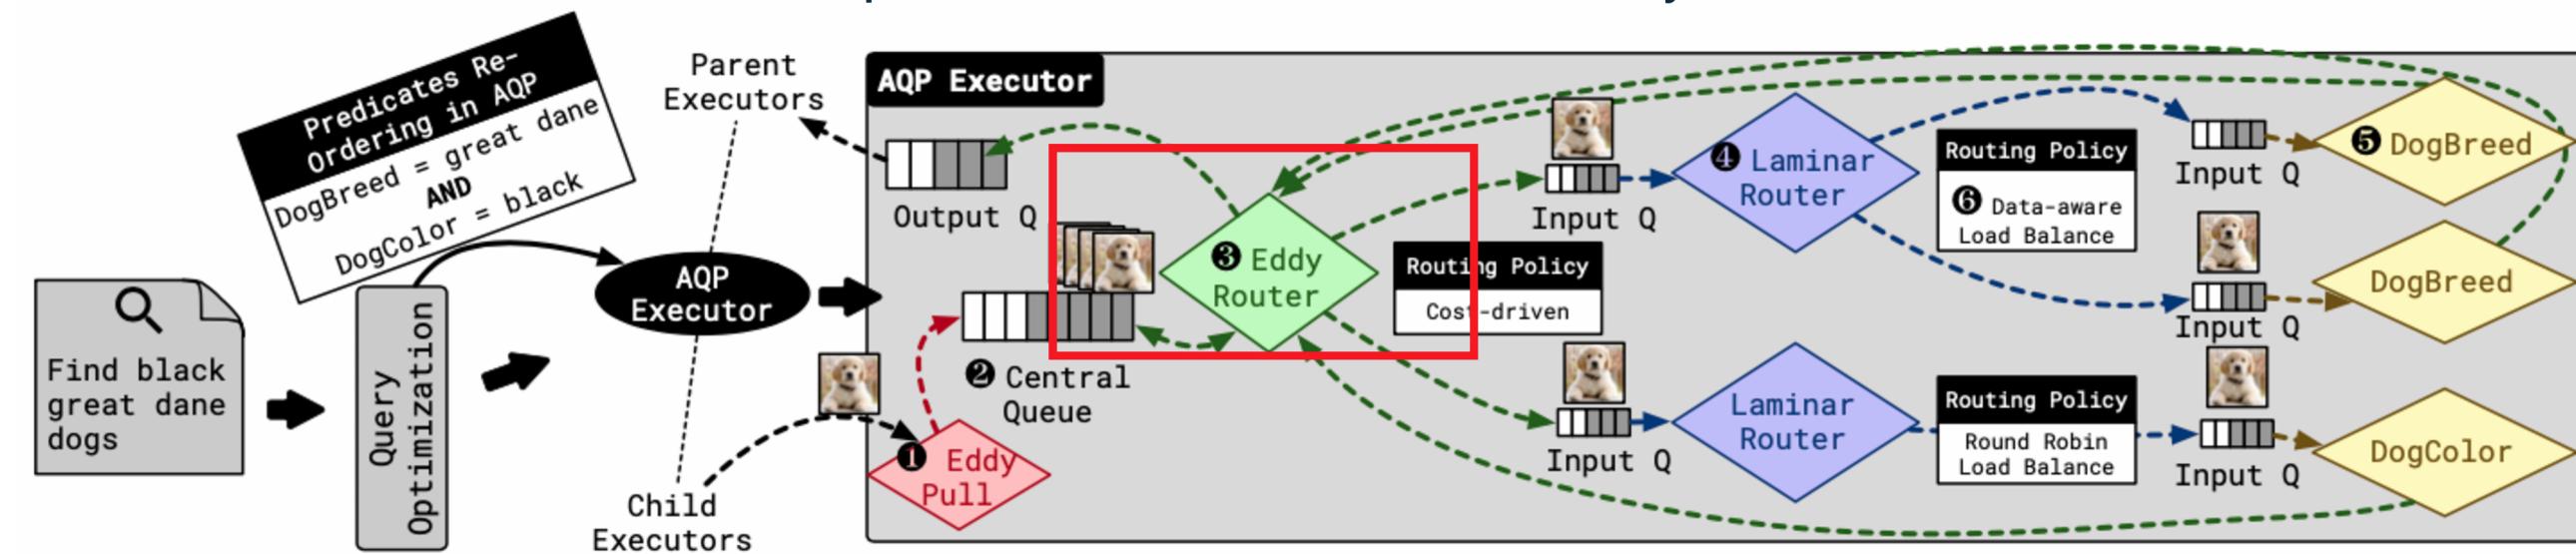
Eddy Architecture under Distributed Setting



Compiling Database Systems



Optimization on ML-centric Query





Continuous Query-based

Pure CQ-based

- Use Stream data characteristics
- Dynamically choose optimal plan
- Query sharing
- Examples
 - NiagaraCQ (2000)
 - CAPE (2004)
 - StreaMon(2004)

Eddy + CQ-based

- Eddy based dynamic routing
- UC Berkeley's Telegraph Projects
- Examples
 - CACQ (2002)
 - Psoup (2003)
 - TelegraphCQ(2003)



Continuous Query-based



Join Optimization

NiagaraCQ

Use **join signatures** to share join execution across multiple queries



Computational Cost

(+) **Cache** group query plans and recently accessed files



Performance Evaluation

(+) Works well with **overlapped queries**

(-) Drops in **multiple complex queries**, especially in UDF scenarios

CACQ

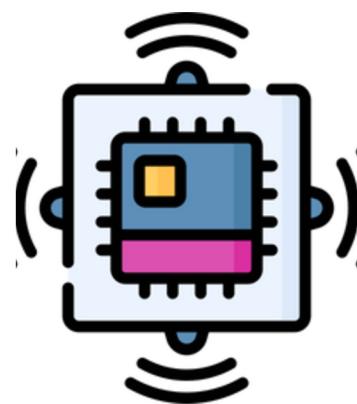
Uses **Stems** to enable pipelined multiway joins

(+) Use **grouped filter** to reduce computation

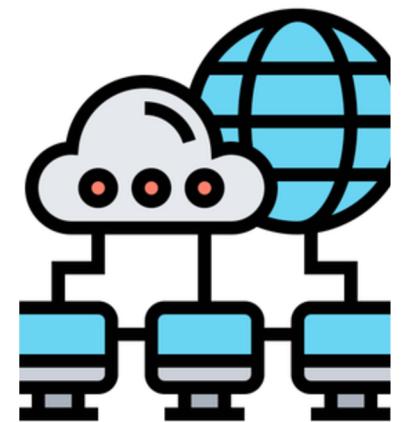
Works well with both scenarios



Continuous Query-based



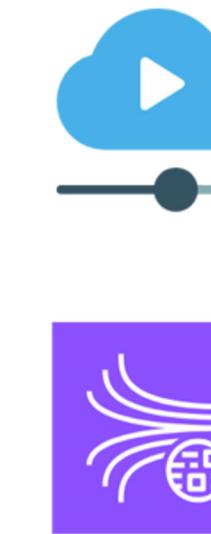
Sensor
networks



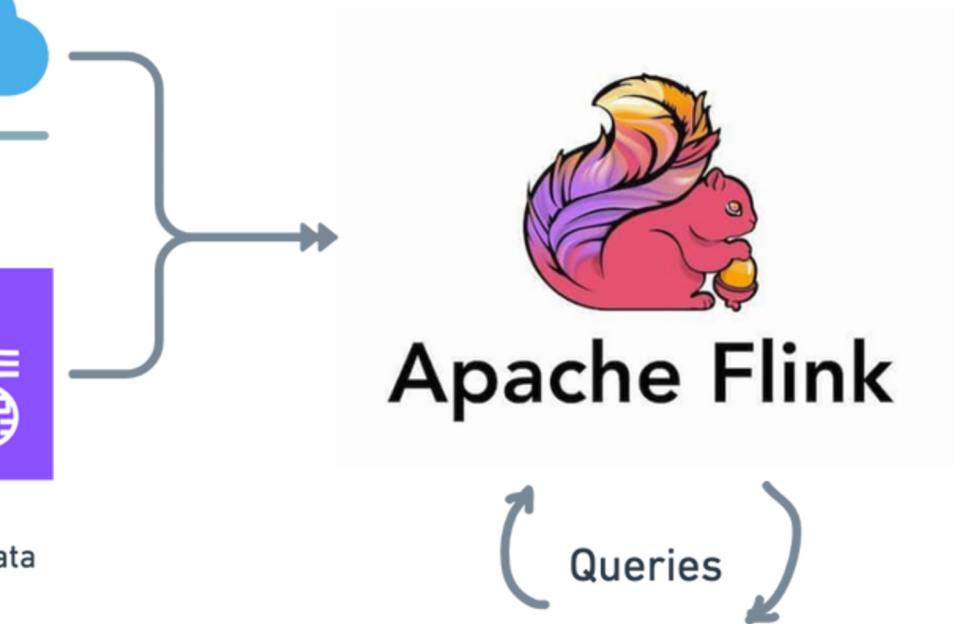
Network
monitors



Stock price
alerts



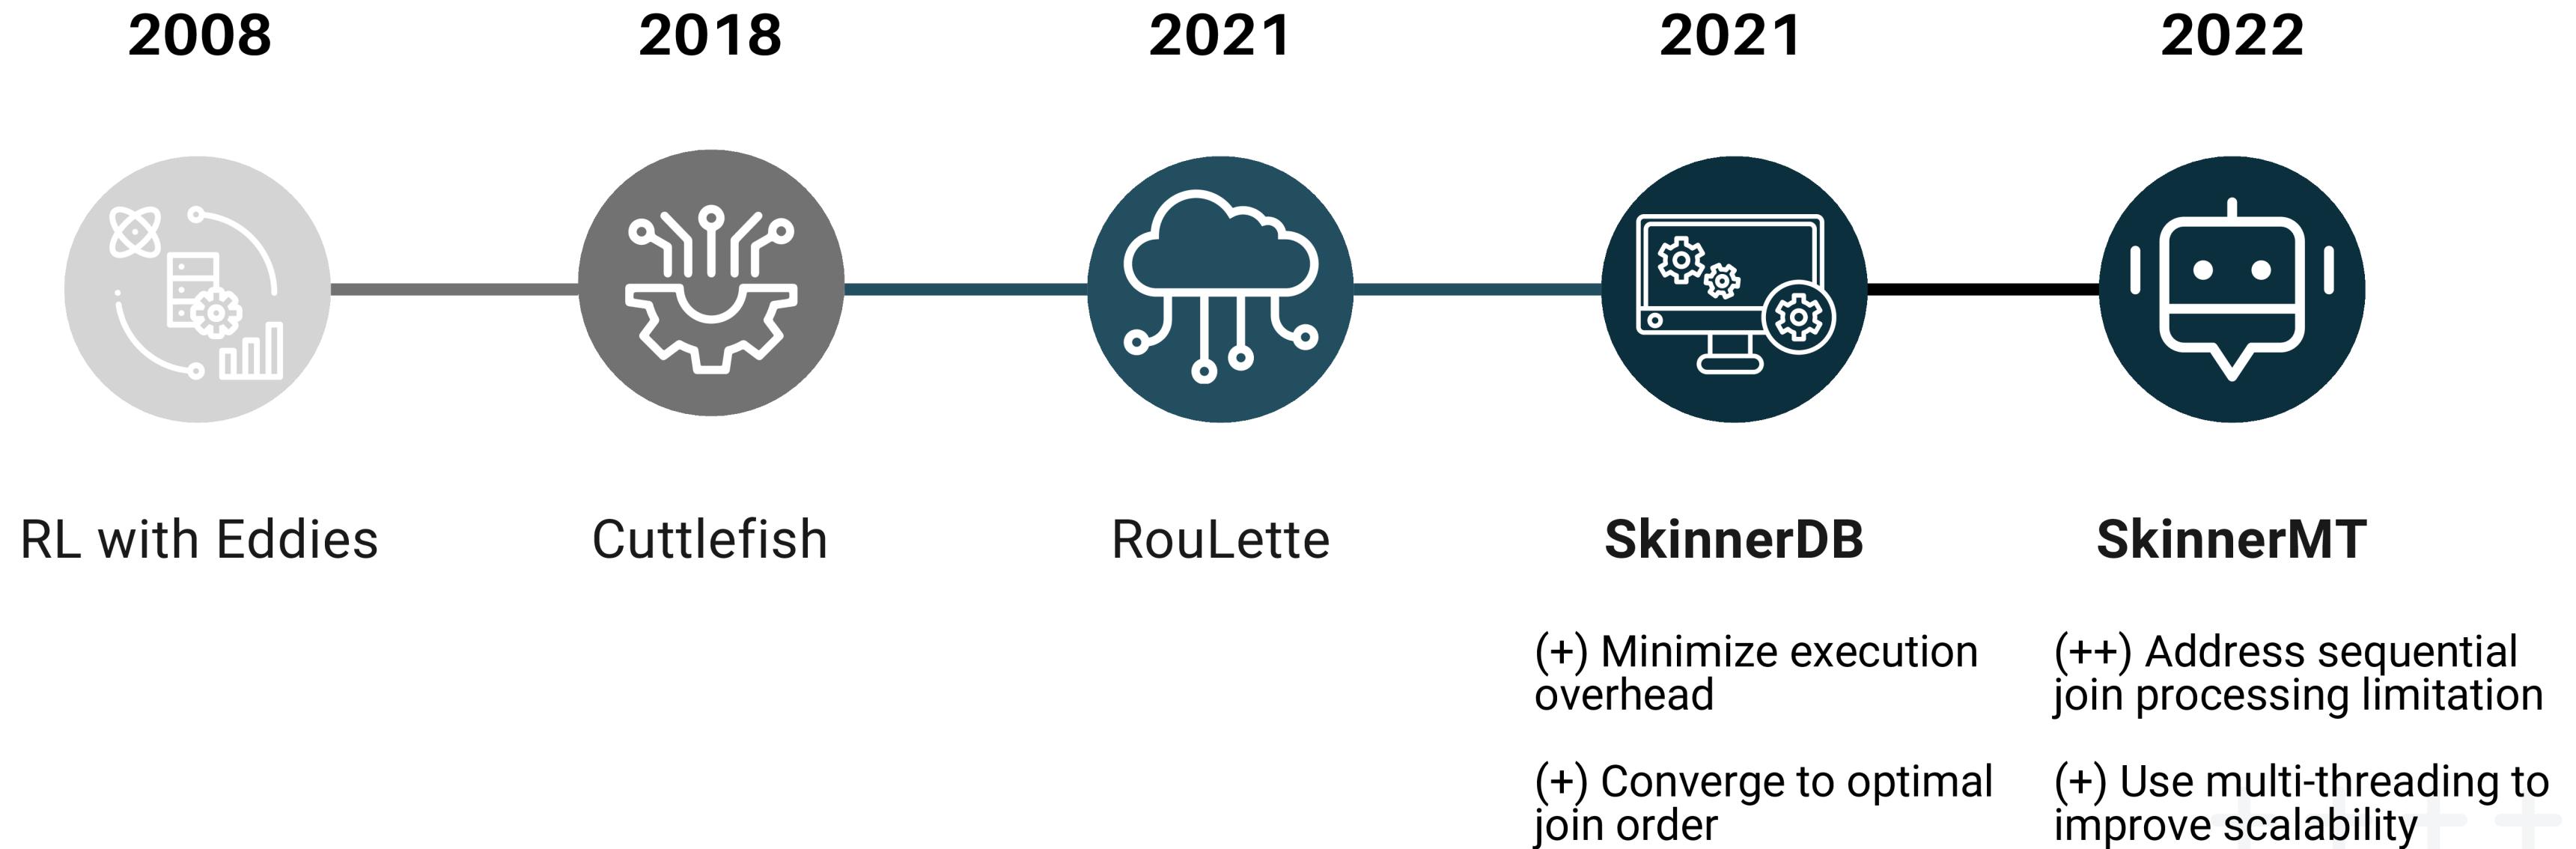
Stream Data



Nowadays



Reinforcement Learning-based





Cross Category Comparison

Plan-based & CQ-based

Prediction and algorithm of statistics

Using cost modeller for estimation

Need to determine and choose re-optimization plans

vs

Information collection

Cost estimation

Plan choosing

Routing-based & RL-based

Information gathered during exploration

Using greedy approach for cost estimation

Re-optimization of plan happens automatically



Conclusion & Future Direction

Scalability

- Trade-off between granularity and runtime overhead
- Support distributed systems

Standardized Benchmark

- Build standardized benchmark for system comparison

Security and Reliability

- Support disaster recovery and fault tolerance

Evolution of Machine Learning and AI

- Incorporate ML and AI algorithms to develop intelligent query processing systems



THANKS

FOR YOUR ATTENTION





Reference

- [1] P. G. Selinger, M. M. Astrahan, D. D. Chamberlin, R. A. Lorie, and T. G. Price, "Access path selection in a relational database management system," in Proceedings of the 1979 ACM SIGMOD international conference on Management of data, pp. 23–34, 1979.
- [2] M. Stillger, G. M. Lohman, V. Markl, and M. Kandil, "Leo-db2's learning optimizer," in VLDB, vol. 1, pp. 19–28, 2001.
- [3] R. Motwani, J. Widom, A. Arasu, B. Babcock, S. Babu, M. Datar, G. Manku, C. Olston, J. Rosenstein, and R. Varma, "Query processing, resource management, and approximation in a data stream management system," in CIDR 2003, Stanford InfoLab, 2002.
- [4] Y. E. Ioannidis and S. Christodoulakis, "On the propagation of errors in the size of join results," in Proceedings of the 1991 ACM SIGMOD International Conference on Management of data, pp. 268–277, 1991.
- [5] J. M. Hellerstein, R. Avnur, A. Chou, C. Hidber, C. Olston, V. Raman, T. Roth, and P. J. Haas, "Interactive data analysis: The control project," Computer, vol. 32, no. 8, pp. 51–59, 1999.
- [6] Z. G. Ives, A. Y. Levy, D. S. Weld, D. Florescu, and M. Friedman, "Adaptive query processing for internet applications," 2000.
- [7] S. Babu and P. Bizarro, "Adaptive query processing in the looking glass," in Proceedings of the Second Biennial Conference on Innovative Data Systems Research (CIDR), Jan. 2005.
- [8] T. Schmidt, P. Fent, and T. Neumann, "Efficiently compiling dynamic code for adaptive query processing," in ADMS@ VLDB, pp. 11–22, 2022.
- [9] N. Kabra and D. J. DeWitt, "Efficient mid-query re-optimization of sub-optimal query execution plans," in Proceedings of the 1998 ACM SIGMOD international conference on Management of data, pp. 106–117, 1998.
- [10] V. Markl, V. Raman, D. Simmen, G. Lohman, H. Pirahesh, and M. Cilimdzic, "Robust query processing through progressive optimization," in Proceedings of the 2004 ACM SIGMOD international conference on Management of data, pp. 659–670, 2004.
- [11] J. Zhao, H. Zhang, and Y. Gao, "Efficient query re-optimization with judicious subquery selections," Proceedings of the ACM on Management of Data, vol. 1, no. 2, pp. 1–26, 2023.
- [12] R. L. Cole and G. Graefe, "Optimization of dynamic query evaluation plans," in Proceedings of the 1994 ACM SIGMOD international conference on Management of data, pp. 150–160, 1994.
- [13] P. Bizarro, N. Bruno, and D. J. DeWitt, "Progressive parametric query optimization," IEEE Transactions on Knowledge and Data Engineering, vol. 21, no. 4, pp. 582–594, 2008.
- [14] G. Antoshenkov and M. Ziauddin, "Query processing and optimization in oracle rdb," The VLDB Journal, vol. 5, pp. 229–237, 1996.
- [15] A. Deshpande, Z. Ives, V. Raman, et al., "Adaptive query processing," Foundations and Trends® in Databases, vol. 1, no. 1, pp. 1–140, 2007.
- [16] R. Avnur and J. M. Hellerstein, "Eddies: Continuously adaptive query processing," in Proceedings of the 2000 ACM SIGMOD international conference on Management of data, pp. 261–272, 2000.
- [17] V. Raman, A. Deshpande, and J. Hellerstein, "Using state modules for adaptive query processing," in Proceedings 19th International Conference on Data Engineering (Cat. No.03CH37405), pp. 353–364, 2003.
- [18] A. Deshpande and J. M. Hellerstein, "Lifting the burden of history from adaptive query processing," in Proceedings of the Thirtieth International Conference on Very Large Data Bases - Volume 30, VLDB '04, p. 948–959, VLDB Endowment, 2004.
- [19] B. Răducanu, P. Boncz, and M. Zukowski, "Micro adaptivity in vectorwise," in Proceedings of the 2013 ACM SIGMOD International Conference on Management of Data, SIGMOD '13, (New York, NY, USA), p. 1231–1242, Association for Computing Machinery, 2013.
- [20] D. Terry, D. Goldberg, D. Nichols, and B. Oki, "Continuous queries over append-only databases," Acm Sigmod Record, vol. 21, no. 2, pp. 321–330, 1992.
- [21] S. Madden, M. Shah, J. M. Hellerstein, and V. Raman, "Continuously adaptive continuous queries over streams," in Proceedings of the 2002 ACM SIGMOD international conference on Management of data, pp. 49–60, 2002.
- [22] S. Chandrasekaran and M. J. Franklin, "Psoup: a system for streaming queries over streaming data," The VLDB Journal, vol. 12, pp. 140–156, 2003.
- [23] E. A. Rundensteiner, L. Ding, T. Sutherland, Y. Zhu, B. Pielech, and N. Mehta, "Cape: Continuous query engine with heterogeneous-grained adaptivity," in Proceedings of the Thirtieth international conference on Very large data bases-Volume 30, pp. 1353–1356, 2004.
- [24] M. A. Shah, J. M. Hellerstein, S. Chandrasekaran, and M. J. Franklin, "Flux: An adaptive partitioning operator for continuous query systems," in Proceedings 19th International Conference on Data Engineering (Cat. No. 03CH37405), pp. 25–36, IEEE, 2003.
- [25] S. Krishnamurthy, S. Chandrasekaran, O. Cooper, A. Deshpande, M. J. Franklin, J. M. Hellerstein, W. Hong, S. Madden, F. Reiss, and M. A. Shah, "Telegraphcq: An architectural status report," IEEE Data Eng. Bull., vol. 26, no. 1, pp. 11–18, 2003.
- [26] J. Chen, D. J. DeWitt, F. Tian, and Y. Wang, "Niagaracq: A scalable continuous query system for internet databases," in Proceedings of the 2000 ACM SIGMOD international conference on Management of data, pp. 379–390, 2000.
- [27] S. Babu and J. Widom, "Streamon: an adaptive engine for stream query processing," in Proceedings of the 2004 ACM SIGMOD international conference on Management of data, pp. 931–932, 2004.
- [28] K. Tzoumas, T. Sellis, and C. S. Jensen, "A reinforcement learning approach for adaptive query processing," History, pp. 1–25, 2008.
- [29] T. Kaftan, M. Balazinska, A. Cheung, and J. Gehrke, "Cuttlefish: A lightweight primitive for adaptive query processing," arXiv preprint arXiv:1802.09180, 2018.
- [30] I. Trummer, J. Wang, Z. Wei, D. Maram, S. Moseley, S. Jo, J. Antonakakis, and A. Rayabhari, "Skinnerdb: Regret-bounded query evaluation via reinforcement learning," ACM Trans. Database Syst., vol. 46, sep 2021.
- [31] P. Sioulas and A. Ailamaki, "Scalable multi-query execution using reinforcement learning," in Proceedings of the 2021 International Conference on Management of Data, pp. 1651–1663, 2021.
- [32] Z. Wei and I. Trummer, "Skinnermt: Parallelizing for efficiency and robustness in adaptive query processing on multicore platforms," Proc. VLDB Endow., vol. 16, p. 905–917, dec 2022.
- [33] G. T. Kakkar, J. Cao, A. Sengupta, J. Arulraj, and H. Kim, "Hydro: Adaptive query processing of ml queries," arXiv preprint arXiv:2403.14902, 2024.
- [34] A. Gounaris, E. Tsamoura, and Y. Manolopoulos, "Adaptive query processing in distributed settings," Intelligent Systems Reference Library, vol. 36, 01 2013.

